

DDS

EXPLAINED

Angelo Corsaro, PhD

Chief Technology Officer

ADLINK Tech. Inc.

angelo.corsaro@adlinktech.com




Leading EDGE COMPUTING

What is DDS?

DDS in 131 Characters

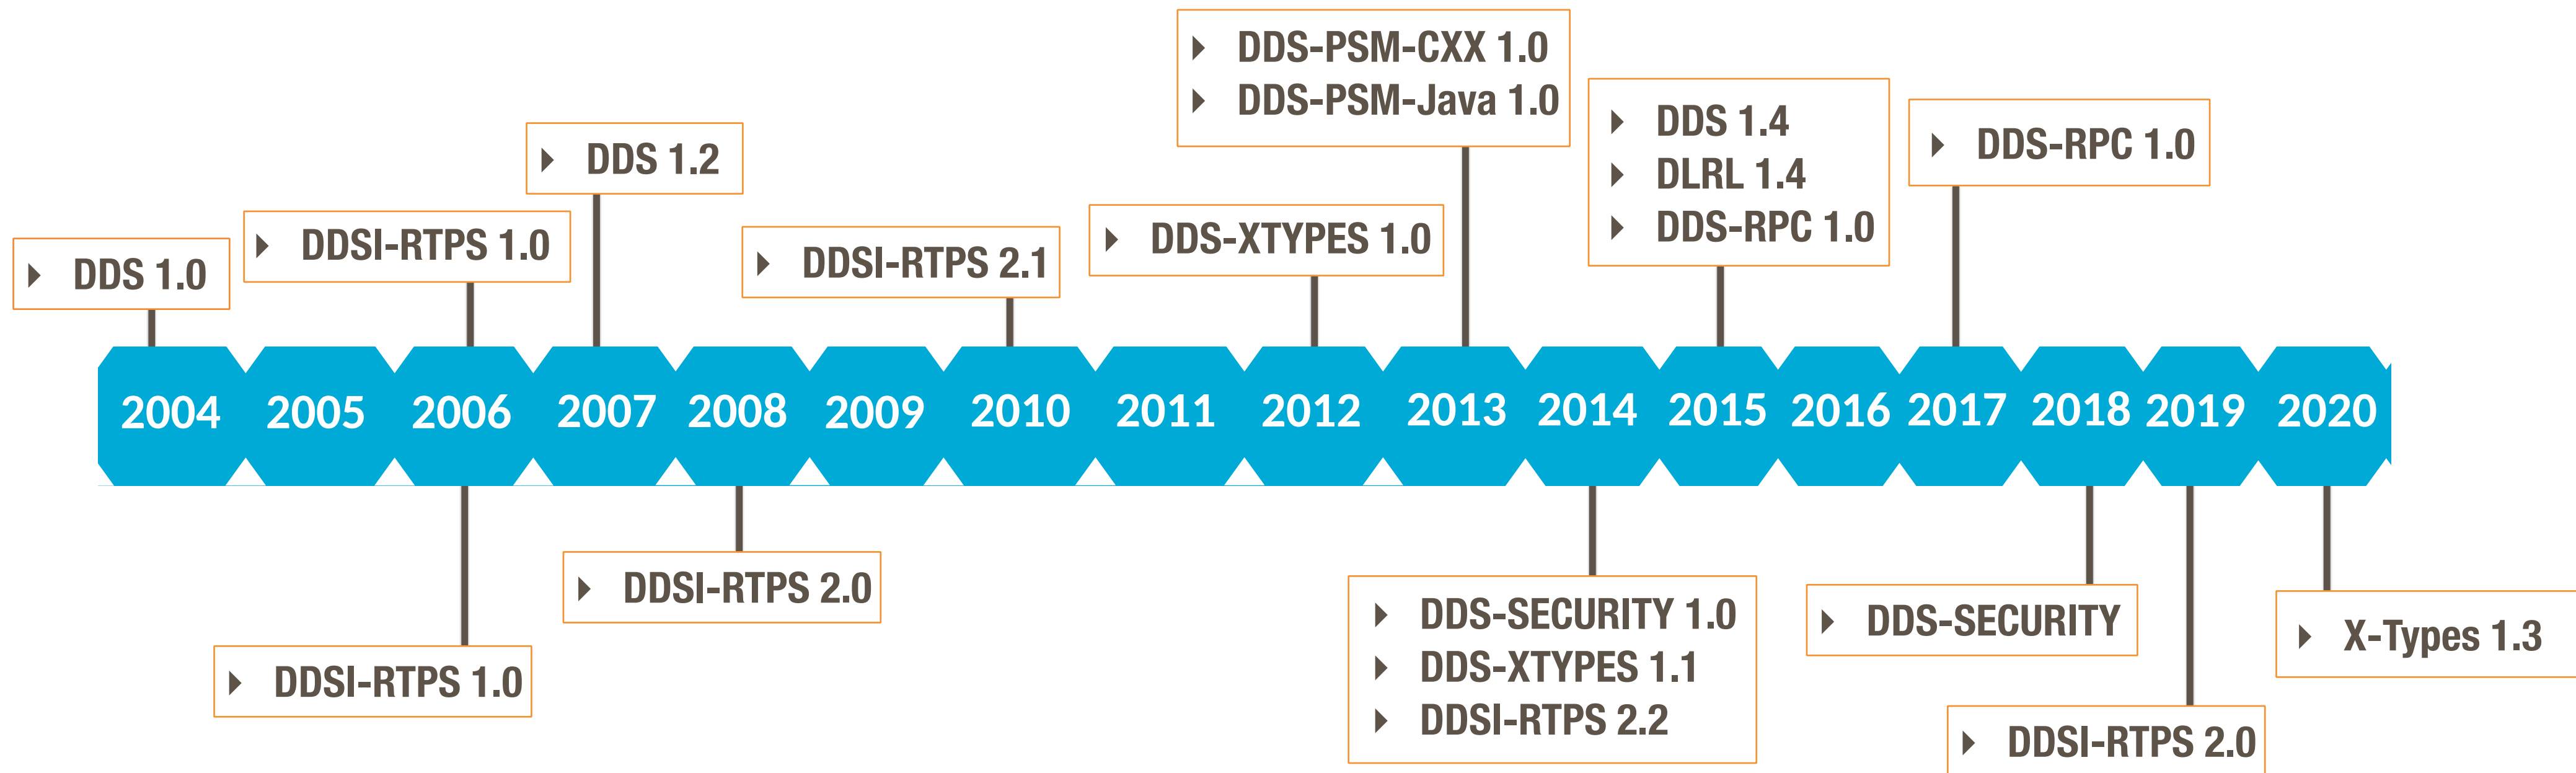
DDS is a standard technology for ubiquitous, interoperable, secure, platform independent, and real-time **data sharing** across network connected devices



The DDS Standard

Relevant Standards

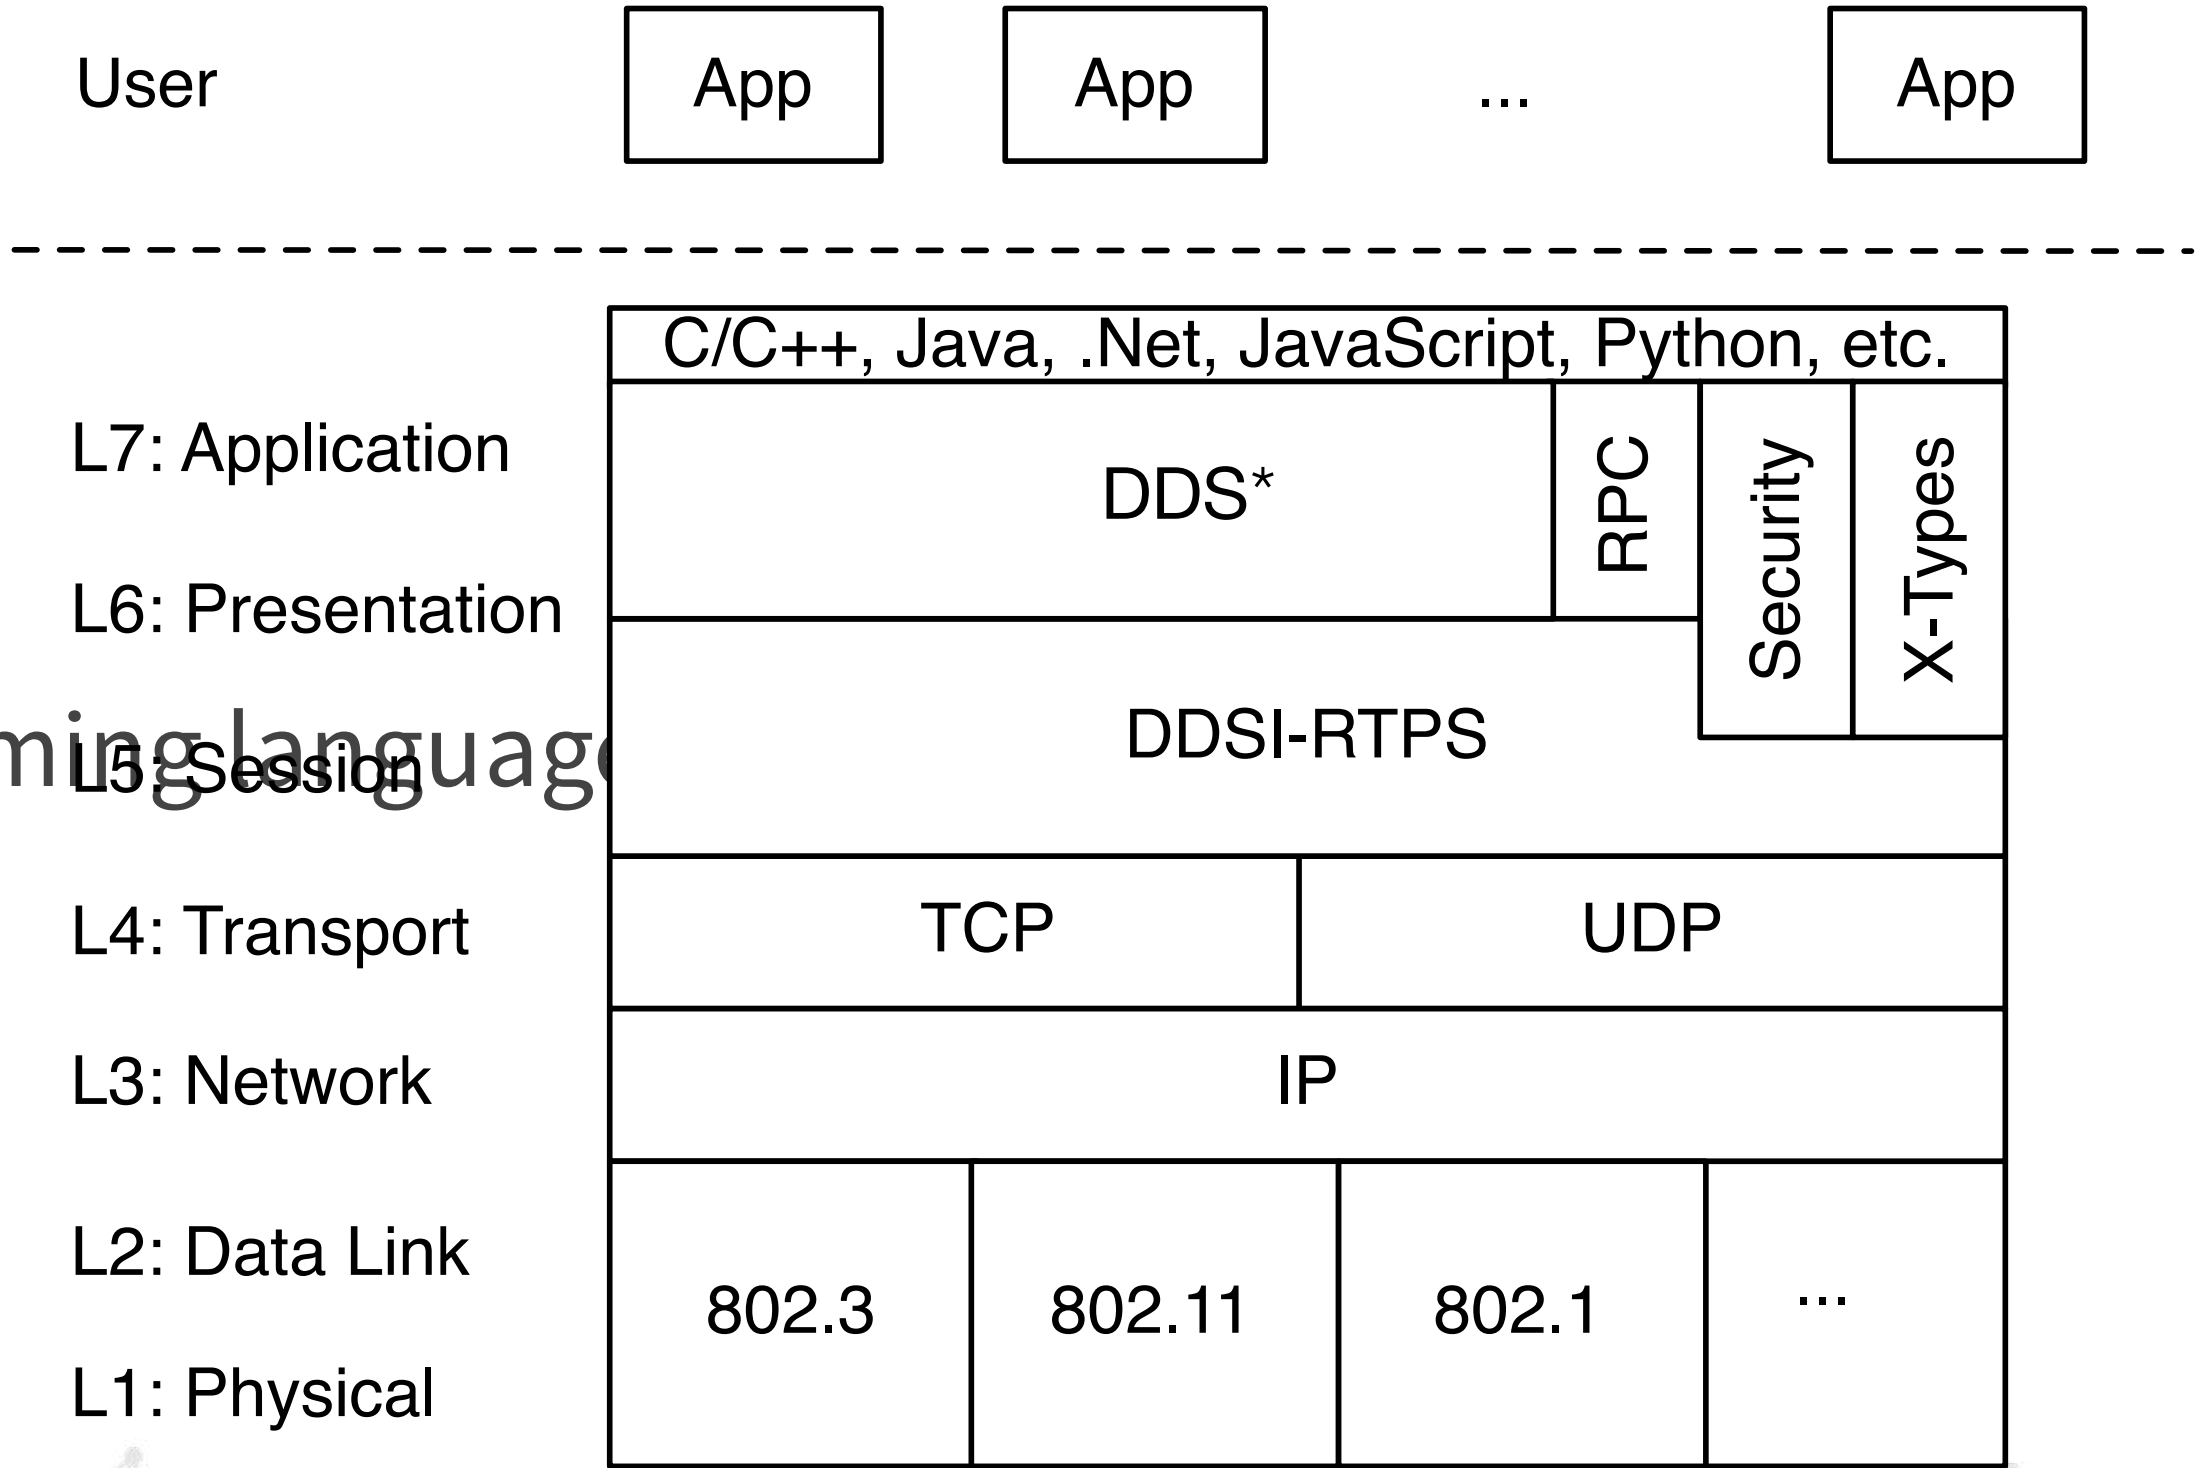
DDS Evolution



DDS Standards

Data Distribution Service (DDS)

Defines a high level API for programming language independent data sharing



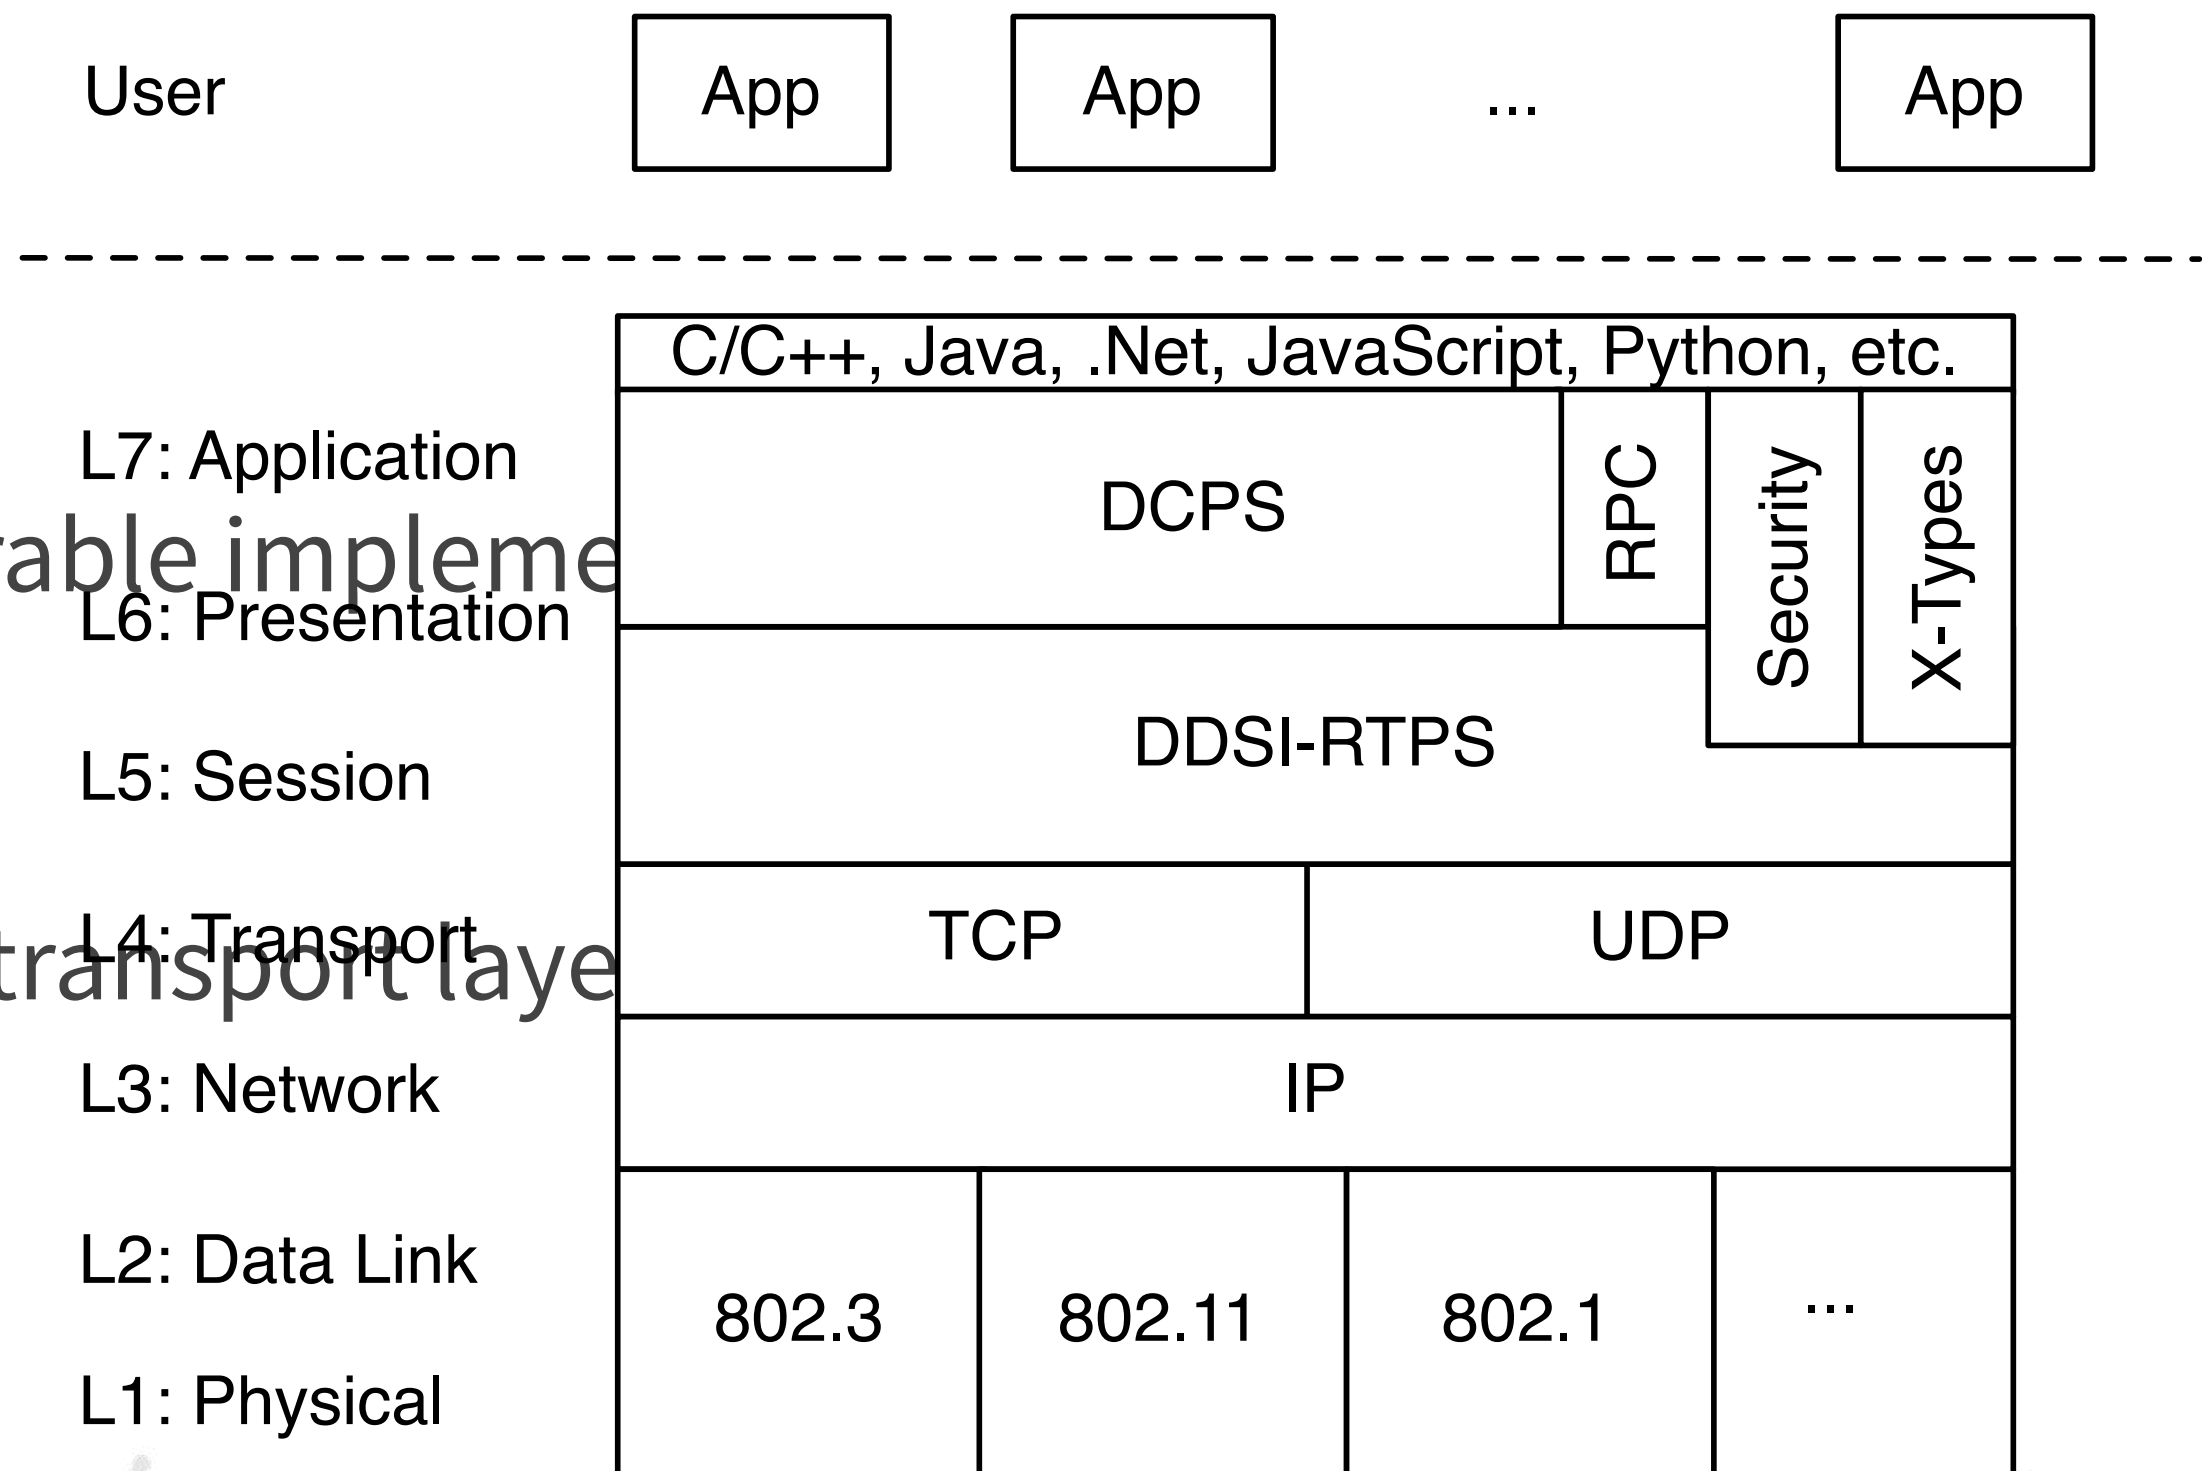
(*) This used to be called DCPS as originally the DDS standard was composed by two layers, DCPS (Data Centric Publish/Subscribe and DLRL (Data Local Reconstruction Layer).

DDS Standards

DDS Interoperability Protocol (DDSI-RTPS)

Defines a wire protocol for interoperable implementations.

This protocol assumes a best-effort transport layer provided by DDSI.



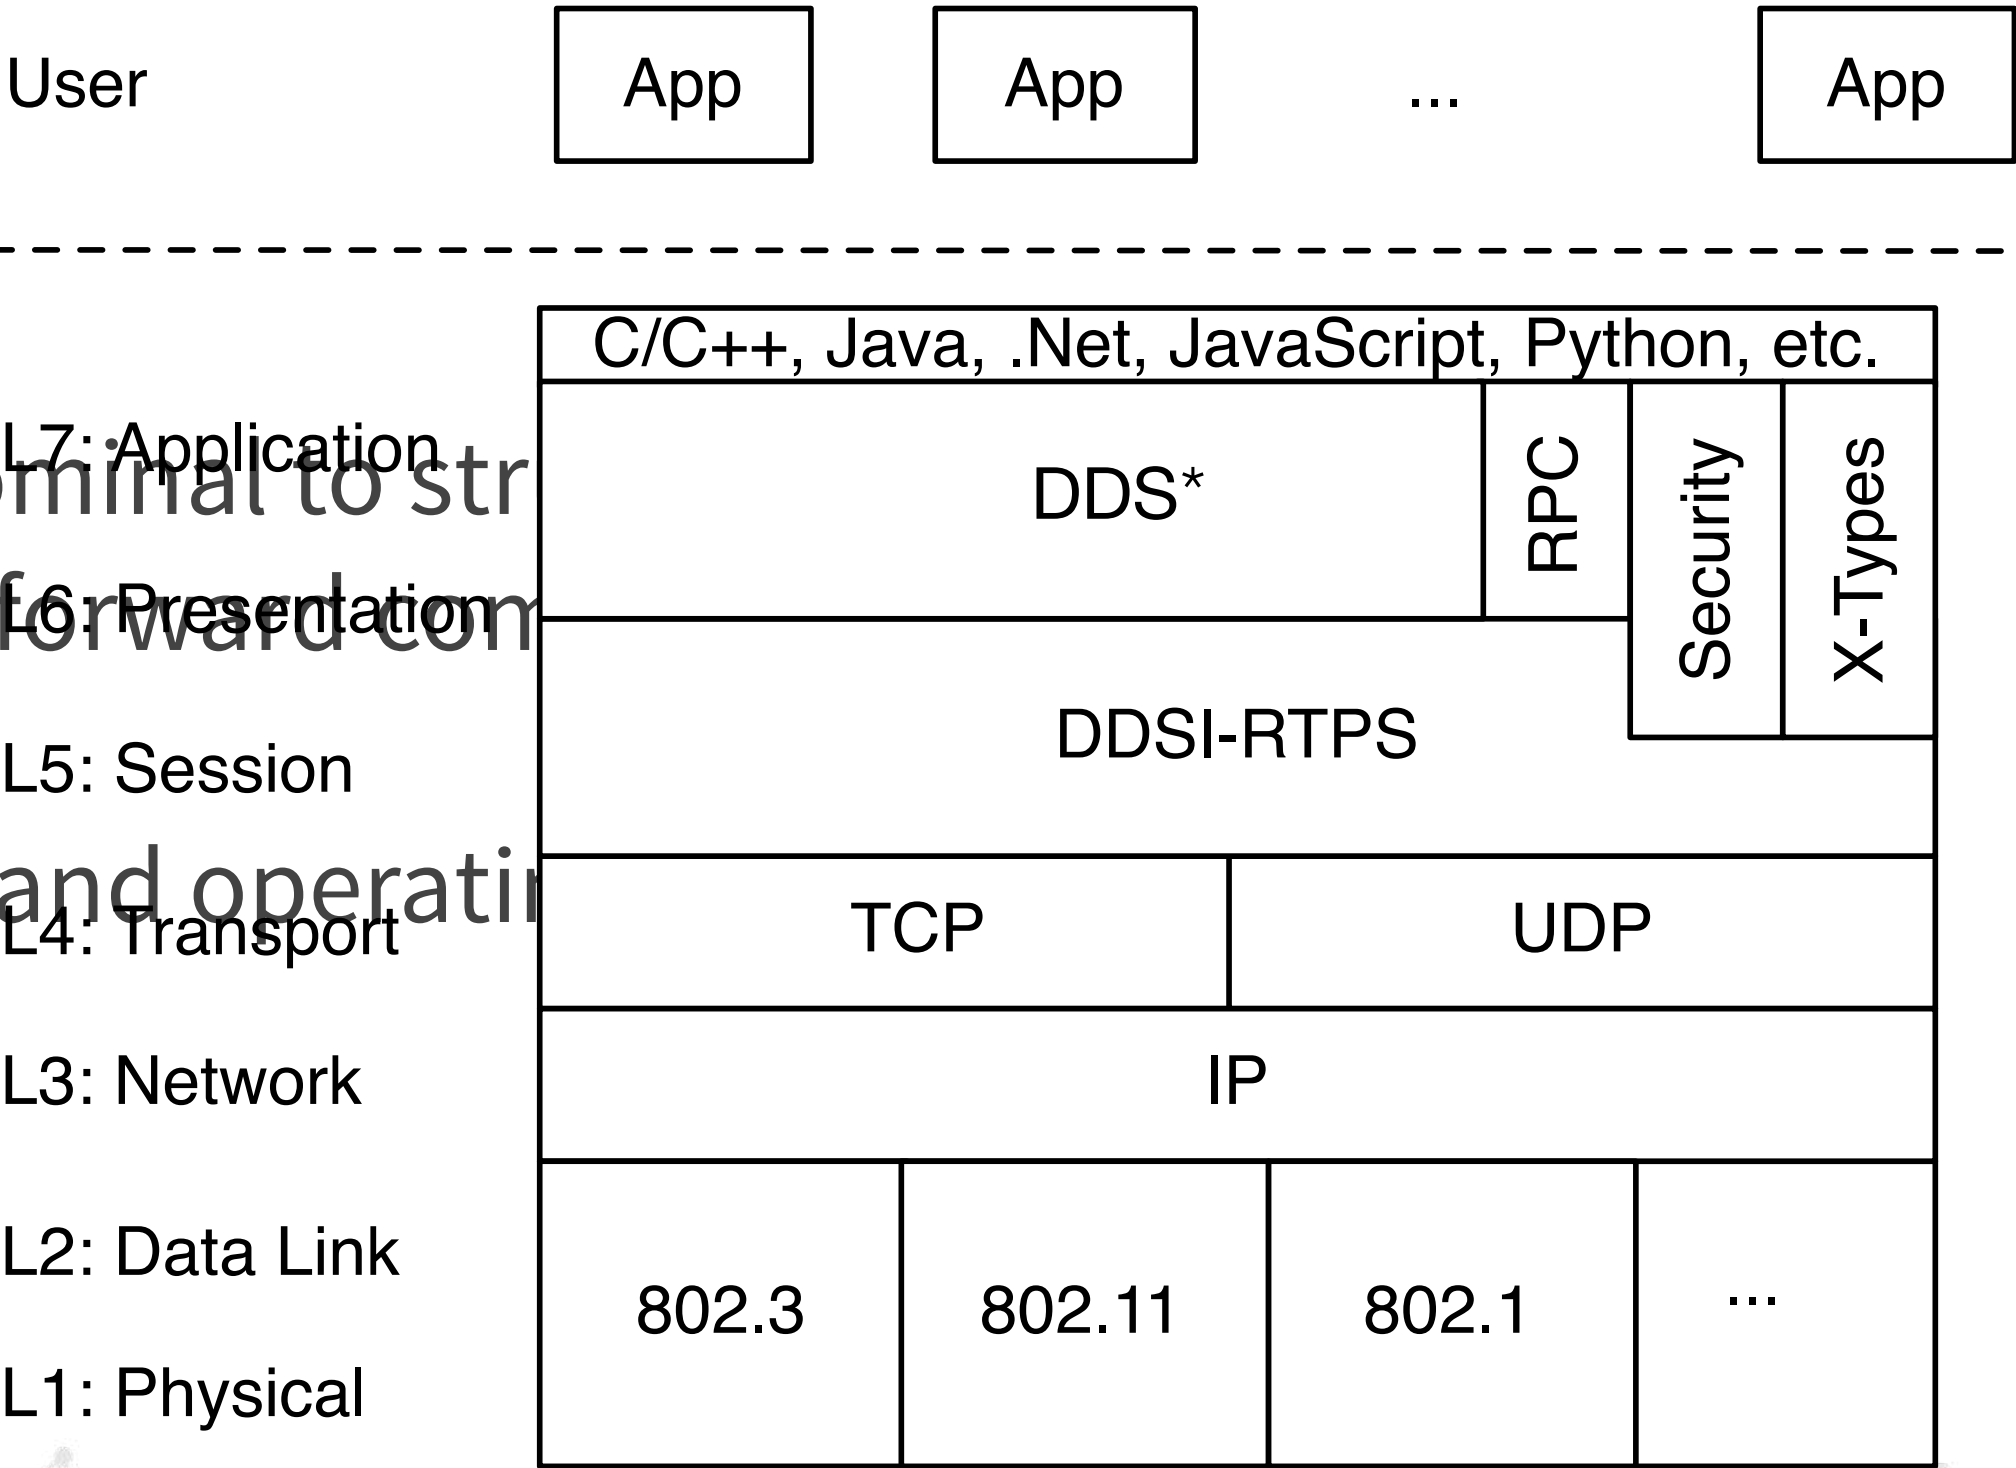
(*) This used to be called DCPS as originally the DDS standard was composed by two layers, DCPS (Data Centric Publish/Subscribe and DLRL (Data Local Reconstruction Layer).

DDS Standards

eXtensible Types (DDS-XTypes)

Extends the DDS type system from nominal to structural
very good support for evolutions and forward compatibility

Defines APIs for dynamically defining and operating

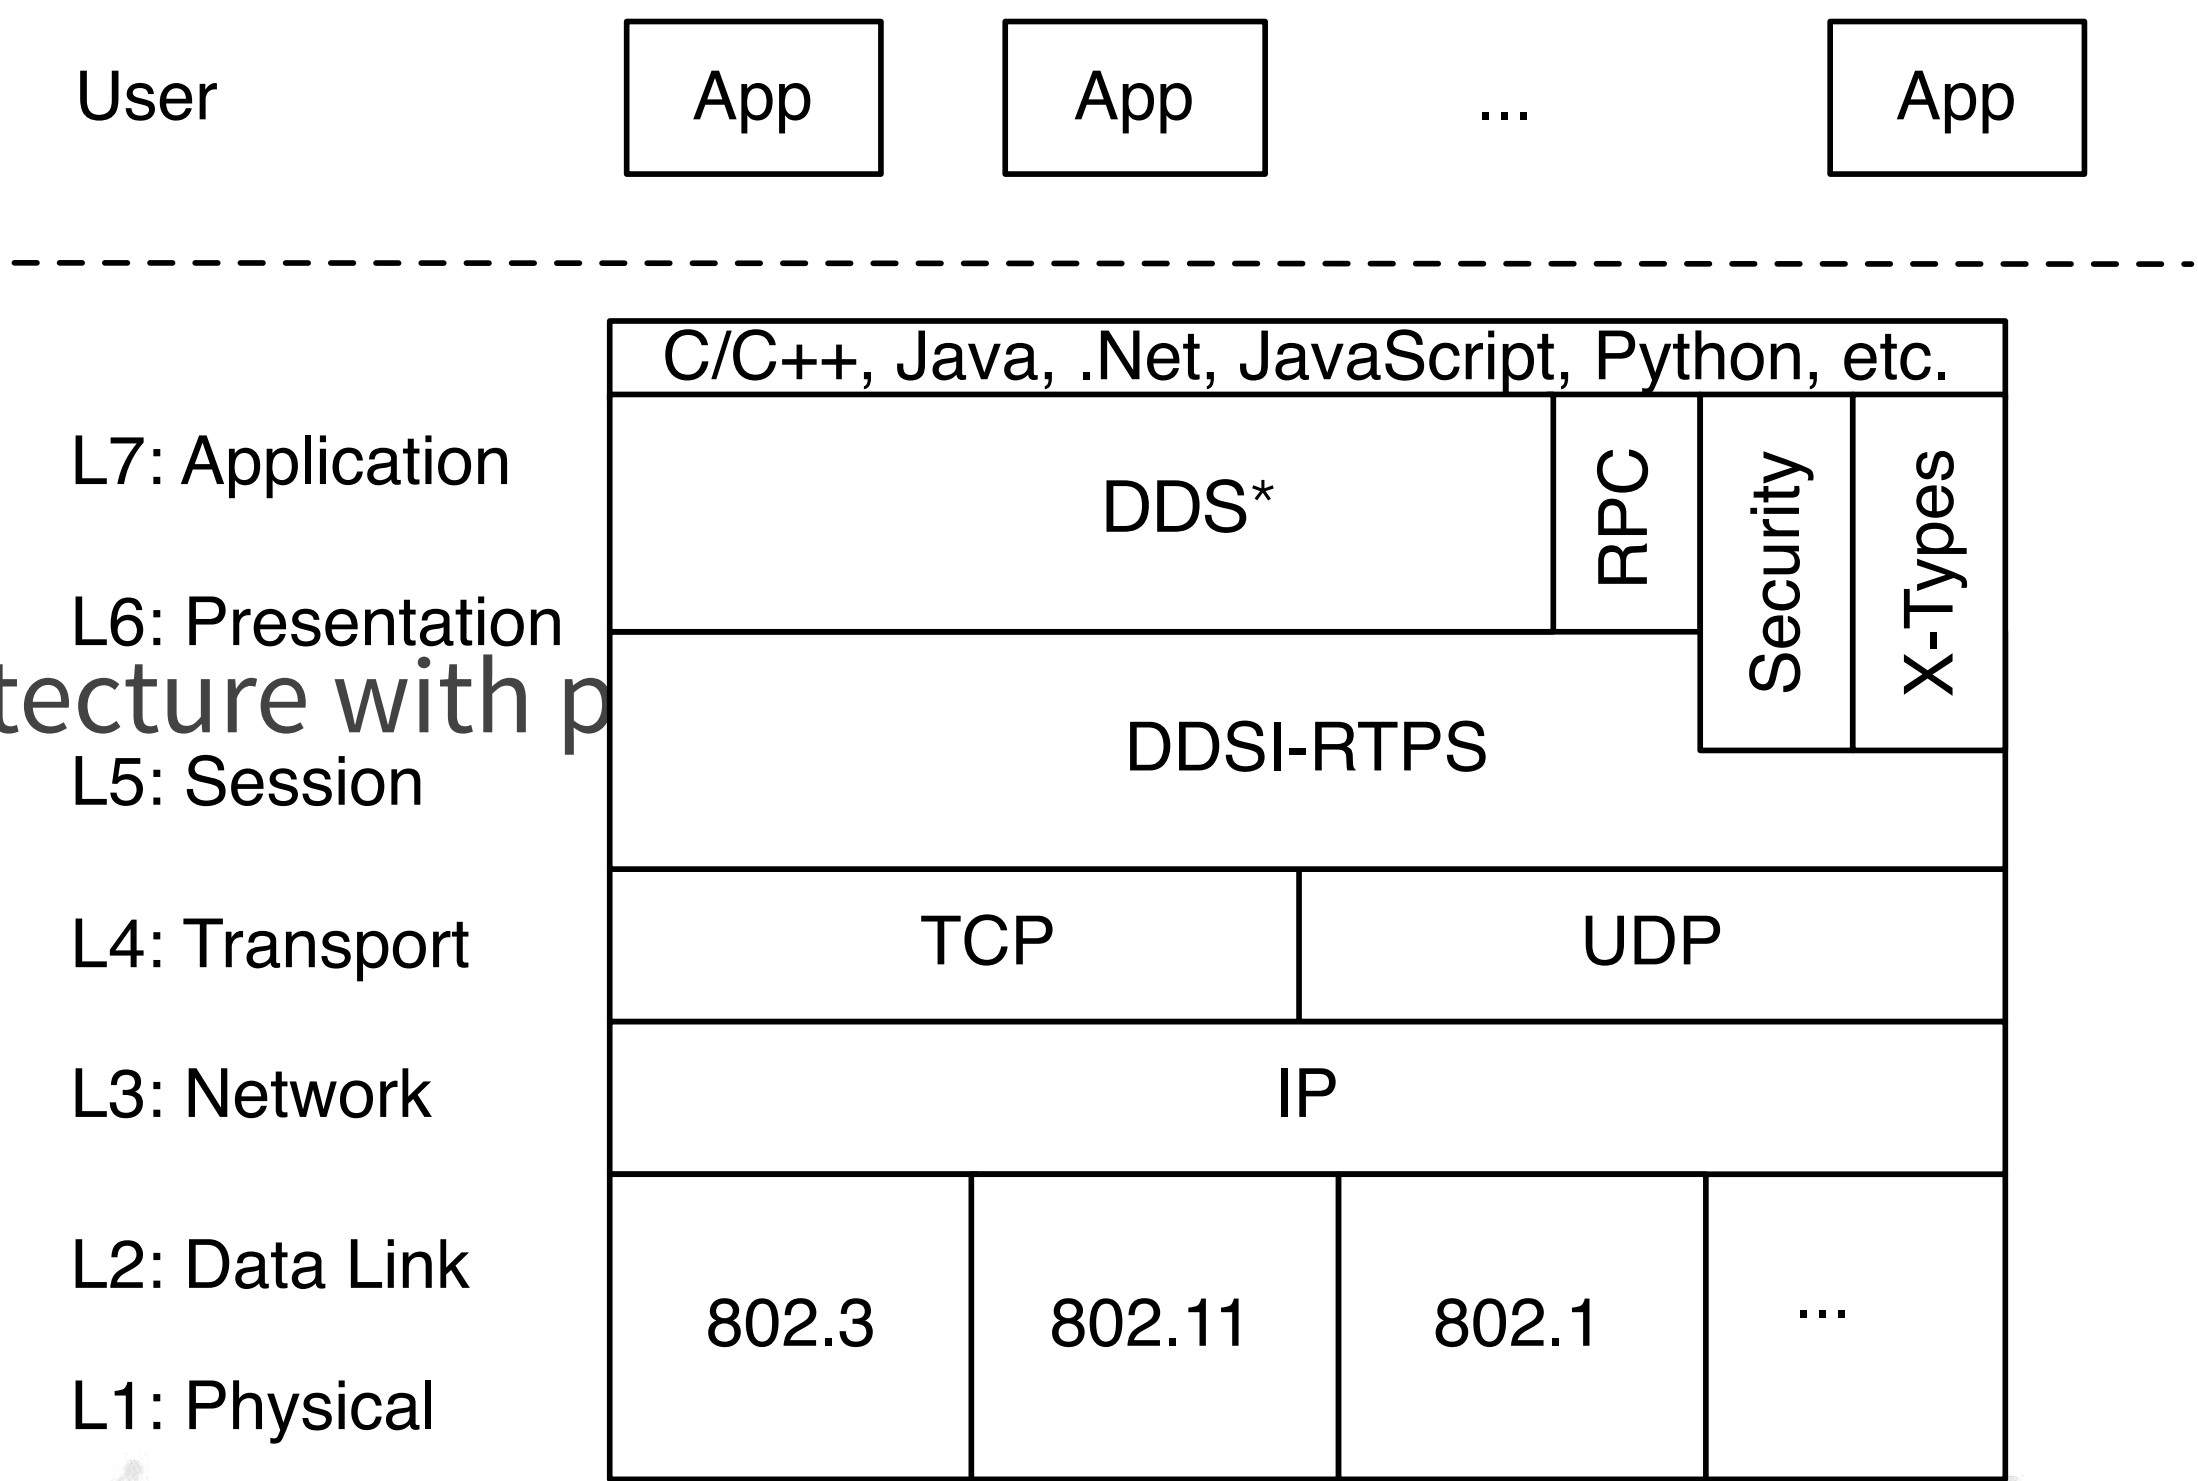


(*) This used to be called DCPS as originally the DDS standard was composed by two layers, DCPS (Data Centric Publish/Subscribe and DLRL (Data Local Reconstruction Layer).

DDS Standards

Security (DDS-Security)

Defines a data-centric security architecture with p
Access Control, Crypto and Logging.

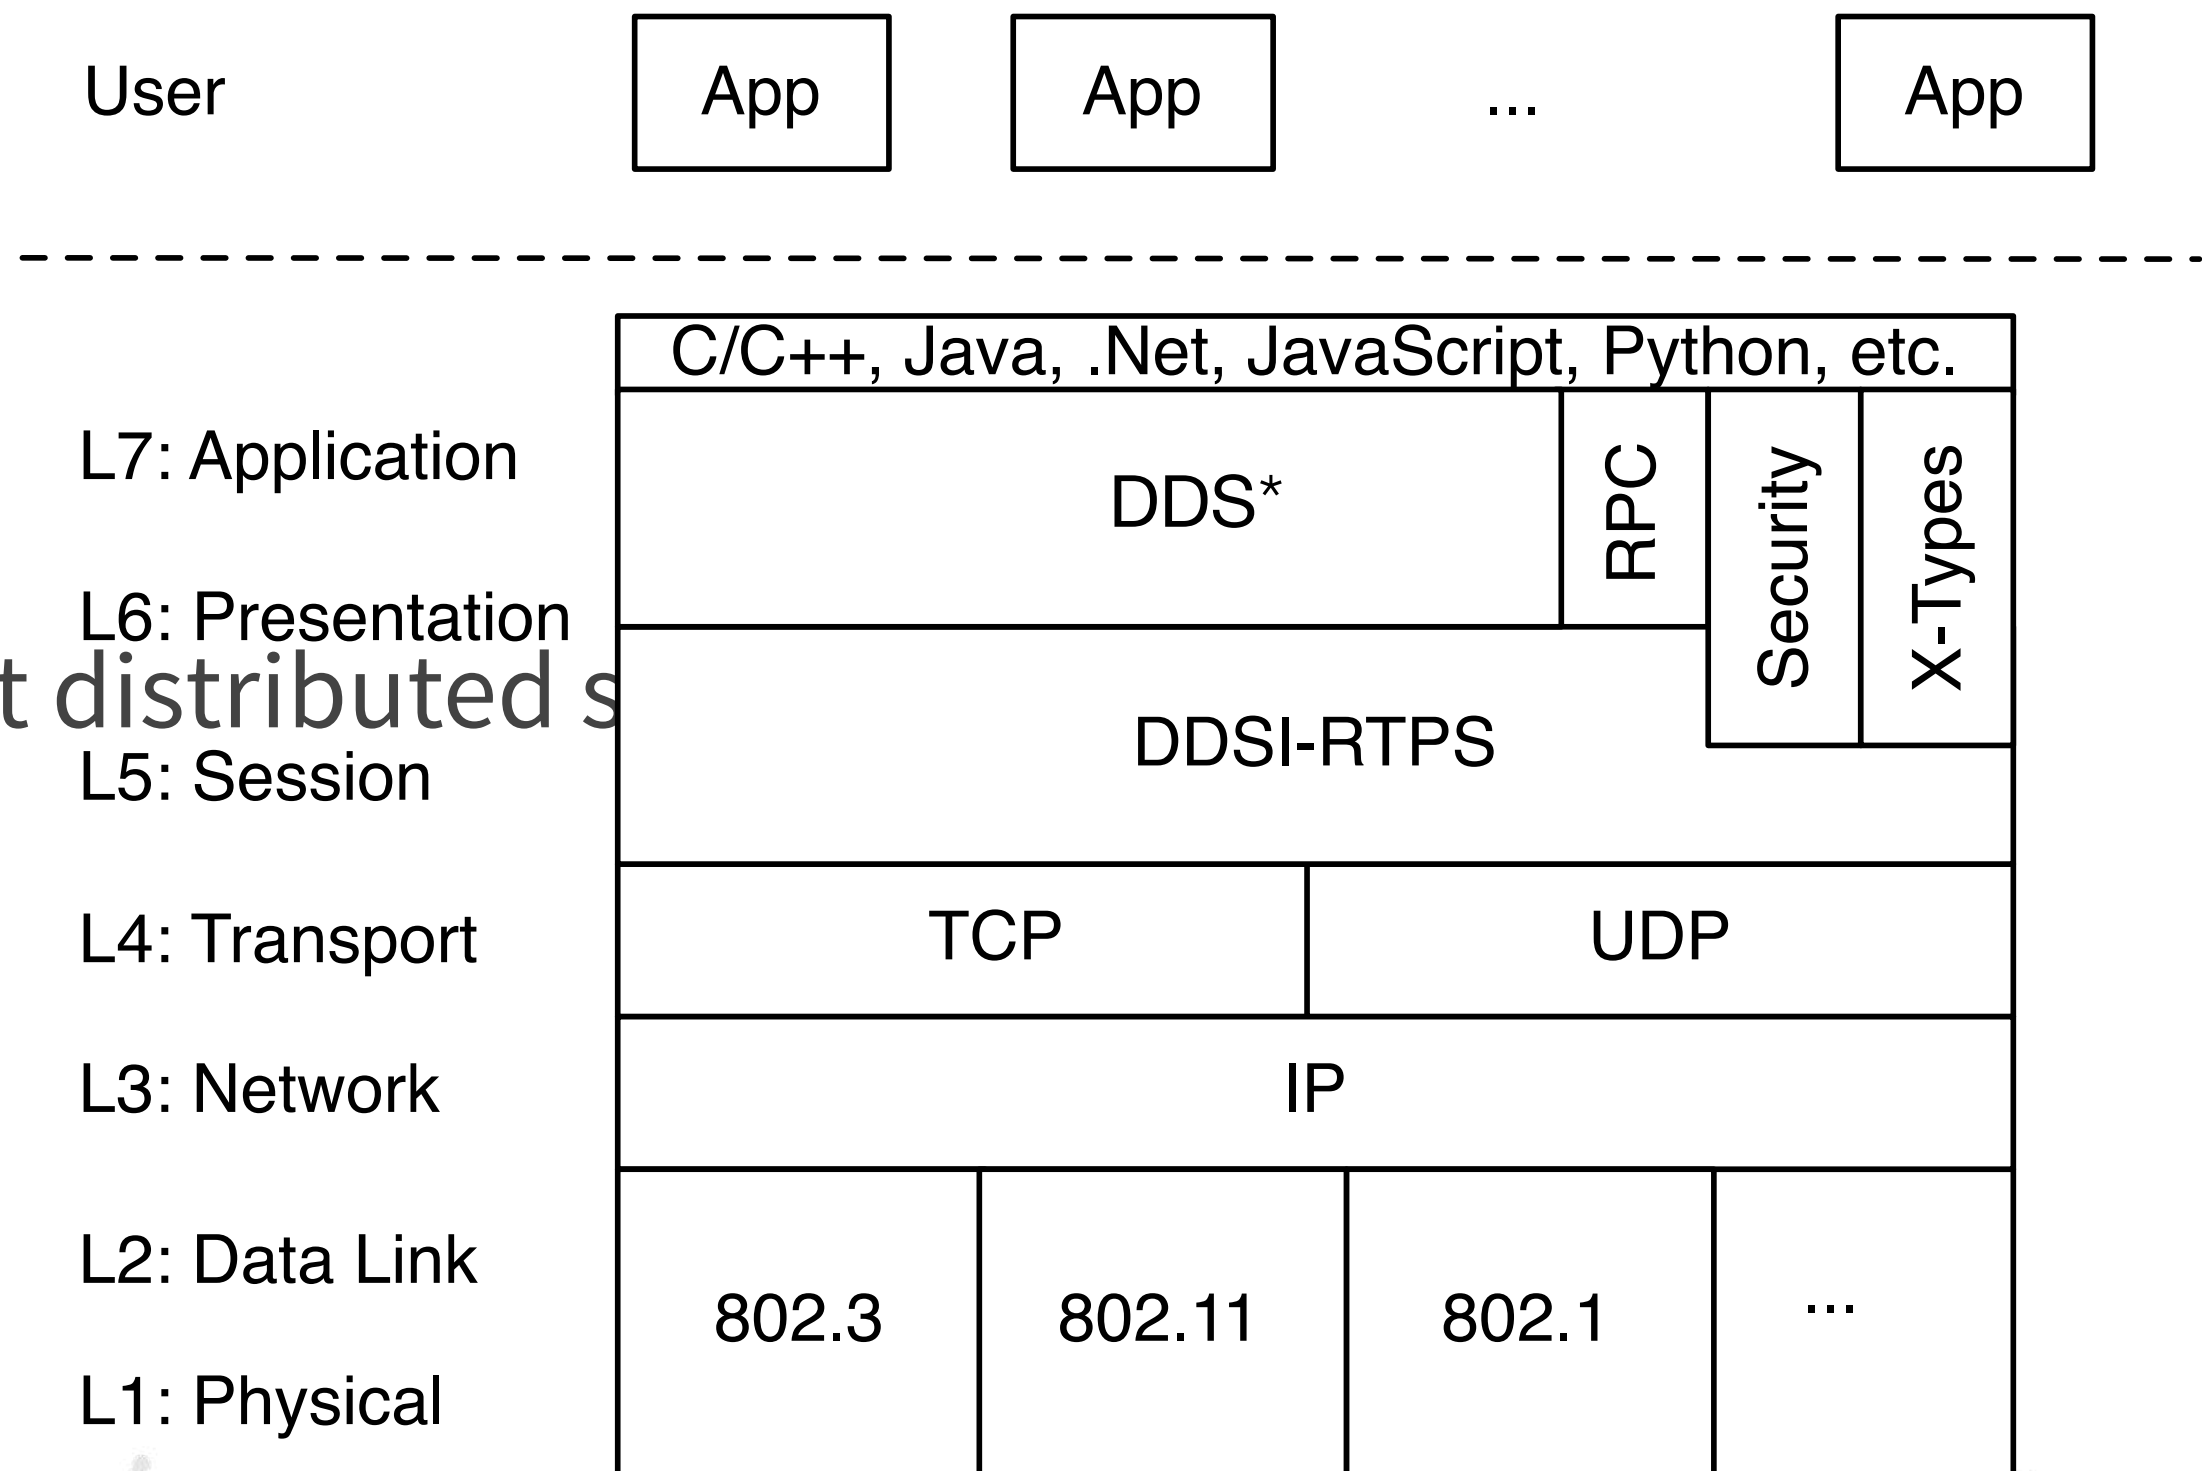


(*) This used to be called DCPS as originally the DDS standard was composed by two layers, DCPS (Data Centric Publish/Subscribe and DLRL (Data Local Reconstruction Layer).

DDS Standards

Remote Procedure Calls (DDS-RPC)

Extends DDS abstractions to support distributed systems
remote operation invocations.



(*) This used to be called DCPS as originally the DDS standard was composed by two layers, DCPS (Data Centric Publish/Subscribe and DLRL (Data Local Reconstruction Layer).

Who Uses DDS?

AUTONOMOUS VEHICLES

DDS is used for data sharing within and across the vehicle.

The environment is highly heterogeneous and requires dynamic pairing along with coordination of fast moving vehicles



SMART CITIES

DDS is used as the integration technology for data sources and sinks.

DDS is also often used as a control plane to control and provision equipment



SMART GRID

DDS is used to integrate and normalise data sharing among the various elements of a smart grid at scale

Duke's Energy COW showed how only with DDS it was possible to distribute the phase alignment signal at scale with the required 20ms periodicity



SMART GREEN HOUSES

DDS leveraged to
virtualise I/O and
provide better
decoupling between I/
O, Control and
Management functions
of the system



COMBAT MANAGEMENT SYSTEMS

DDS is used at the core of
an incredible number of
Naval Management System

In this context DDS
distributes soft and hard
real-time sensor and
actuator data

TACTICOS

Worlds' favourite Combat Management System
The best got better



NASA LAUNCH SYSTEMS

The NASA Kennedy Space
Centre uses DDS to collect
the Shuttle Launch System
Telemetry.

DDS streams over 400.000
Msgs/sec

In comparison, world-wide,
Twitter deals with less than
7000 msgs/sec



AIR TRAFFIC CONTROL

DDS is used to share
Flight Data within and
across Air Traffic Control
Centres.

These applications have
extremely high
dependability constraints.



NASA'S SMART NAS

The Shadow Mode Assessment Using Realistic Technologies for the National Airspace System (SMART-NAS) Project is an air traffic management simulation capability to explore the NAS integration of alternative concepts, technologies and architectures.



UNMANNED AIR VEHICLES

DDS is used in UAV
in-flight mission
management system
to distribute several
thousands of sensor
targets



SIMULATORS

Airbus Helicopters use DDS for their full-function (level D) helicopter simulators

All data in the system is managed in DDS, with applications converting to the ARINC429 bus used by the avionics



SMART FACTORY


DDS is used in Smart
Factories to provide
horizontal and
vertical data
integration across the
traditional SCADA
layers.



EXTREMELY LARGE TELESCOPE (ELT)

DDS used to control
the 100.000 mirrors
that make up ELT's
optics.

These mirrors are
adjusted 100 times per
second



MEDICAL DEVICES

DDS is used inside
several medical
devices to share data
between the various
stages of data
acquisition, processing
and visualisation



ROBOTICS

DDS is heavily used for data sharing in Robotics and is today at the heart of the Robot Operating System (ROS)



Why are these Applications Using DDS?

**DDS PROVIDES AN EXTREMELY HIGH
LEVEL AND POWERFUL ABSTRACTIONS
ALONG WITH A ROCK SOLID
INFRASTRUCTURE TO BUILD HIGHLY
MODULAR AND DISTRIBUTED SYSTEMS**

**DDS MAKES IT MUCH EASIER TO SOLVE SOME VERY
HARD DISTRIBUTED SYSTEM PROBLEMS, SUCH AS
FAULT-TOLERANCE, SCALABILITY AND
ASYMMETRY**

DDS IS LIKE A **UNIVERSAL GLUE** THAT ALLOWS TO
SEAL TOGETHER HIGHLY HETEROGENEOUS
ENVIRONMENTS WHILE MAINTAINING A SINGLE,
ELEGANT AND EFFICIENT ABSTRACTION

Platform Independence



OS X

HTML



OS




DDS implementation are available on a wide range of platforms, including enterprise desktop, embedded, real-time, mobile, and web.



VxWorks®



WIND RIVER



Polyglot



DDS applications can be written in your favourite programming language.



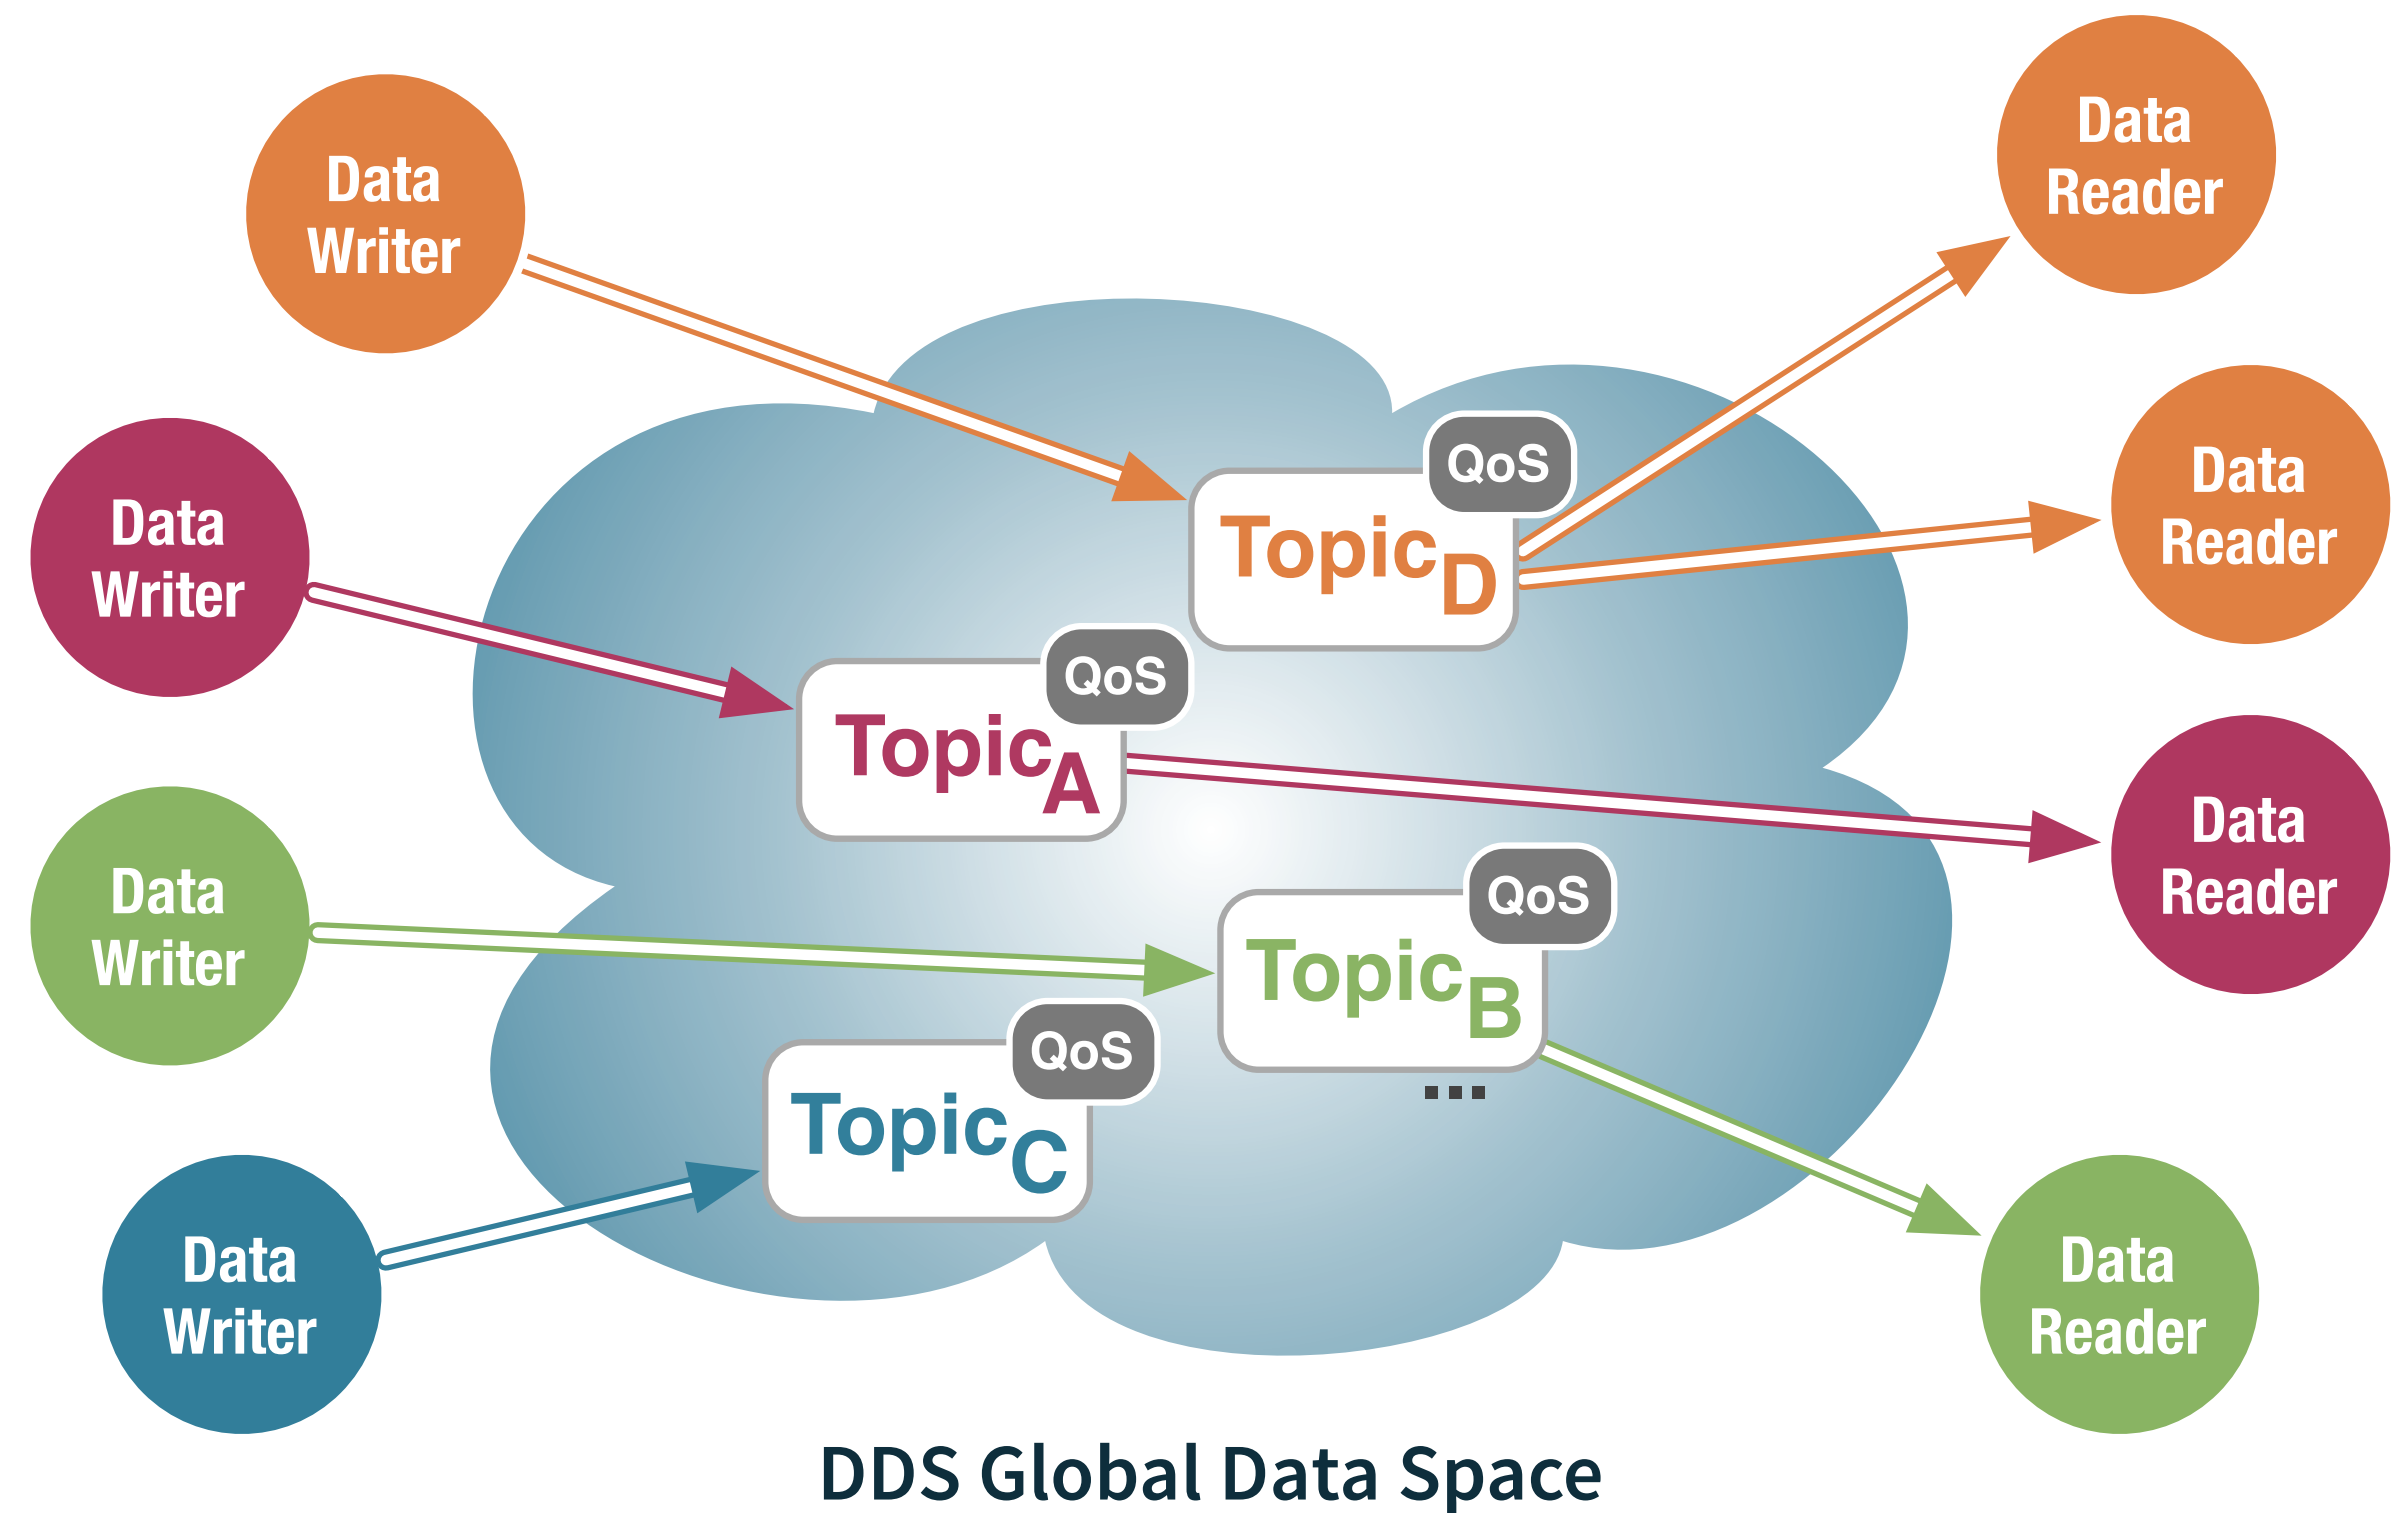
Interoperability across languages is taken care



DDS: Foundations

Virtualised Data Space

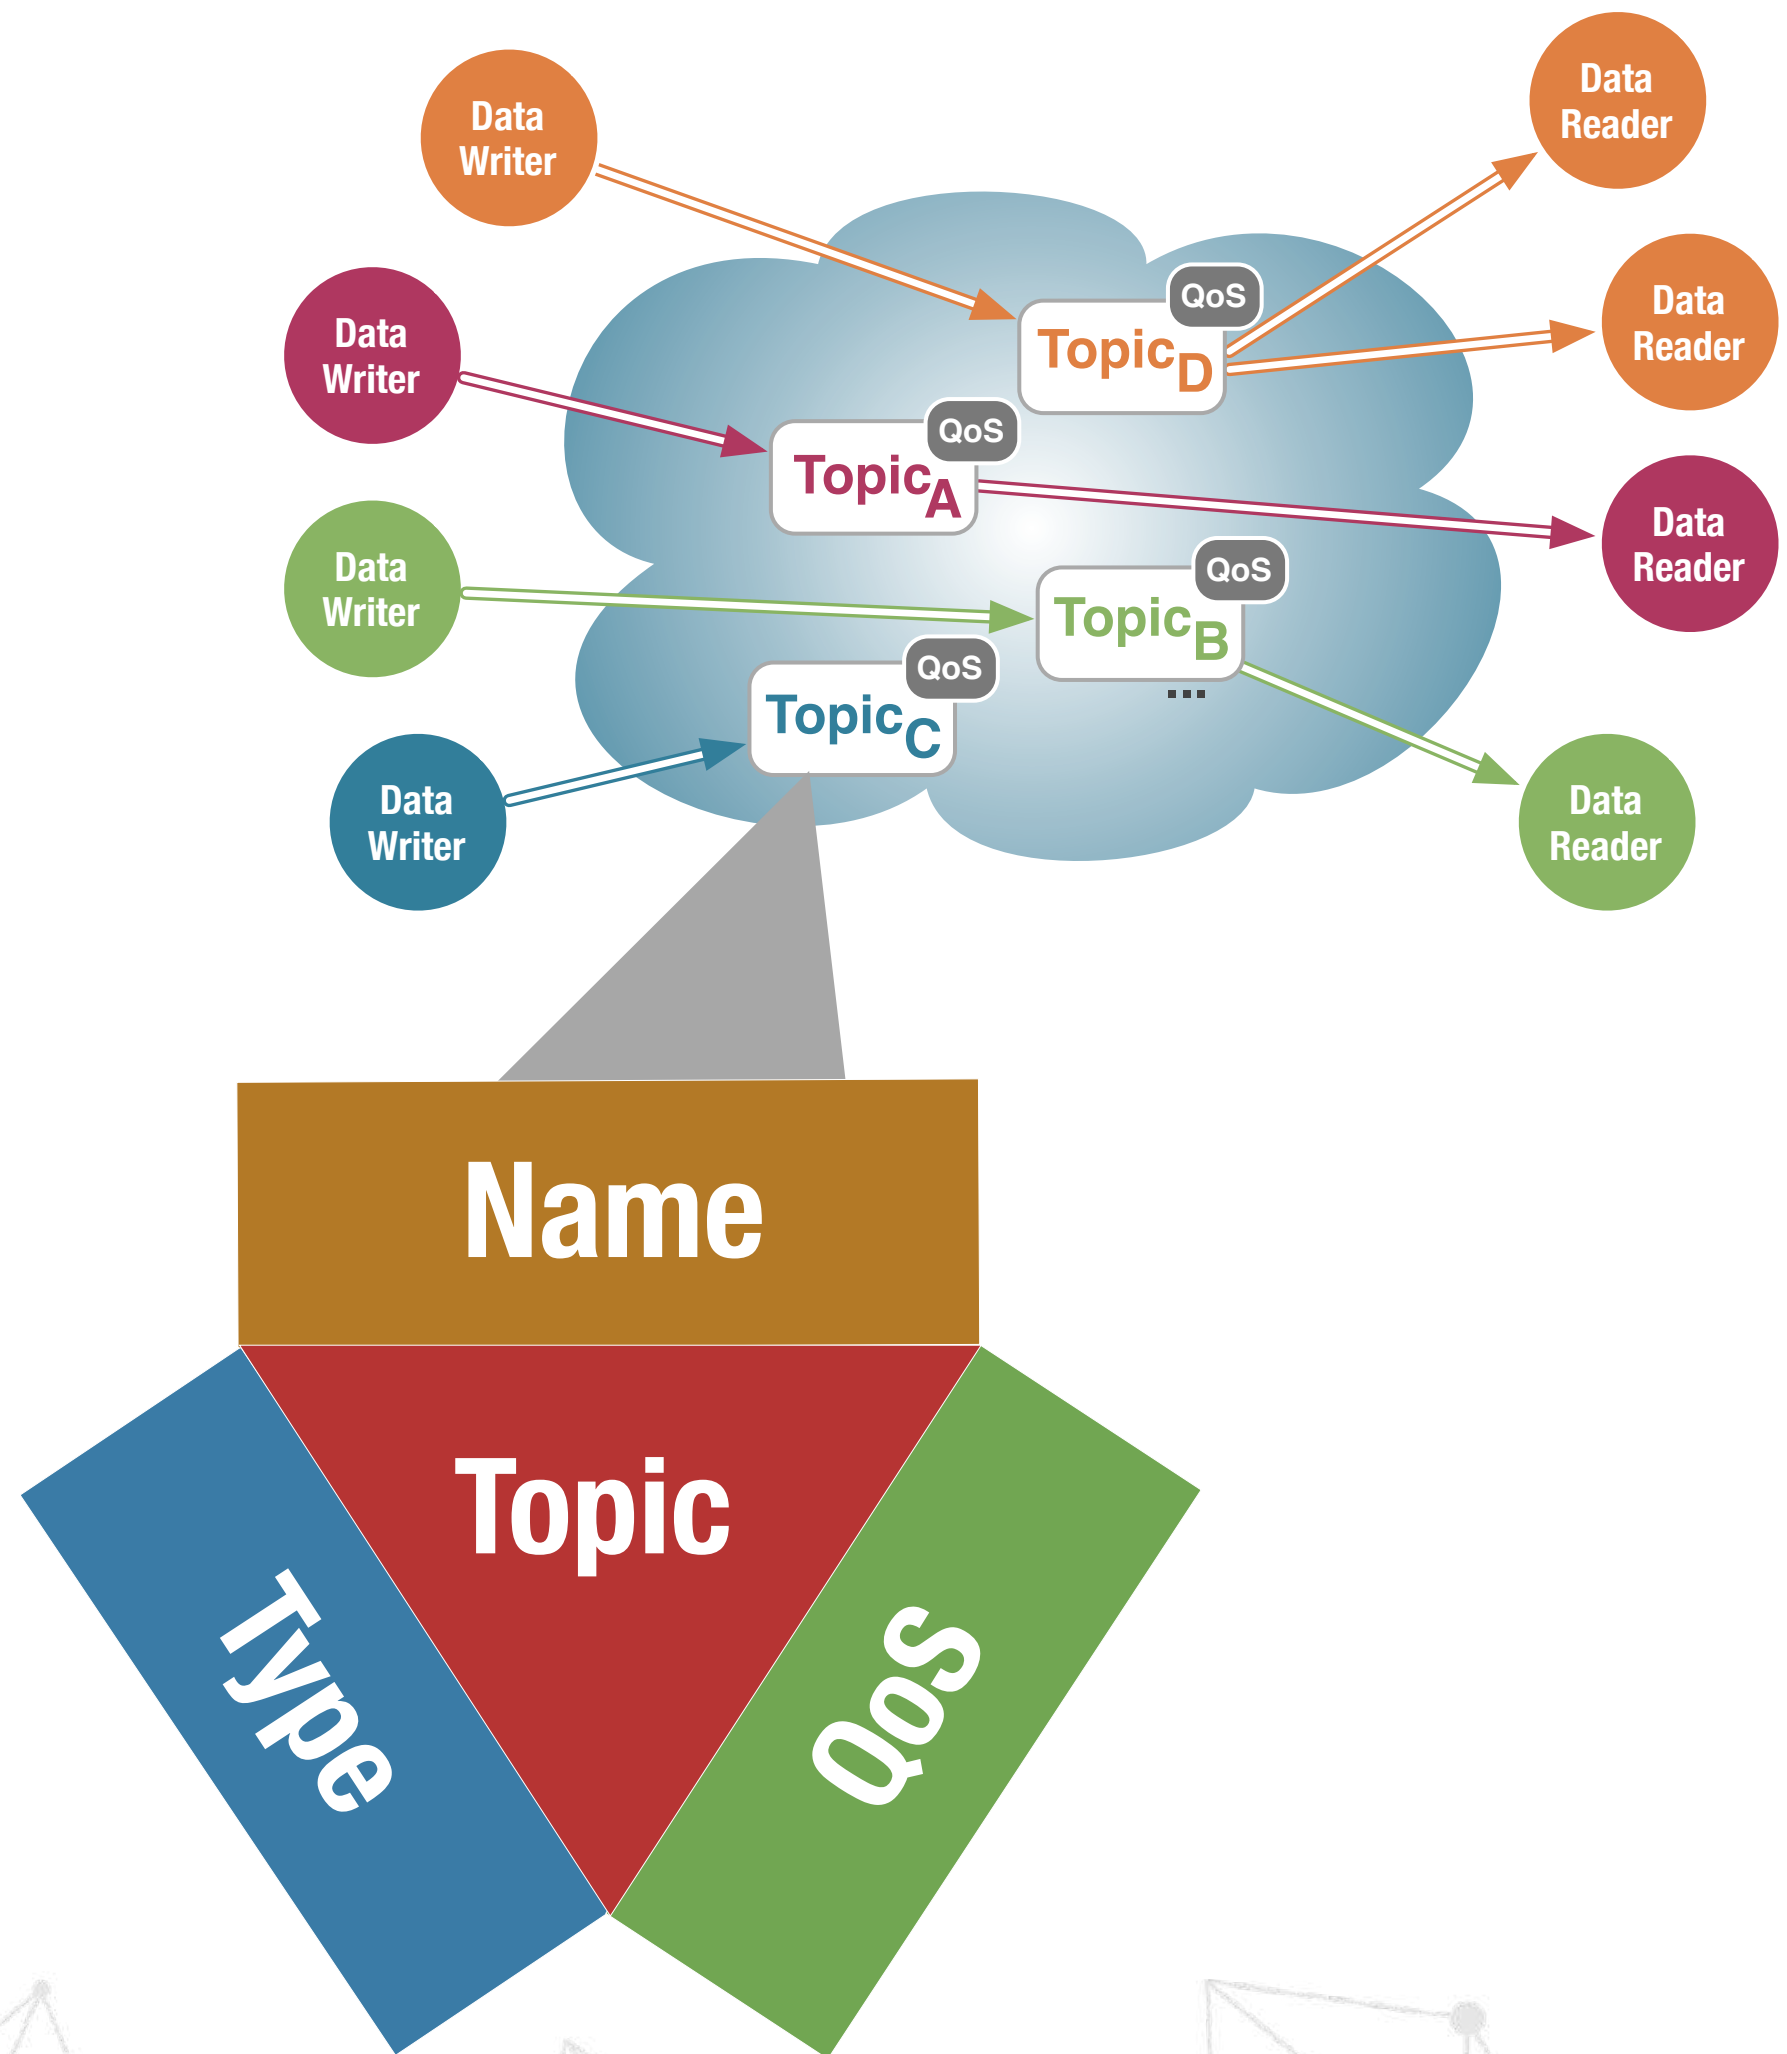
Applications can **autonomously** and **asynchronously** **read** and **write** data enjoying **spatial** and **temporal decoupling**



Topic

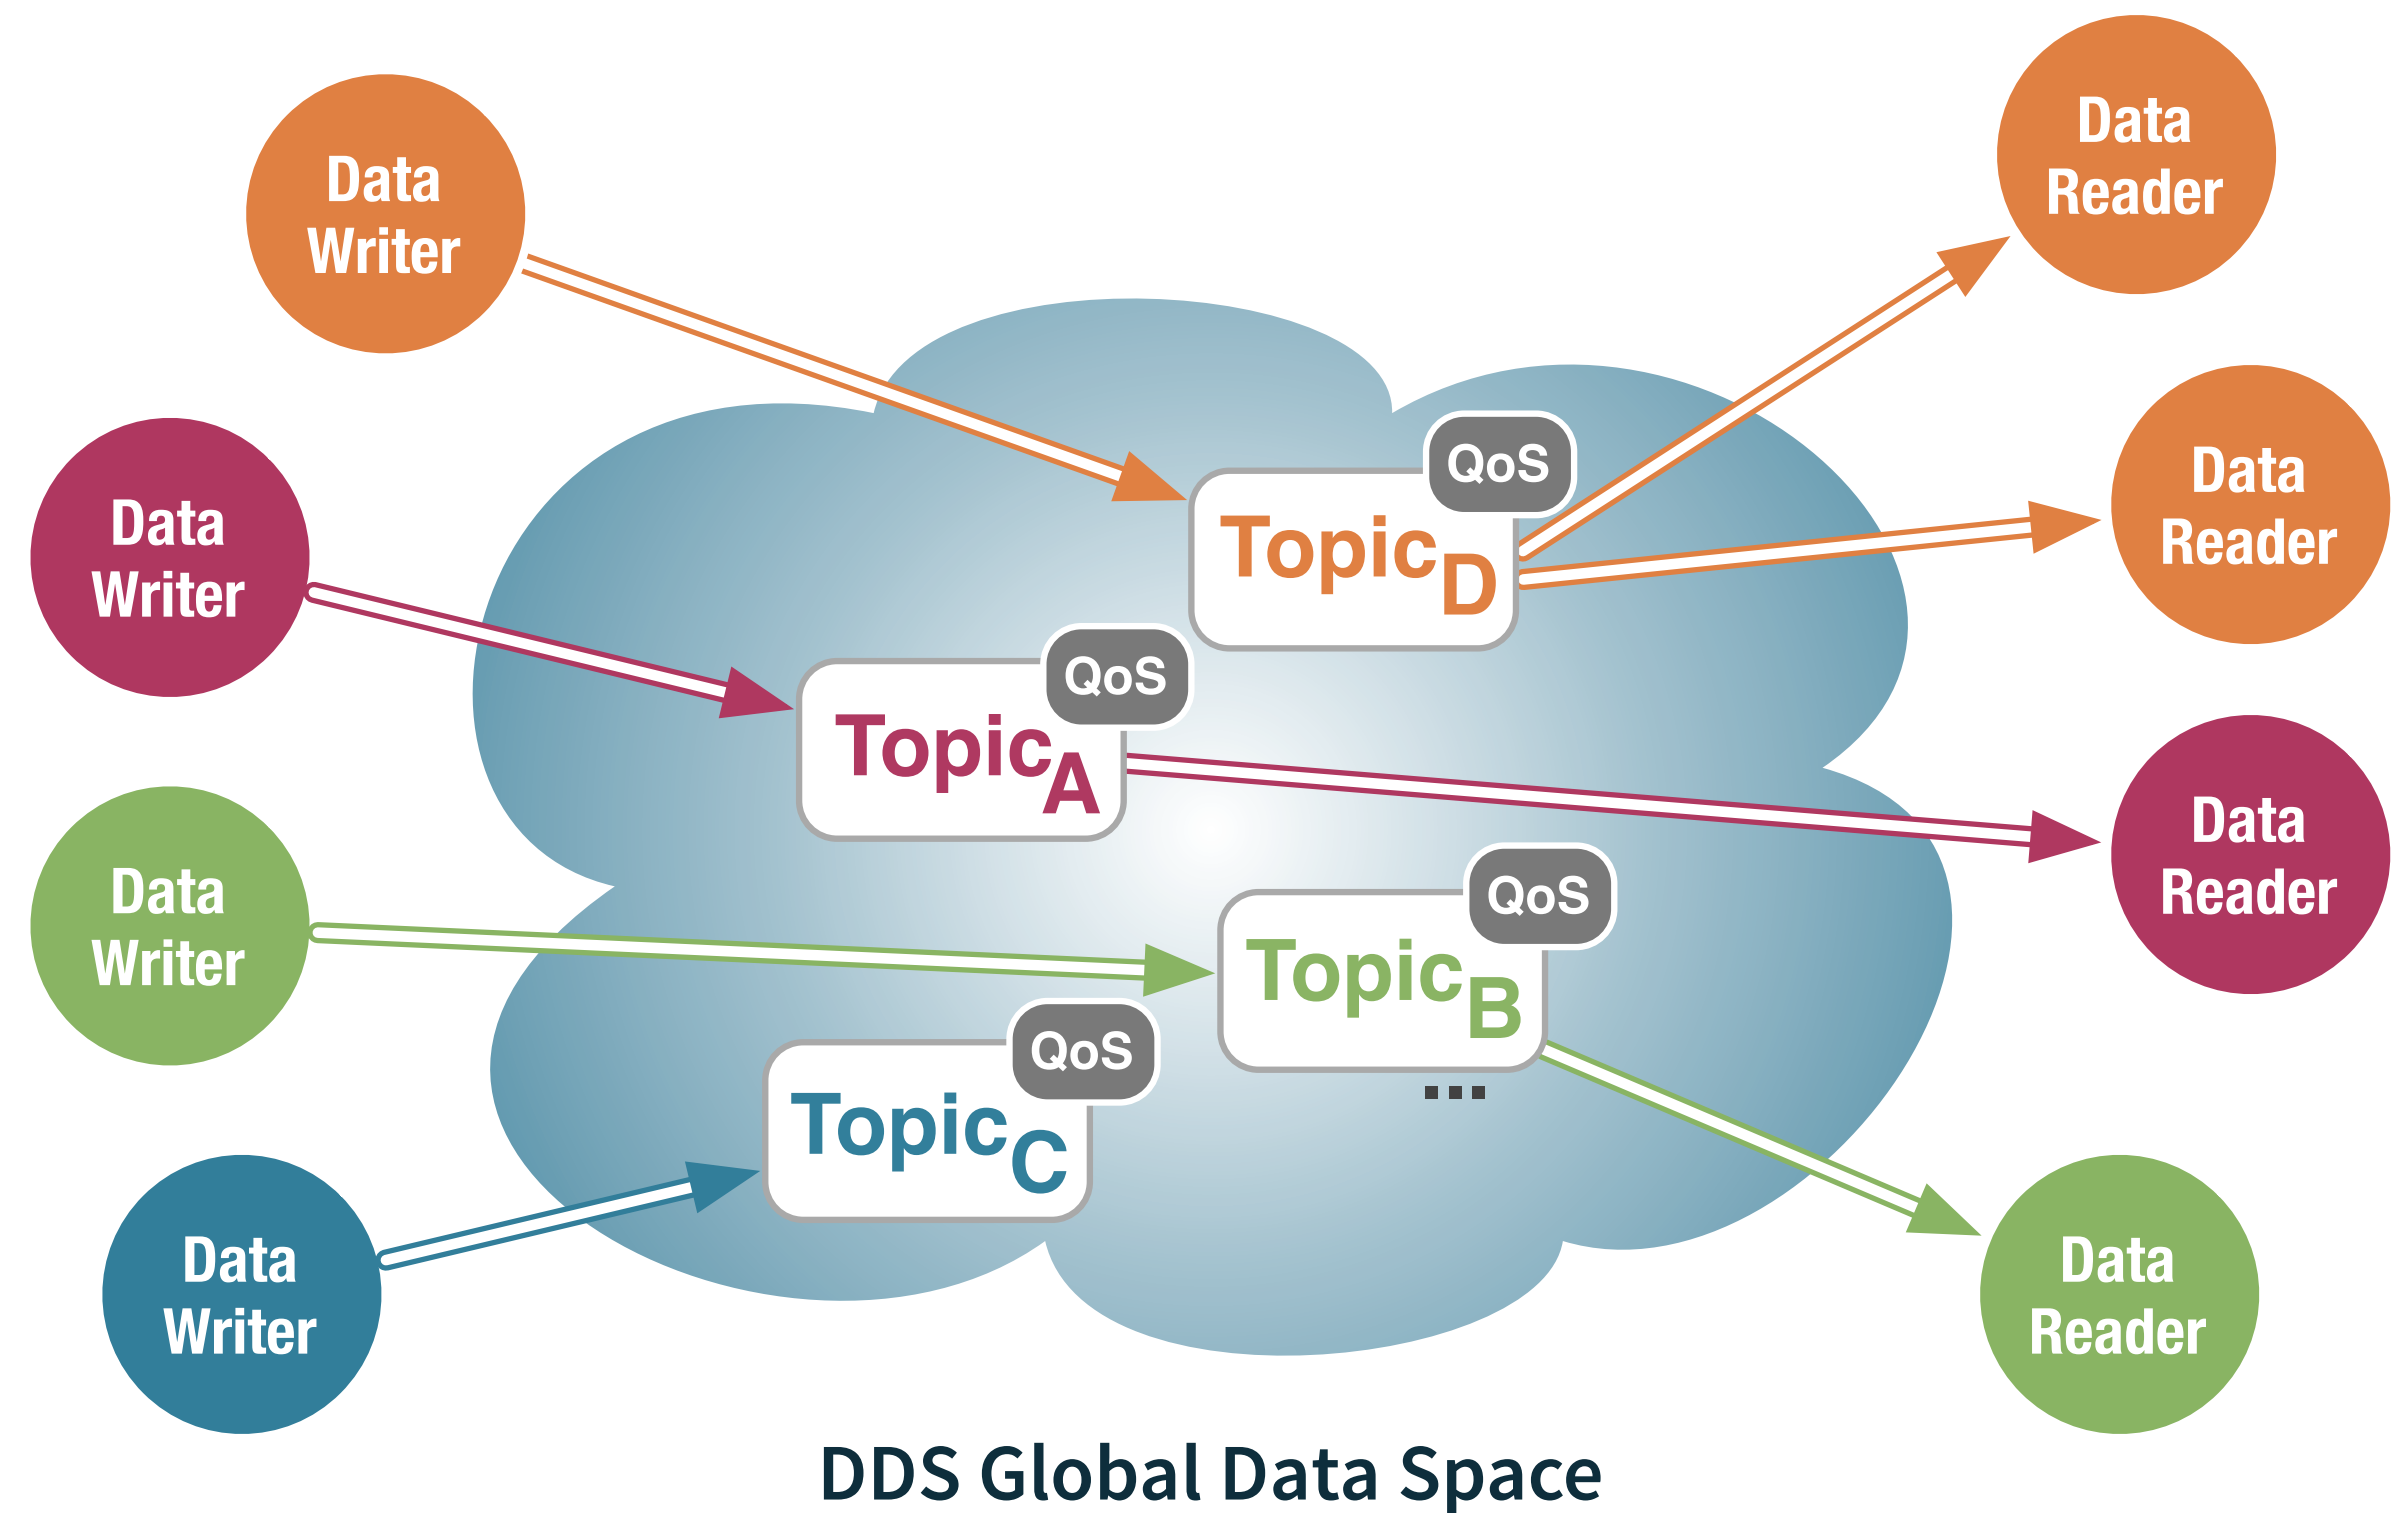
A domain-wide
information's class

A **Topic** defined by means
of a <name, type, qos>



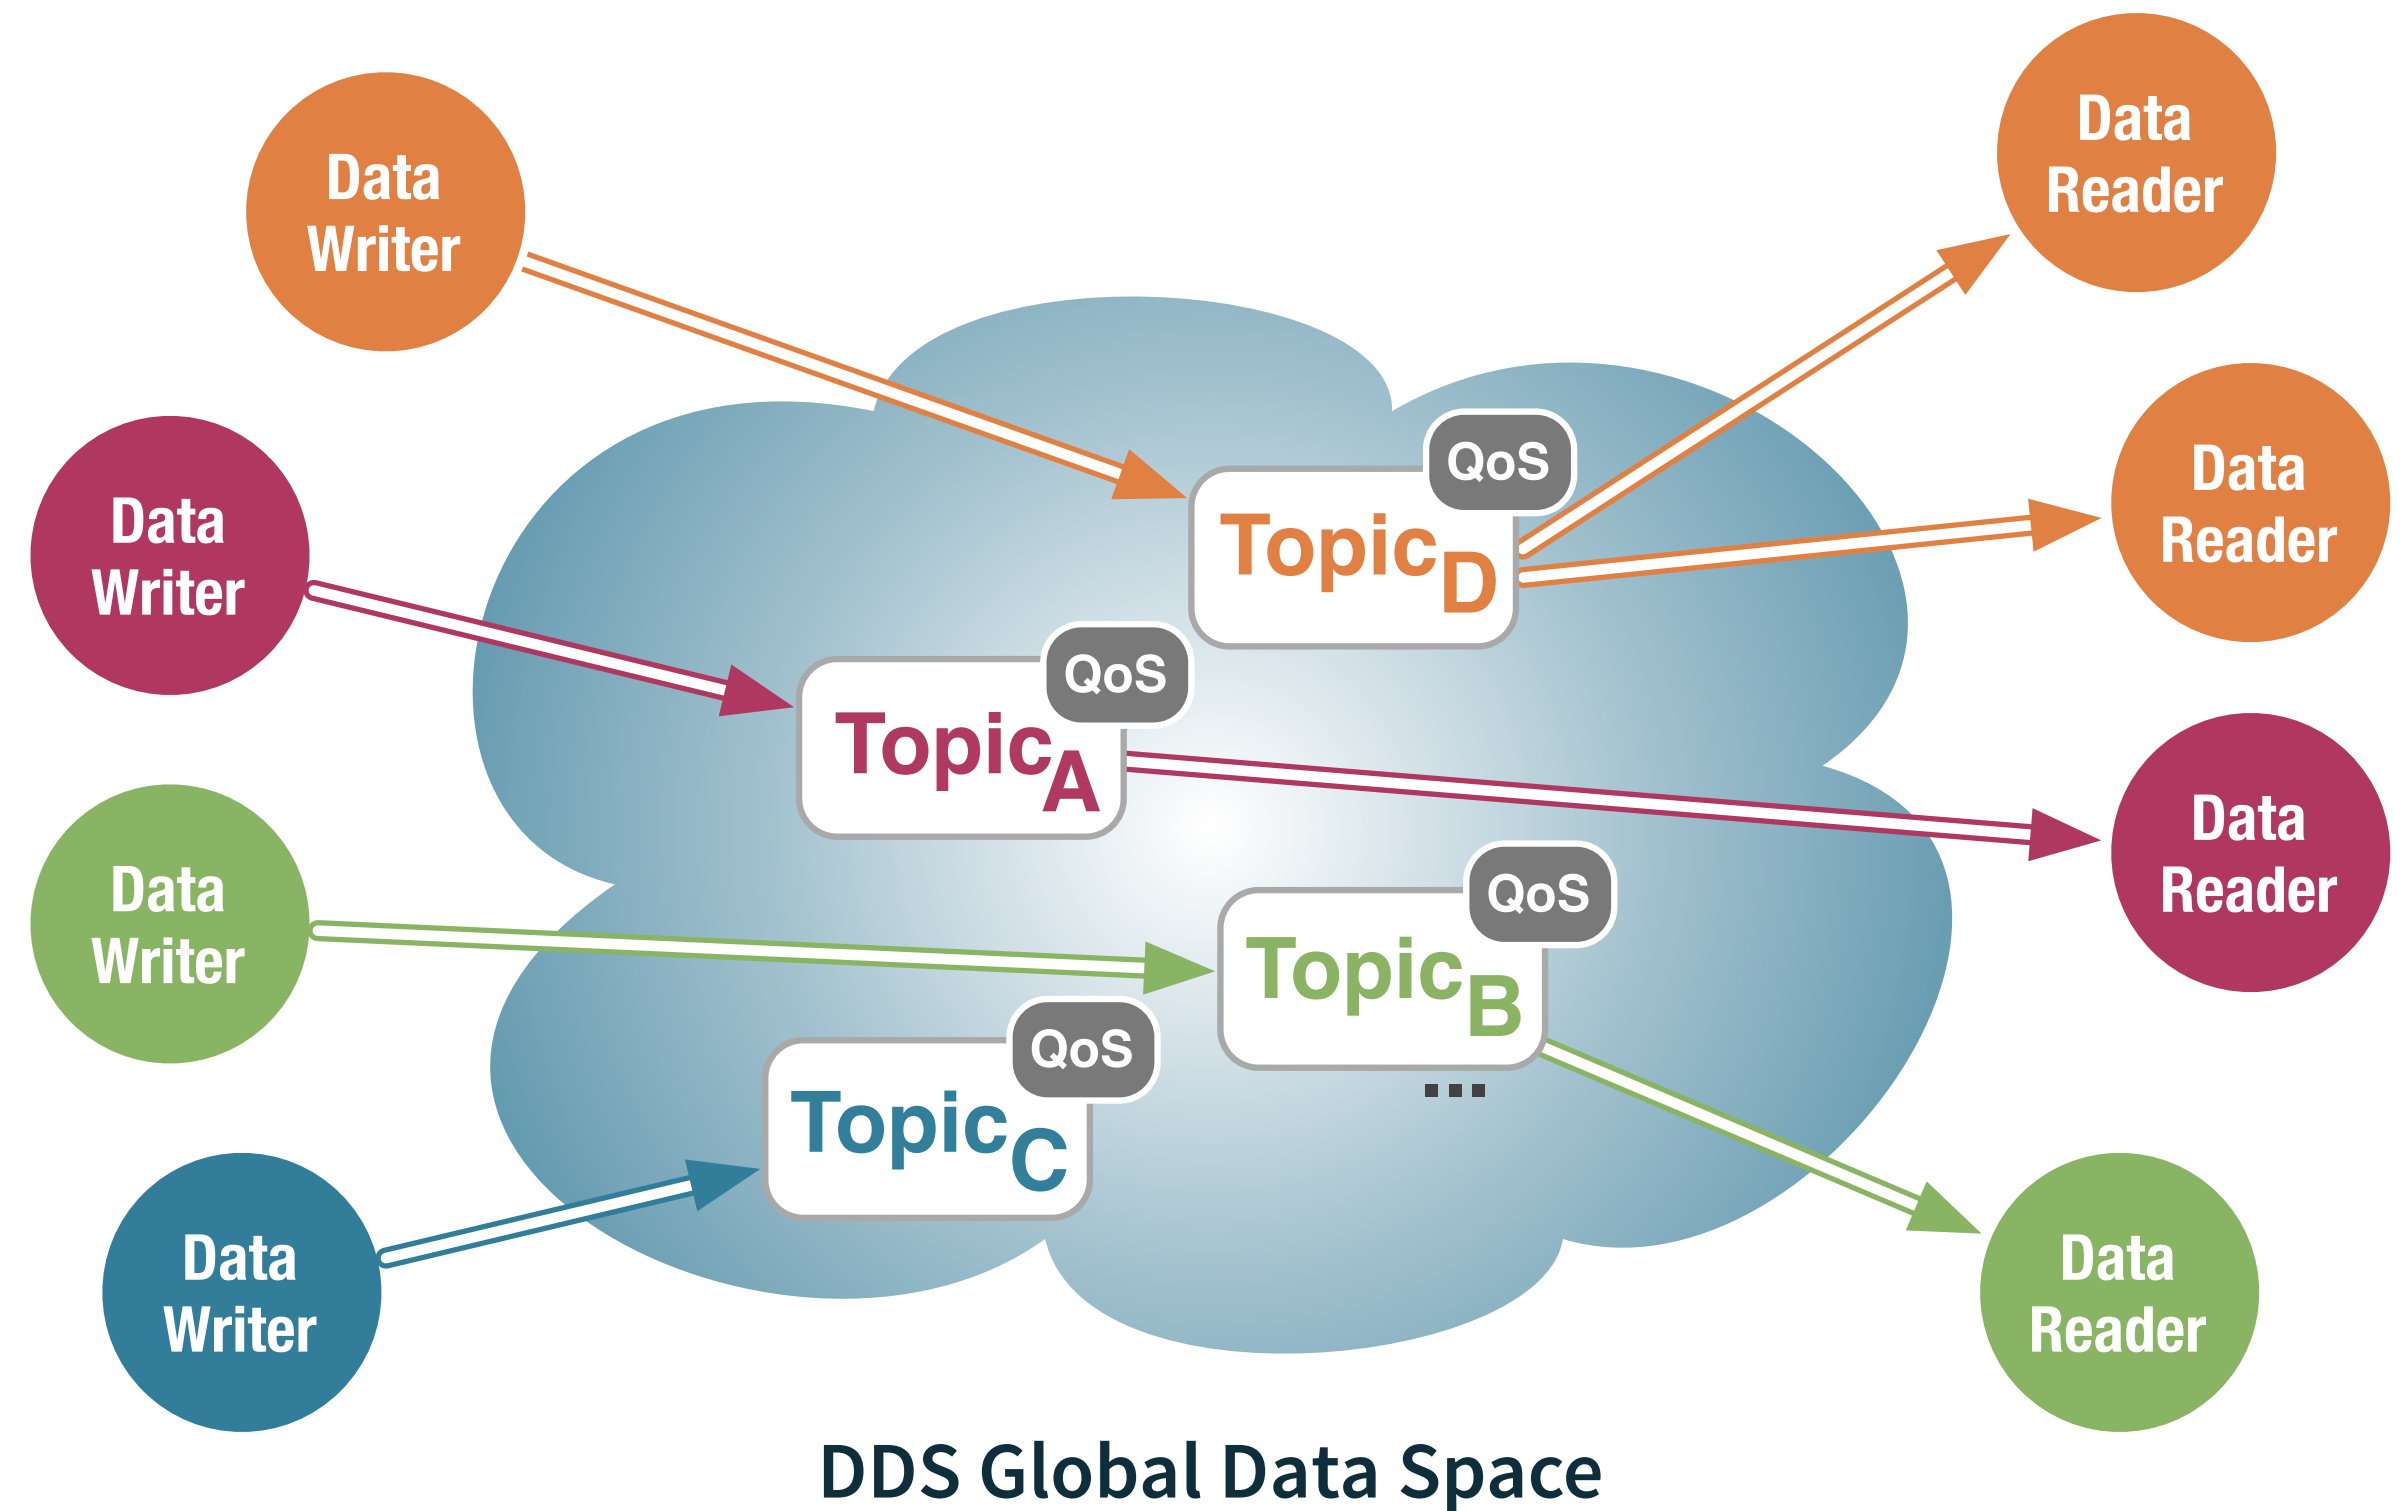
QoS Enabled

QoS policies allow to express **temporal** and **availability constraints** for data



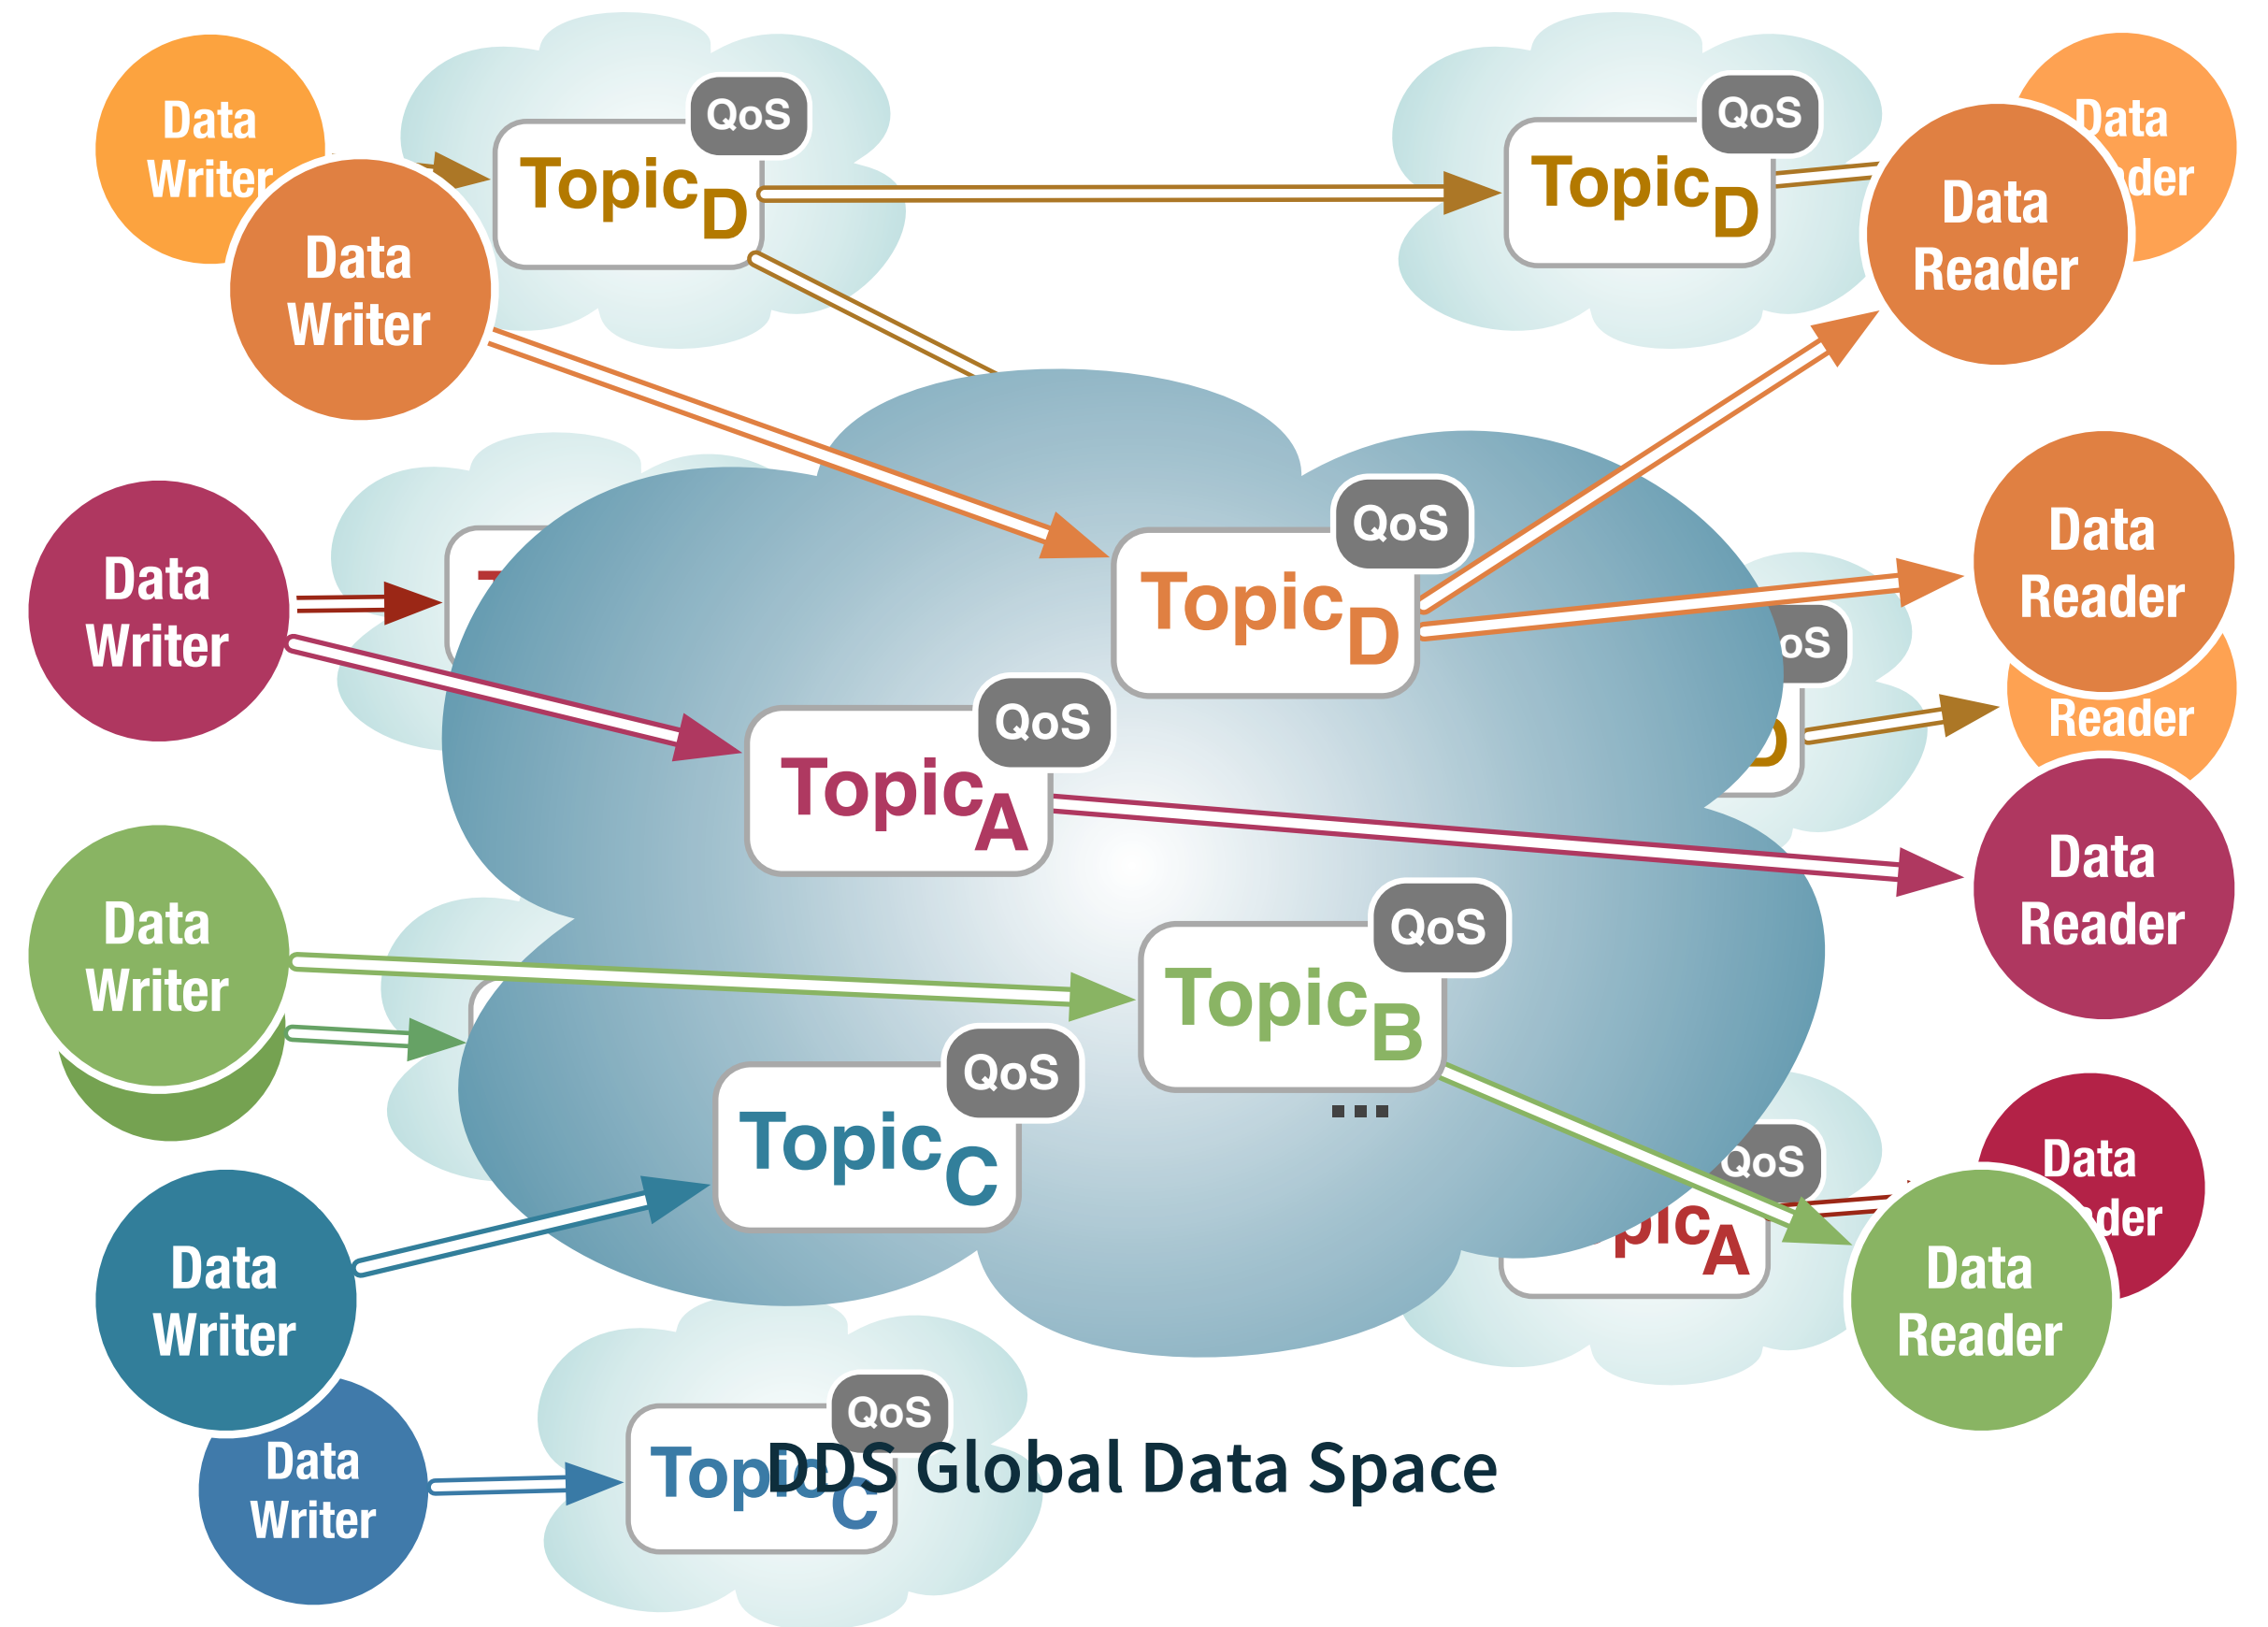
Dynamic Discovery

Built-in dynamic discovery **isolates** applications from **network topology** and **connectivity** details



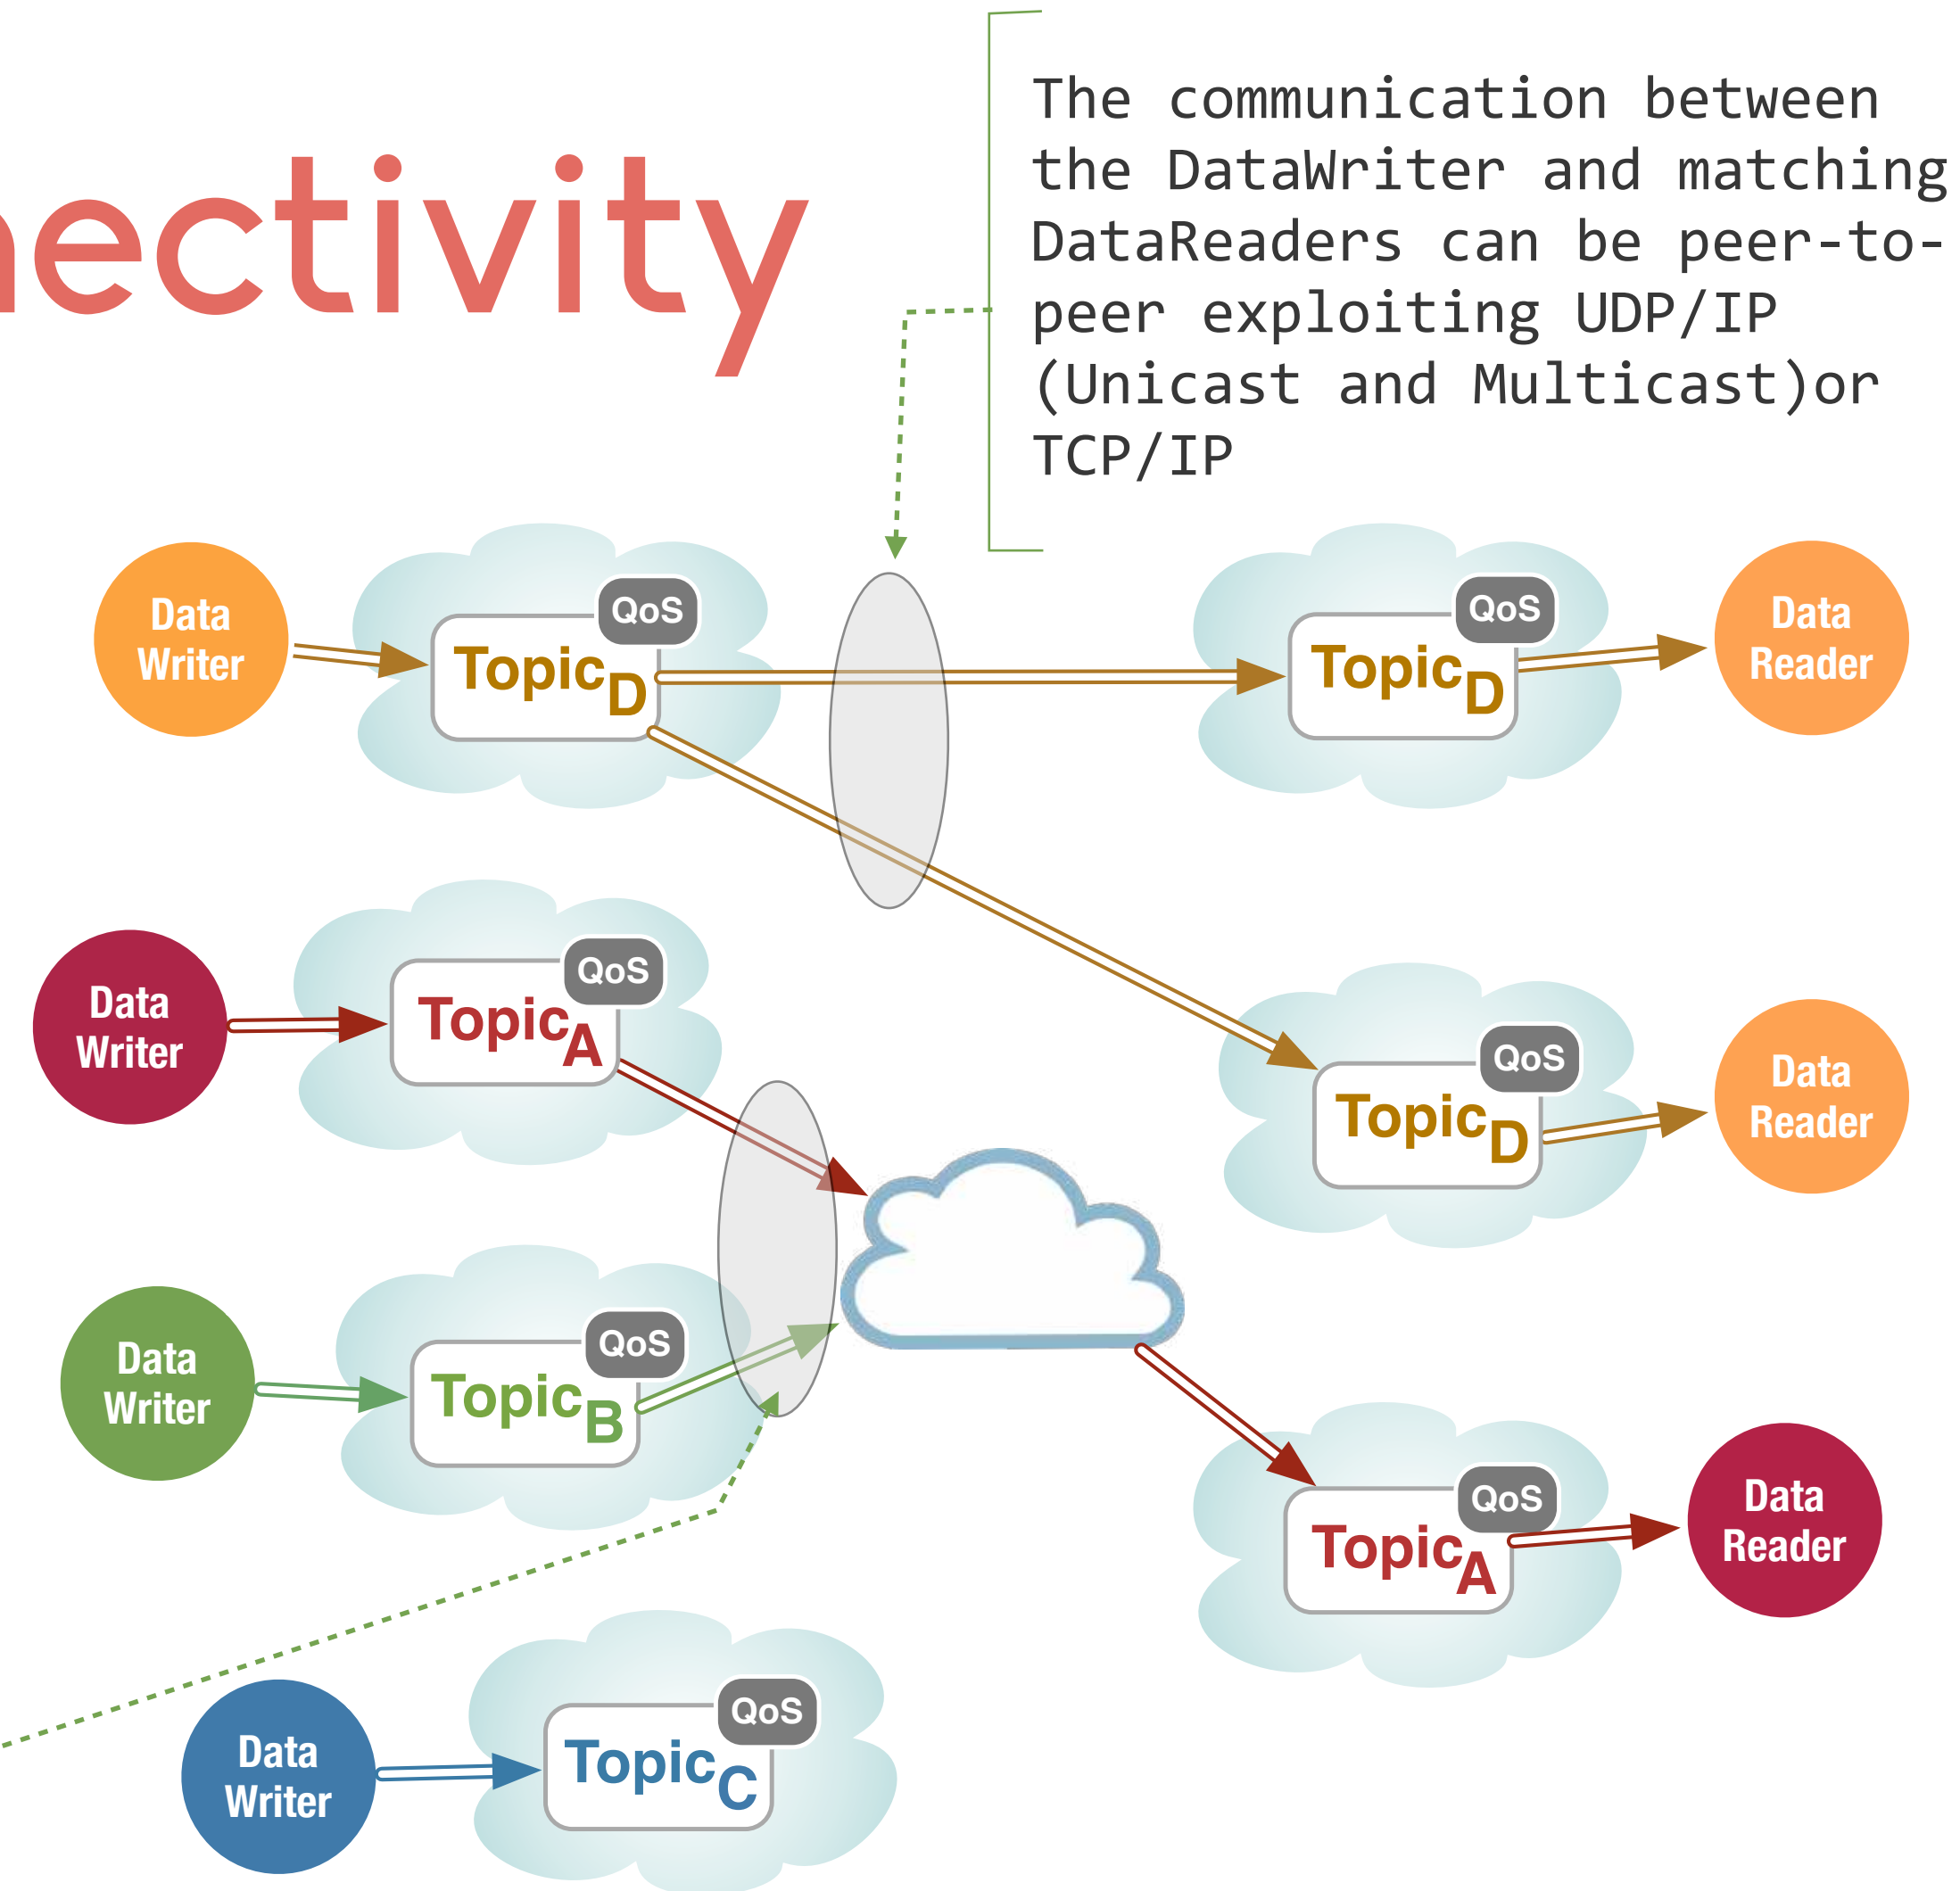
Decentralised Data-Space

No single point of failure or bottleneck



Adaptive Connectivity

Connectivity is **dynamically adapted** to chose the most effective way of sharing data




```
19 using namespace dds;
20 using namespace dds::core;
21 using namespace dds::core::policy;
22 using namespace dds::topic;
23 using namespace dds::pub;
24 using namespace dds::sub;
25
26 using namespace org::eclipse;
27
28 int
29 main(int argc, char* argv[])
30 {
31
32     try {
33         // Parse the command line args
34         tspub_options opt = parse_tspub_args(argc, argv);
35
36         // Create the domain participant
37         dds::domain::DomainParticipant dp(cyclonedds::domain::default_id());
38
39         // Initialize random number generation with a seed
40         srand(clock());
41
42         dds::topic::qos::TopicQos tqos =
43             dp.default_topic_qos() << LatencyBudget(Duration(2,0)) << Deadline(Duration(4,0));
44
45         auto topic = Topic<tutorial::TempSensorType>(dp, "TempSensorTopic", tqos);
46         auto pub = Publisher(dp);
47         auto dw = DataWriter<tutorial::TempSensorType>(pub, topic, tqos);
48
49     }
```

```
[variant] Loaded new set of variants
[kit] Successfully loaded 1 kits from /Users/kydos/.local/share/CMakeTools/cmake-tools-kits.json
[main] Configuring folder: dds-tutorial
```

Lab 1

First App

Cyclone DDS



Eclipse Cyclone DDS was born with the ambition of developing the **best DDS implementation ever**.

We are leveraging the **Open Source Ecosystem** to facilitate **user-driven innovation** and mitigate some of the downsides of Open Standards with Open Source and Open Innovation

We are **working with the community** to **establish** the **Go-To DDS implementation** for all captive markets



<https://github.com/eclipse-cyclonedds/cyclonedds>

Writing Data in C++

```
#include <dds/dds.hpp>

int main(int, char**) {

    dds::domain::DomainParticipant dp(0);
    dds::topic::Topic<Meter> topic("SmartMeter");
    dds::pub::Publisher pub(dp);
    dds::pub::DataWriter<Meter> dw(pub, topic);

    while (!done) {
        auto value = readMeter()
        dw.write(value);
        std::this_thread::sleep_for(SAMPLING_PERIOD);
    }

    return 0;
}
```

```
enum UtilityKind {
    ELECTRICITY,
    GAS,
    WATER
};

struct Meter {
    string sn;
    UtilityKind utility;
    float reading;
    float error;
};

#pragma keylist Meter sn
```


Reading Data in C++

```
#include <dds/dds.hpp>

int main(int, char**) {

    dds::domain::DomainParticipant dp(0);
    dds::topic::Topic<Meter> topic("SmartMeter");
    dds::sub::Subscriber sub(dp);
    dds::sub::DataReader<Meter> dr(dp, topic);

    auto samples = dr.read();
    std::for_each(samples.begin(),
                  samples.end(),
                  [](Sample<Meter>& sample) {
                      std::cout << sample.data() << std::endl;
                  });

    return 0;
}
```

```
enum UtilityKind {
    ELECTRICITY,
    GAS,
    WATER
};

struct Meter {
    string sn;
    UtilityKind utility;
    float reading;
    float error;
};
#pragma keylist Meter sn
```


Console Time!

```
build — kydos@matcha — ../cpp/01/build — -zsh — 80x24
[→ 01 git:(master) * ls
CMakeLists.txt      TempControl.idl      tssub.cpp
README              build                 util.cpp
TempControl-orig.idl tsub.cpp              util.hpp
[→ 01 git:(master) * cd build
[→ build git:(master) * ls
CMakeCache.txt      TempControl.cpp      cmake_install.cmake
CMakeFiles           TempControl.h         isocpp_idlpp
Makefile             TempControlSplDcps.cpp tsub
TempControl-cyclone.c TempControlSplDcps.h  tssub
TempControl-cyclone.h TempControl_DCPS.hpp
[→ build git:(master) * make
[ 0%] Built target isocpp_idlpp
[ 6%] Built target temp_ctrl_idl_isocpp_generate
[ 53%] Built target tssub
[100%] Built target tsub
[→ build git:(master) * ]
```

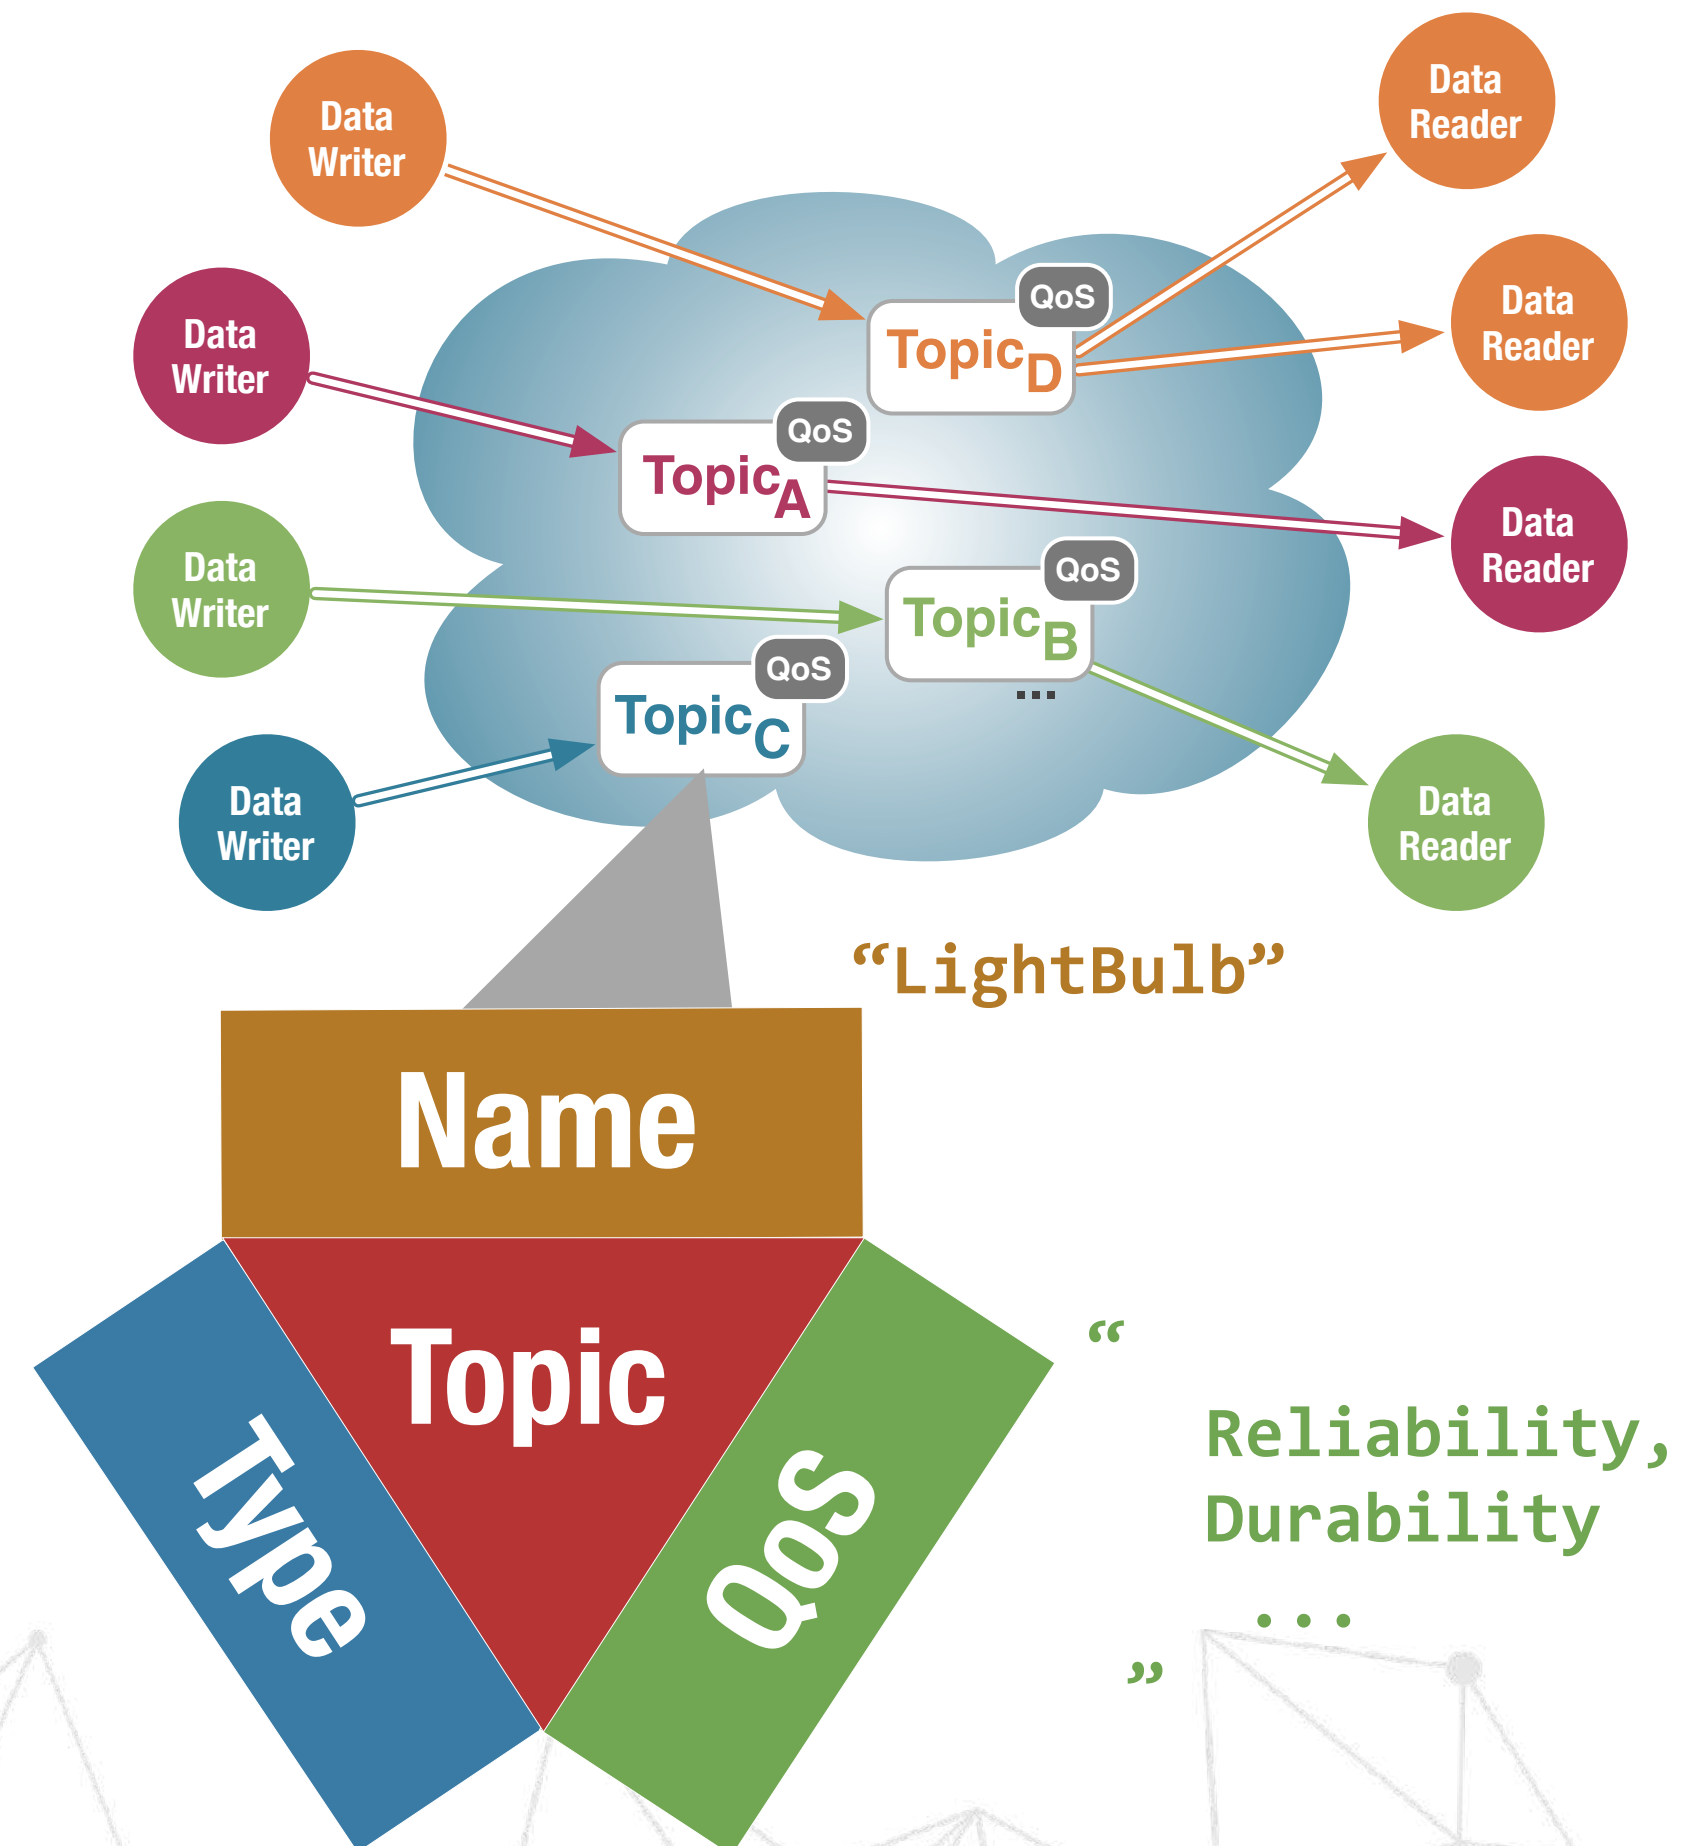

DDS: Going Deeper

Topics

DDS data streams are defined by means of **Topics**

A **Topic** represented is by means of a <name, type, qos>

```
struct LightBulbState {  
    string sn;  
    float luminosity;  
    long hue;  
    boolean on;  
};  
#pragma keylist LightBulbState sn
```

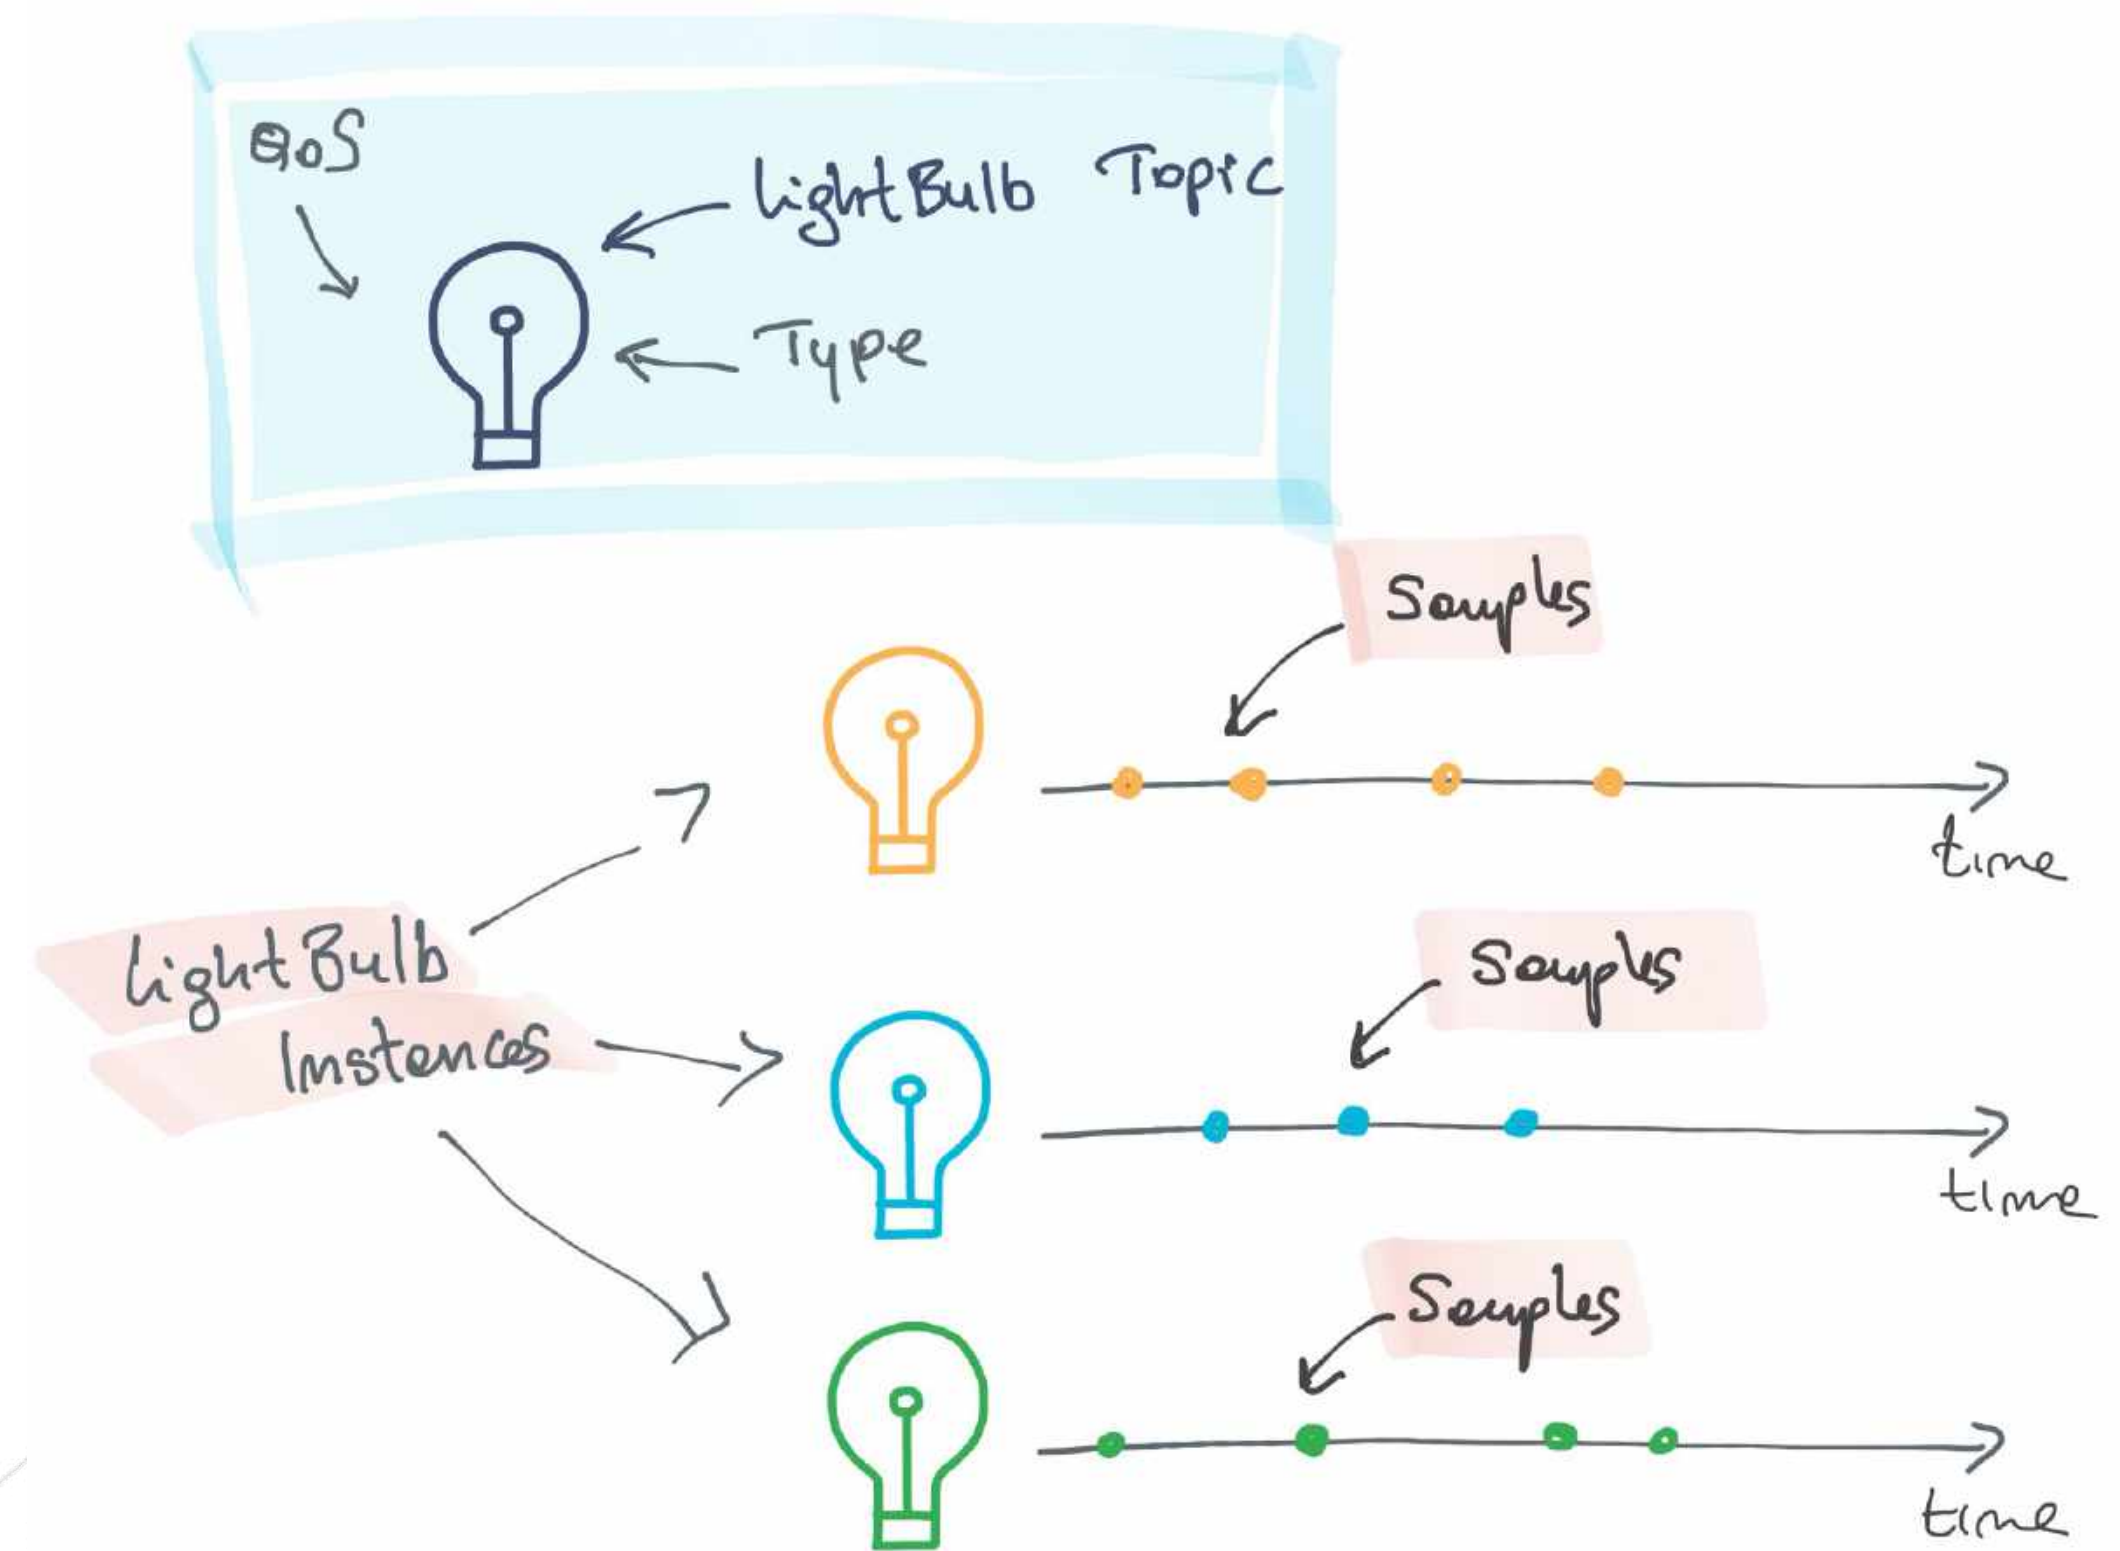


Topic Instances

Topic may **mark some** of their associated type **attributes** as **key-fields**

Each unique key value (tuple of key attributes) identifies a Topic Instance. Each Topic Instance has associated a FIFO ordered stream of samples


DDS provides useful **instance life-cycle management** and samples demultiplexing



Topics and Relations

A **topic** can be seen as defining a **relation**

```
struct LightBulbState {  
    string sn;  
    float luminosity;  
    long hue;  
    boolean on;  
};  
#pragma keylist LightBulbState sn
```




<i><u>sn</u></i>	<i>luminosity</i>	<i>hue</i>	<i>on</i>
<i>a123-21ef</i>	<i>0.5</i>	<i>12750</i>	<i>TRUE</i>
<i>600d-caf3</i>	<i>0.8</i>	<i>46920</i>	<i>FALSE</i>
<i>1234-c001</i>	<i>0.75</i>	<i>25500</i>	<i>TRUE</i>

Information Scopes

DDS information lives within a **domain**

A domain can be thought as organised in **partitions**

Samples belonging to a given **Topic Instance** are read/written from/in one or more **partitions**



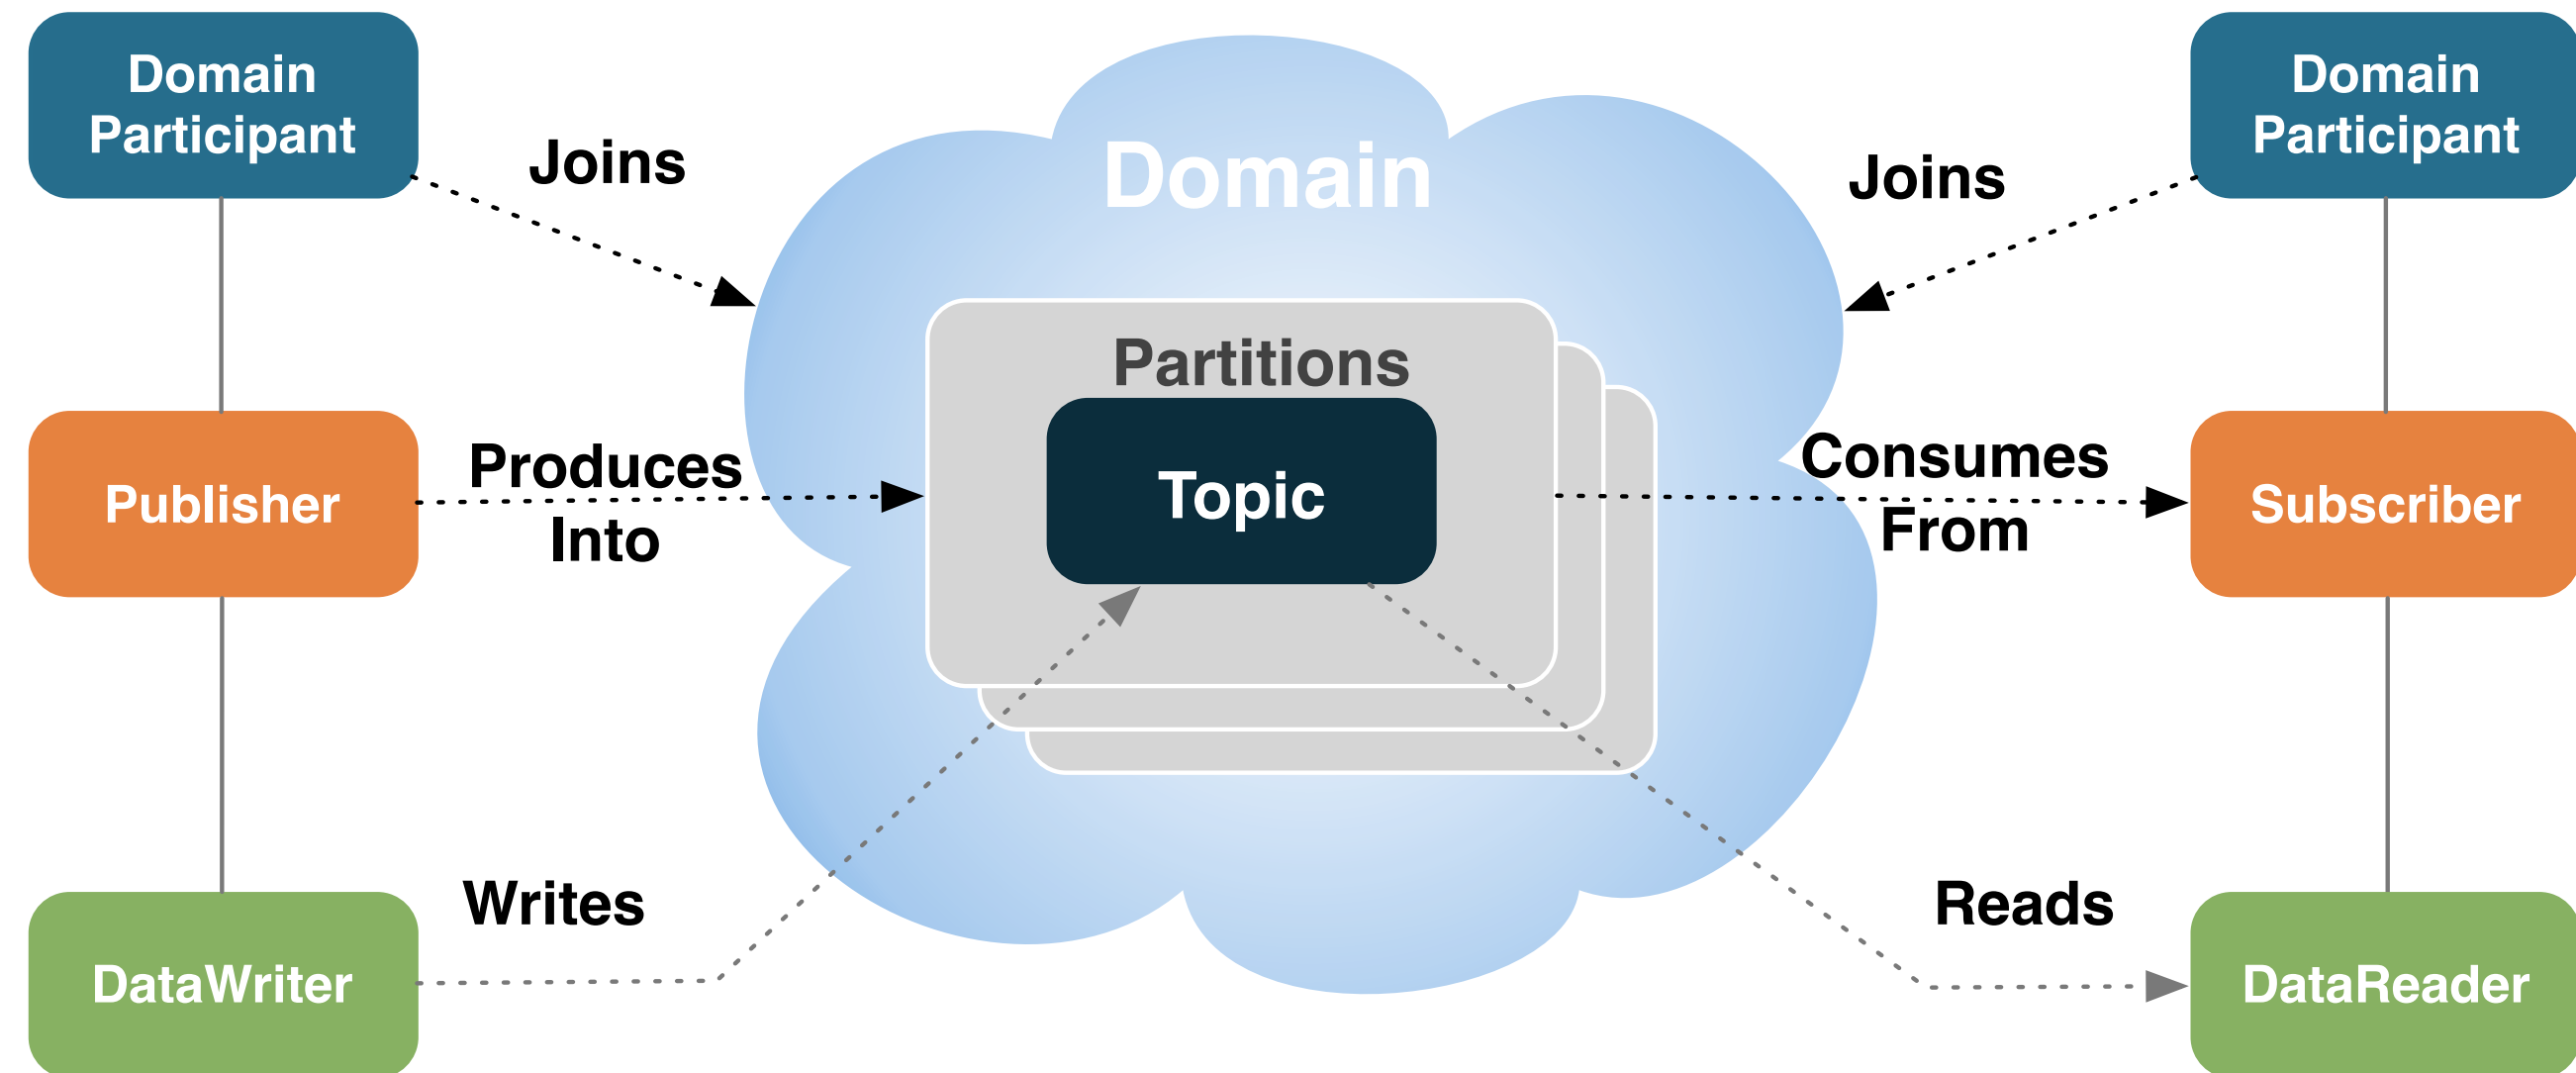
DDS Entities

DDS Entities

DDS provides three different entities to control **where** and **what** data is read/written

The **DomainParticipant**, **Publisher** and **Subscriber** relate to the "**where**"

DataReader and **DataWriter** relate to the "**what**"

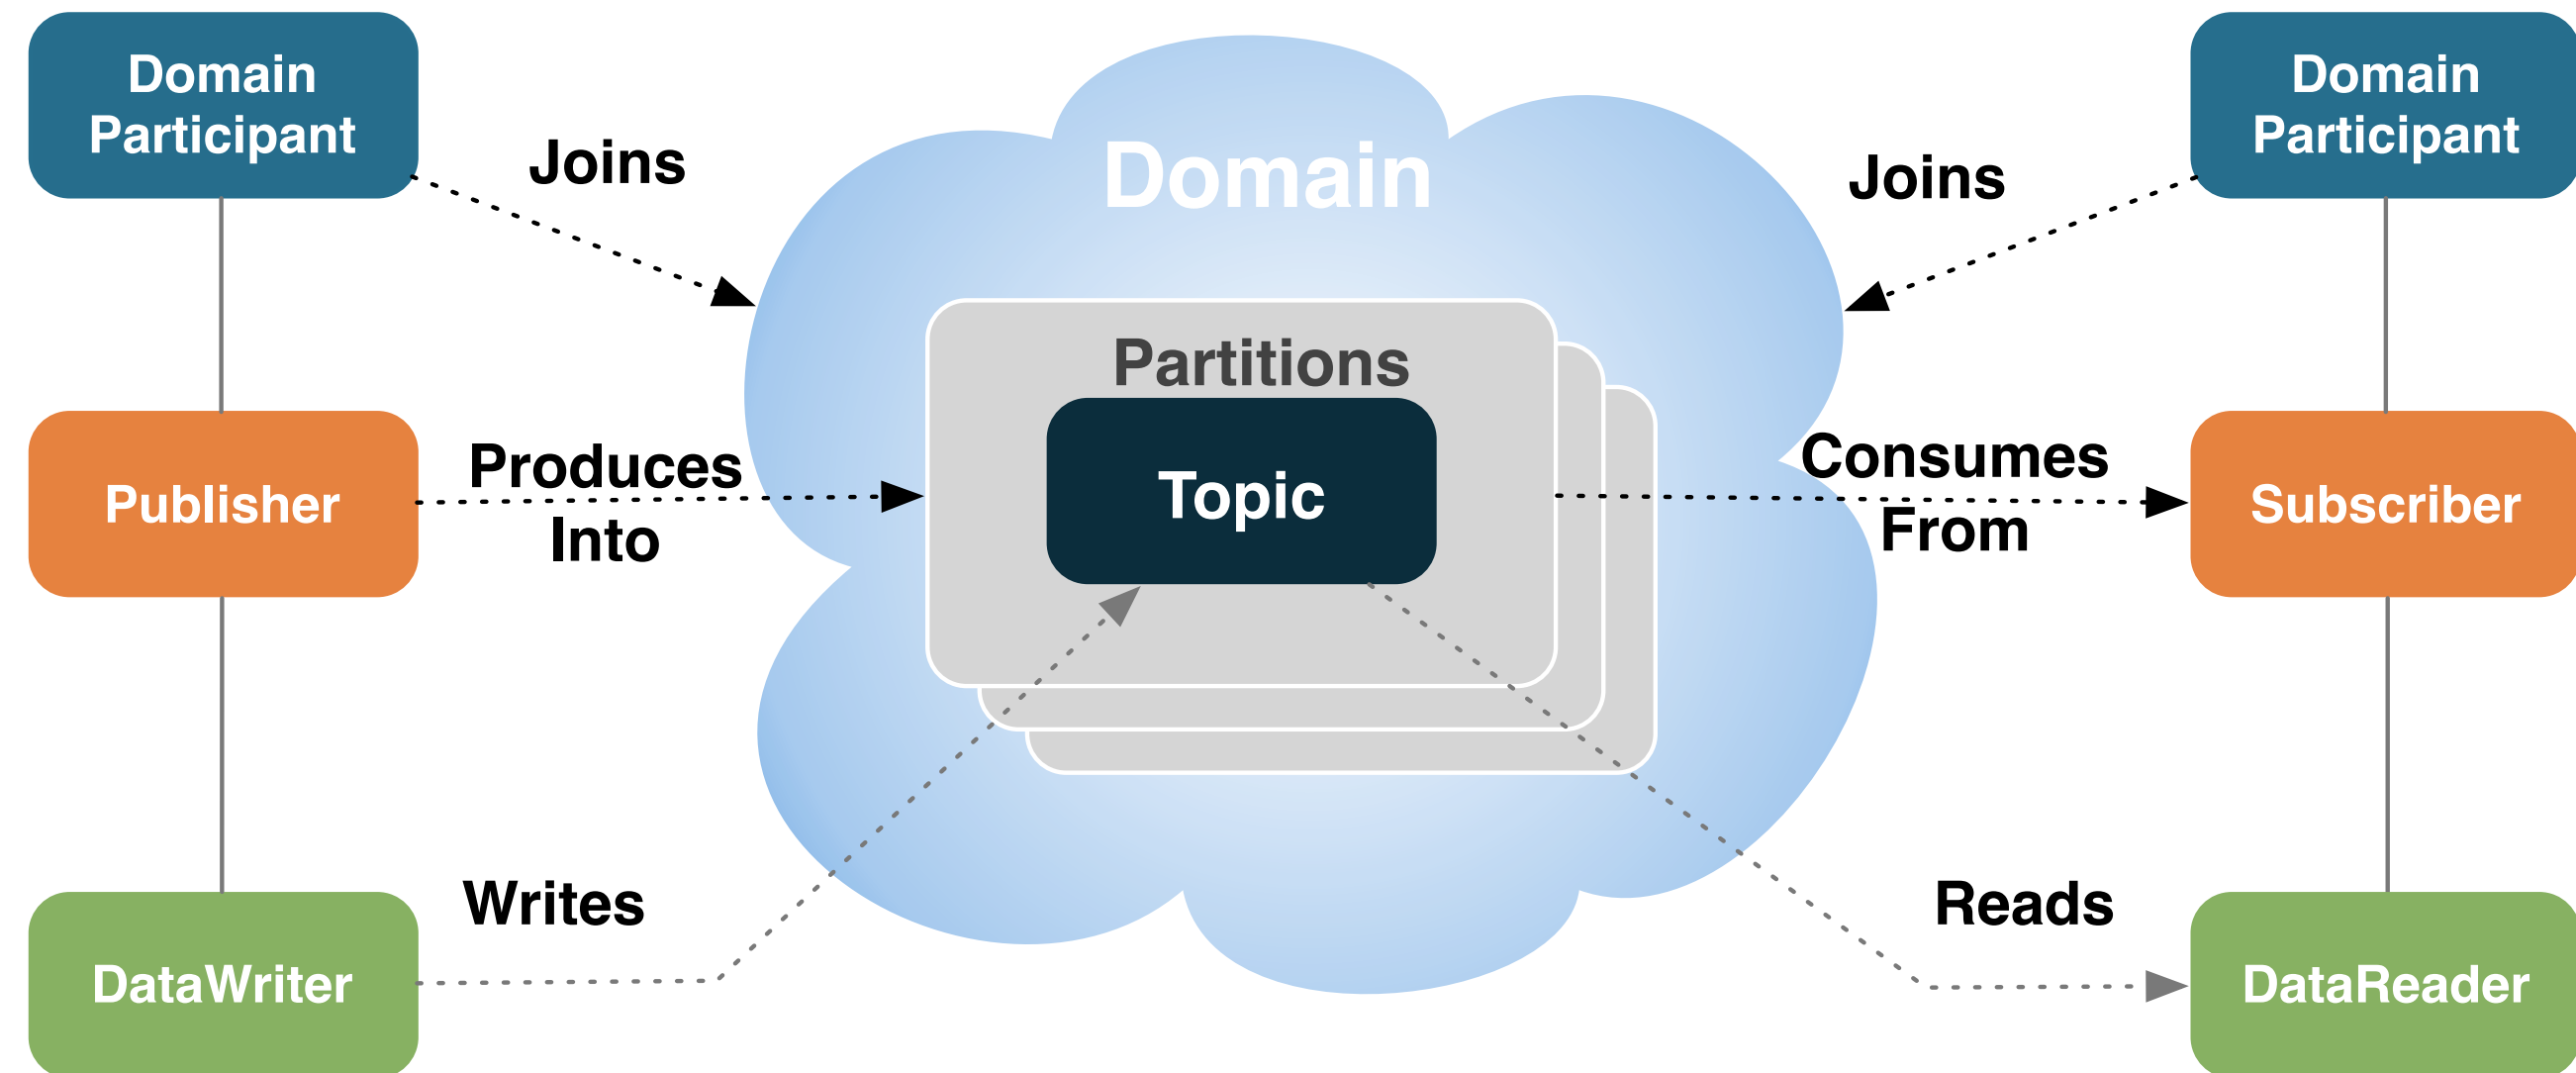


DDS Entities

DDS QoS Policies

control, at a large extent,
to the **"how" data is shared**

QoS Policies also control
resource utilisation



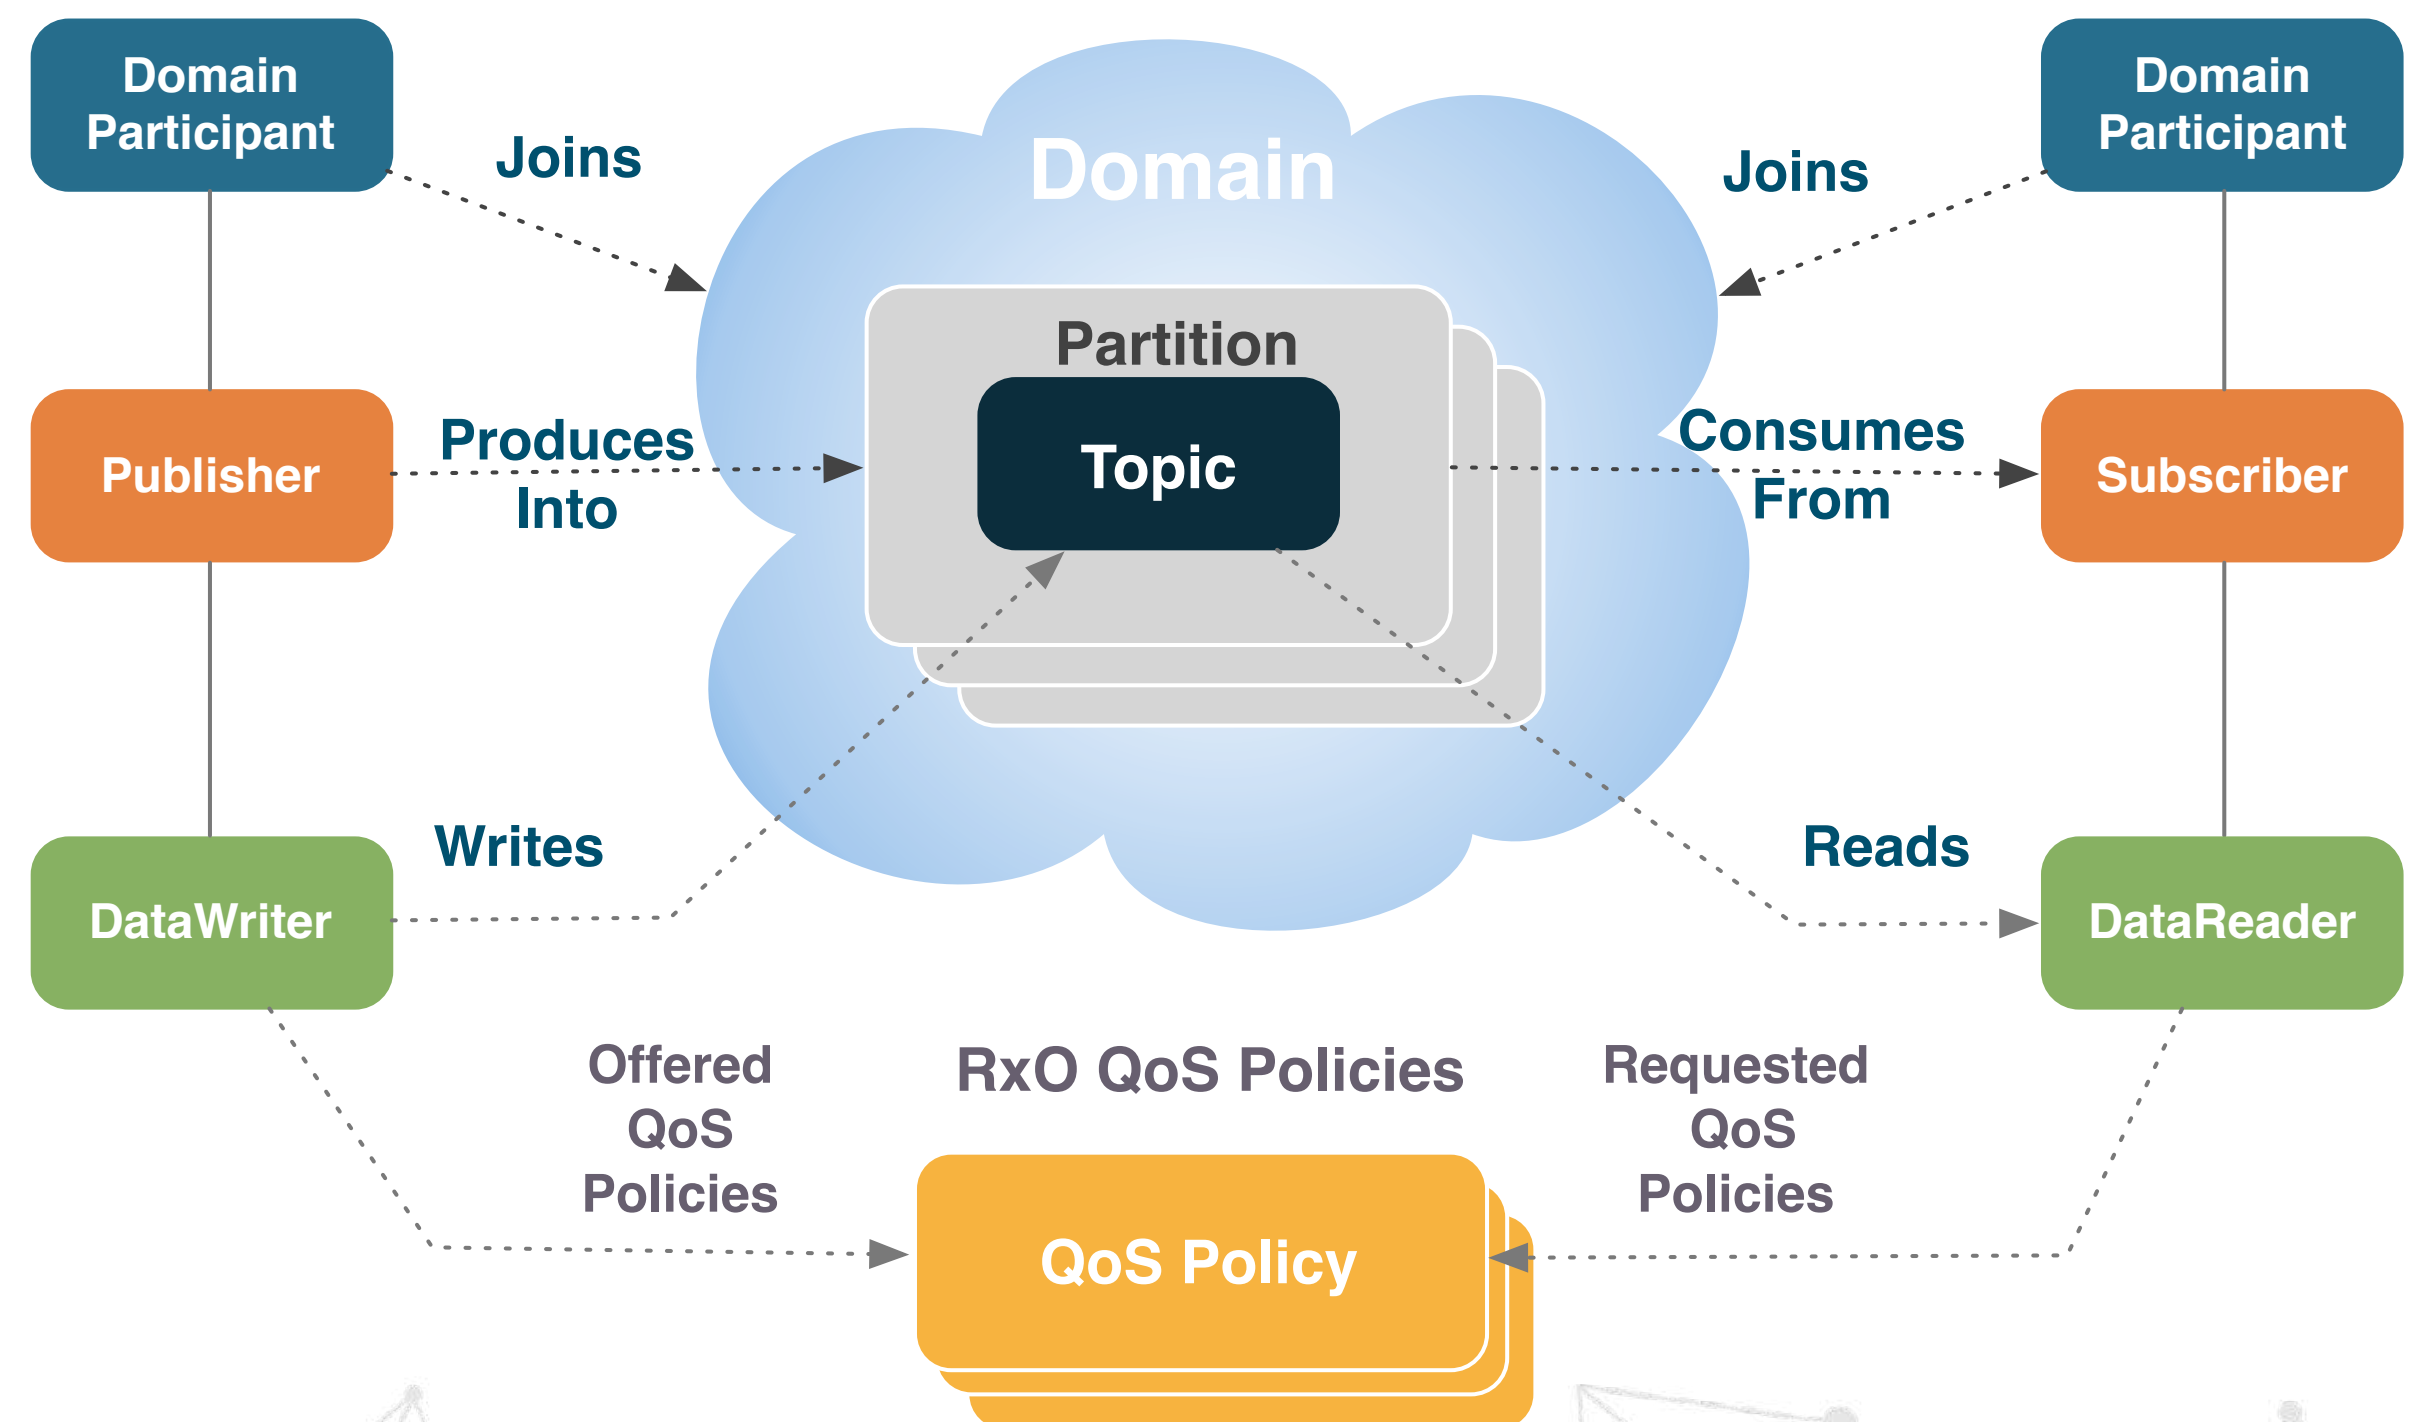
Matching Model

For data to flow from a DataWriter (DW) to one or many DataReader (DR) a few conditions have to apply:

The **DR** and **DW** have to be in the **same domain**

The **partition expression** of the DR's Subscriber and the DW's Publisher should **match** (in terms of regular expression match)

The **QoS Policies offered** by the DW should **exceed or match** those **requested** by the DR

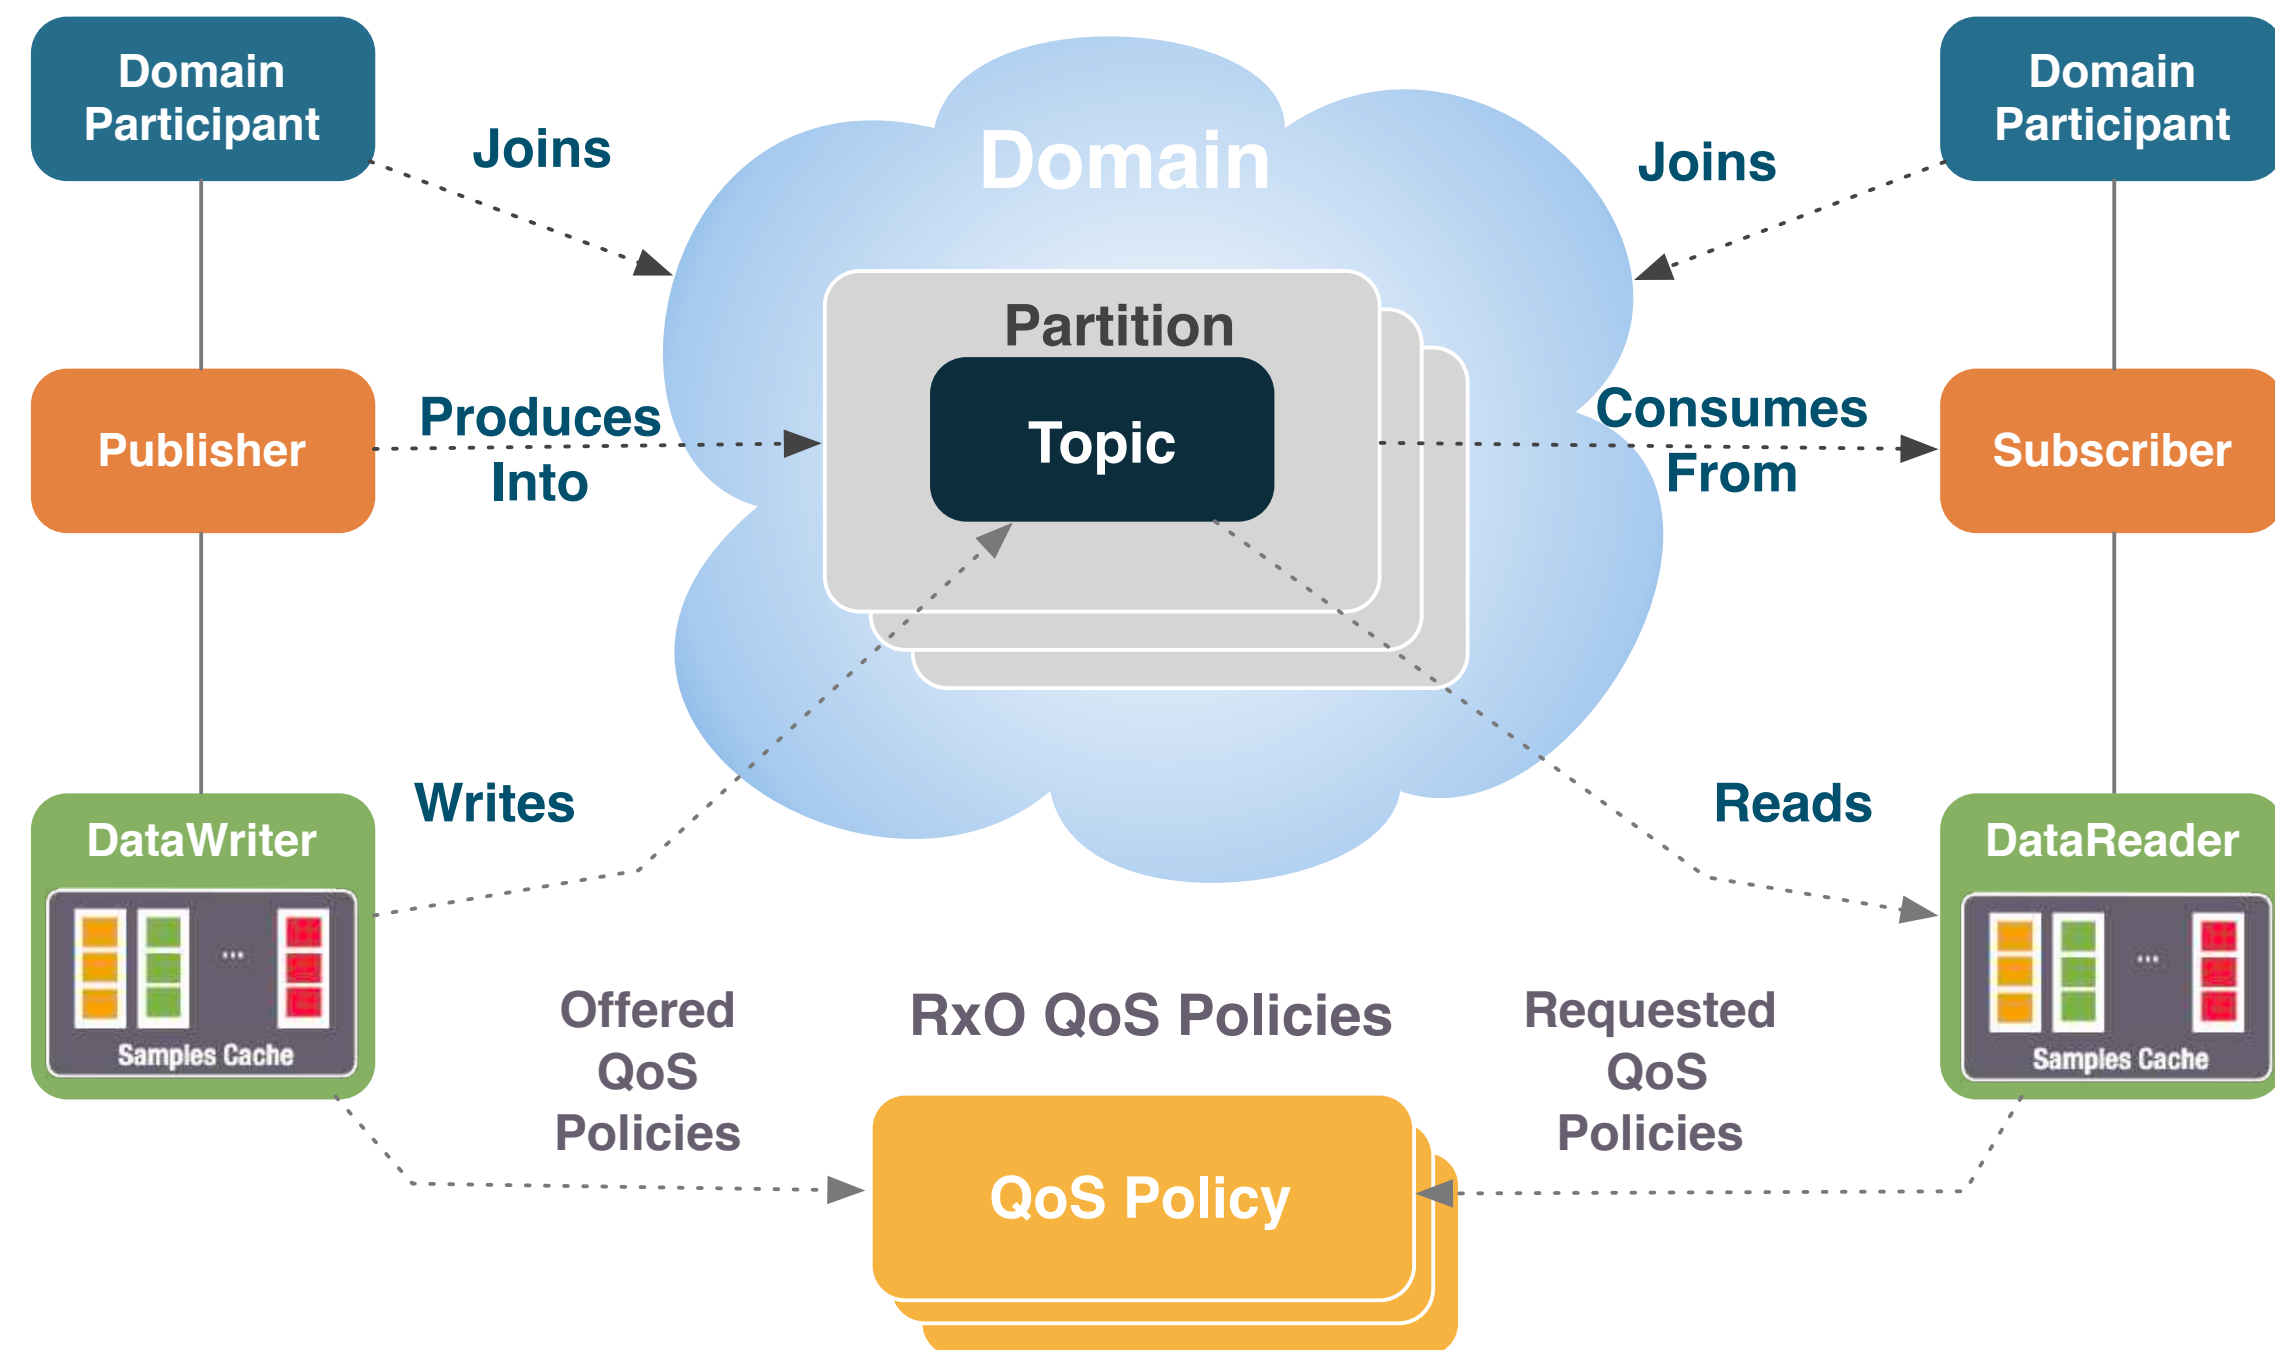


Storage Model

DataWriter and DataReaders have an associated samples cache

In a sense, what DDS does is to project, eventually, the relevant content of the writer cache into matching reader caches

As a consequence of these caches reads and writes are always local and non-blocking*



* **reads** never block and a write will only block, depending on QoS settings if sufficient resources are not available

Content Awareness

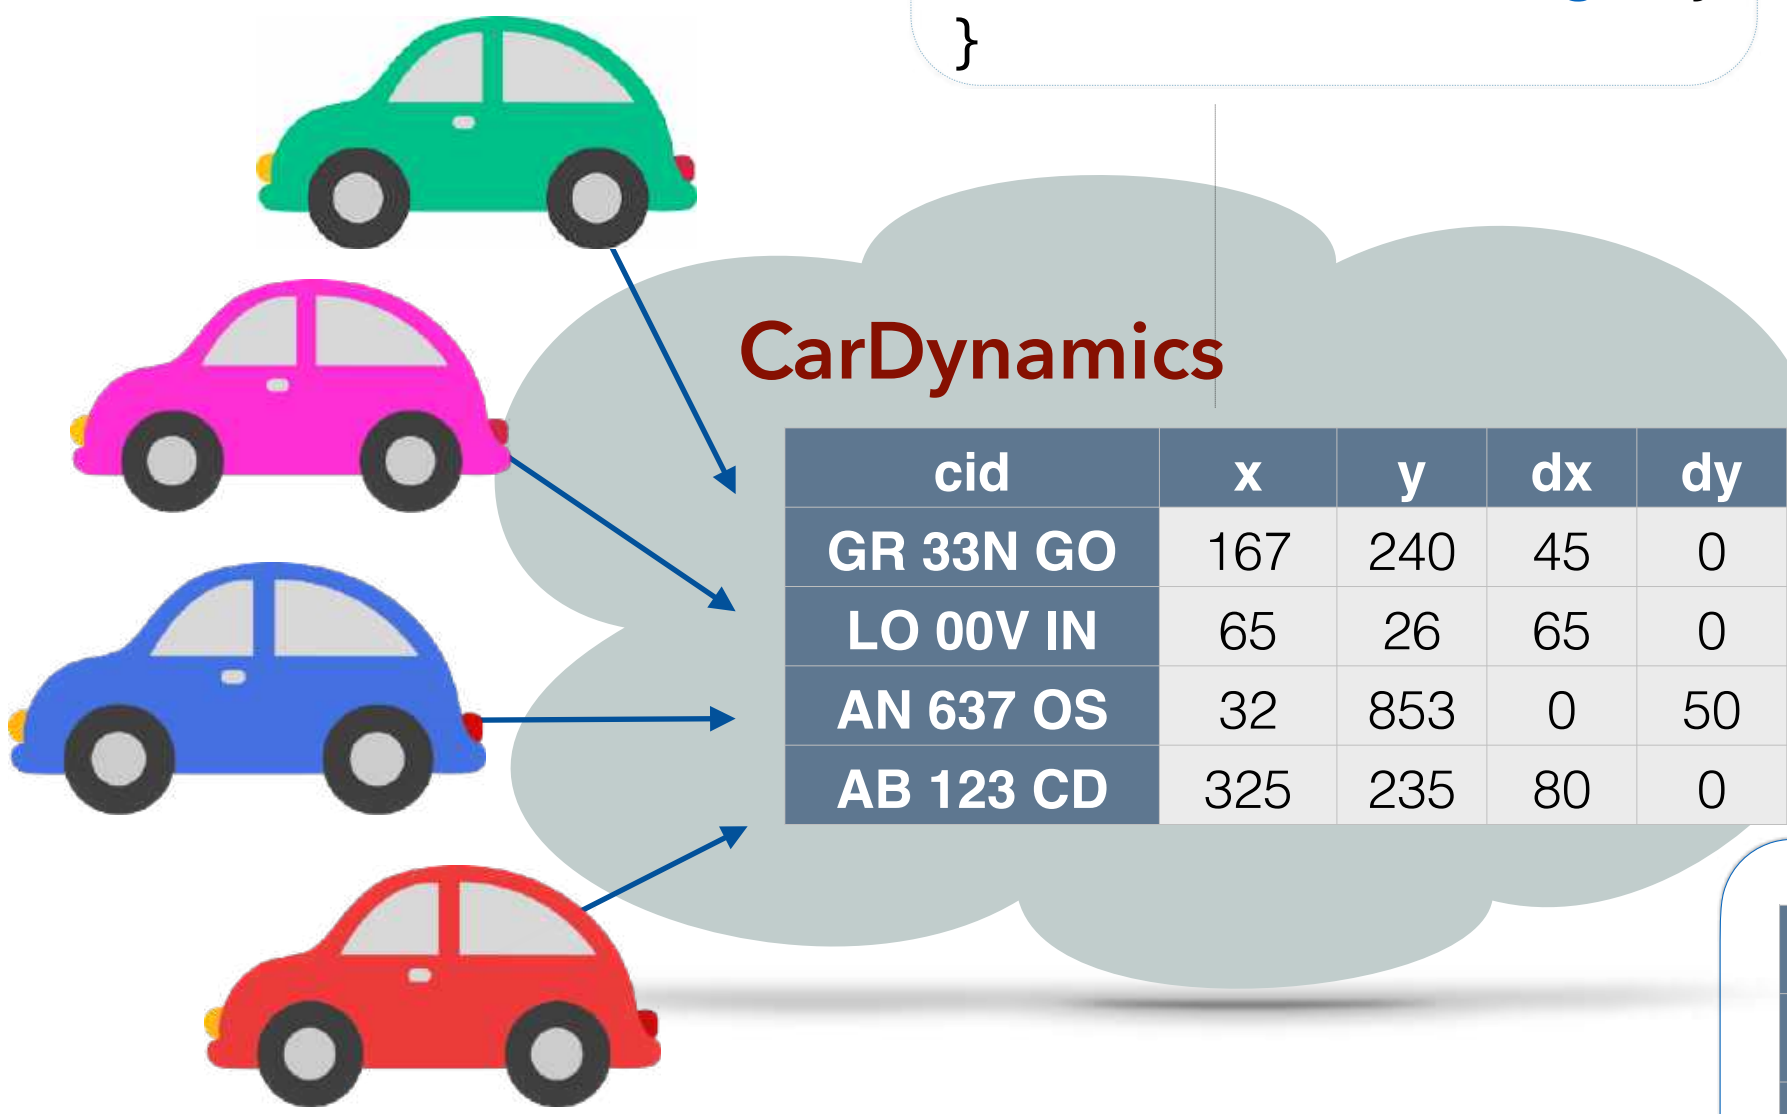
Content Filtering

Content Filters

can be used to **project** on the local cache only the Topic data satisfying a given predicate

Type

```
struct CarDynamics {
  @key
  string  cid;
  long    x;   long    y;
  float   dx;  long    dy;
}
```



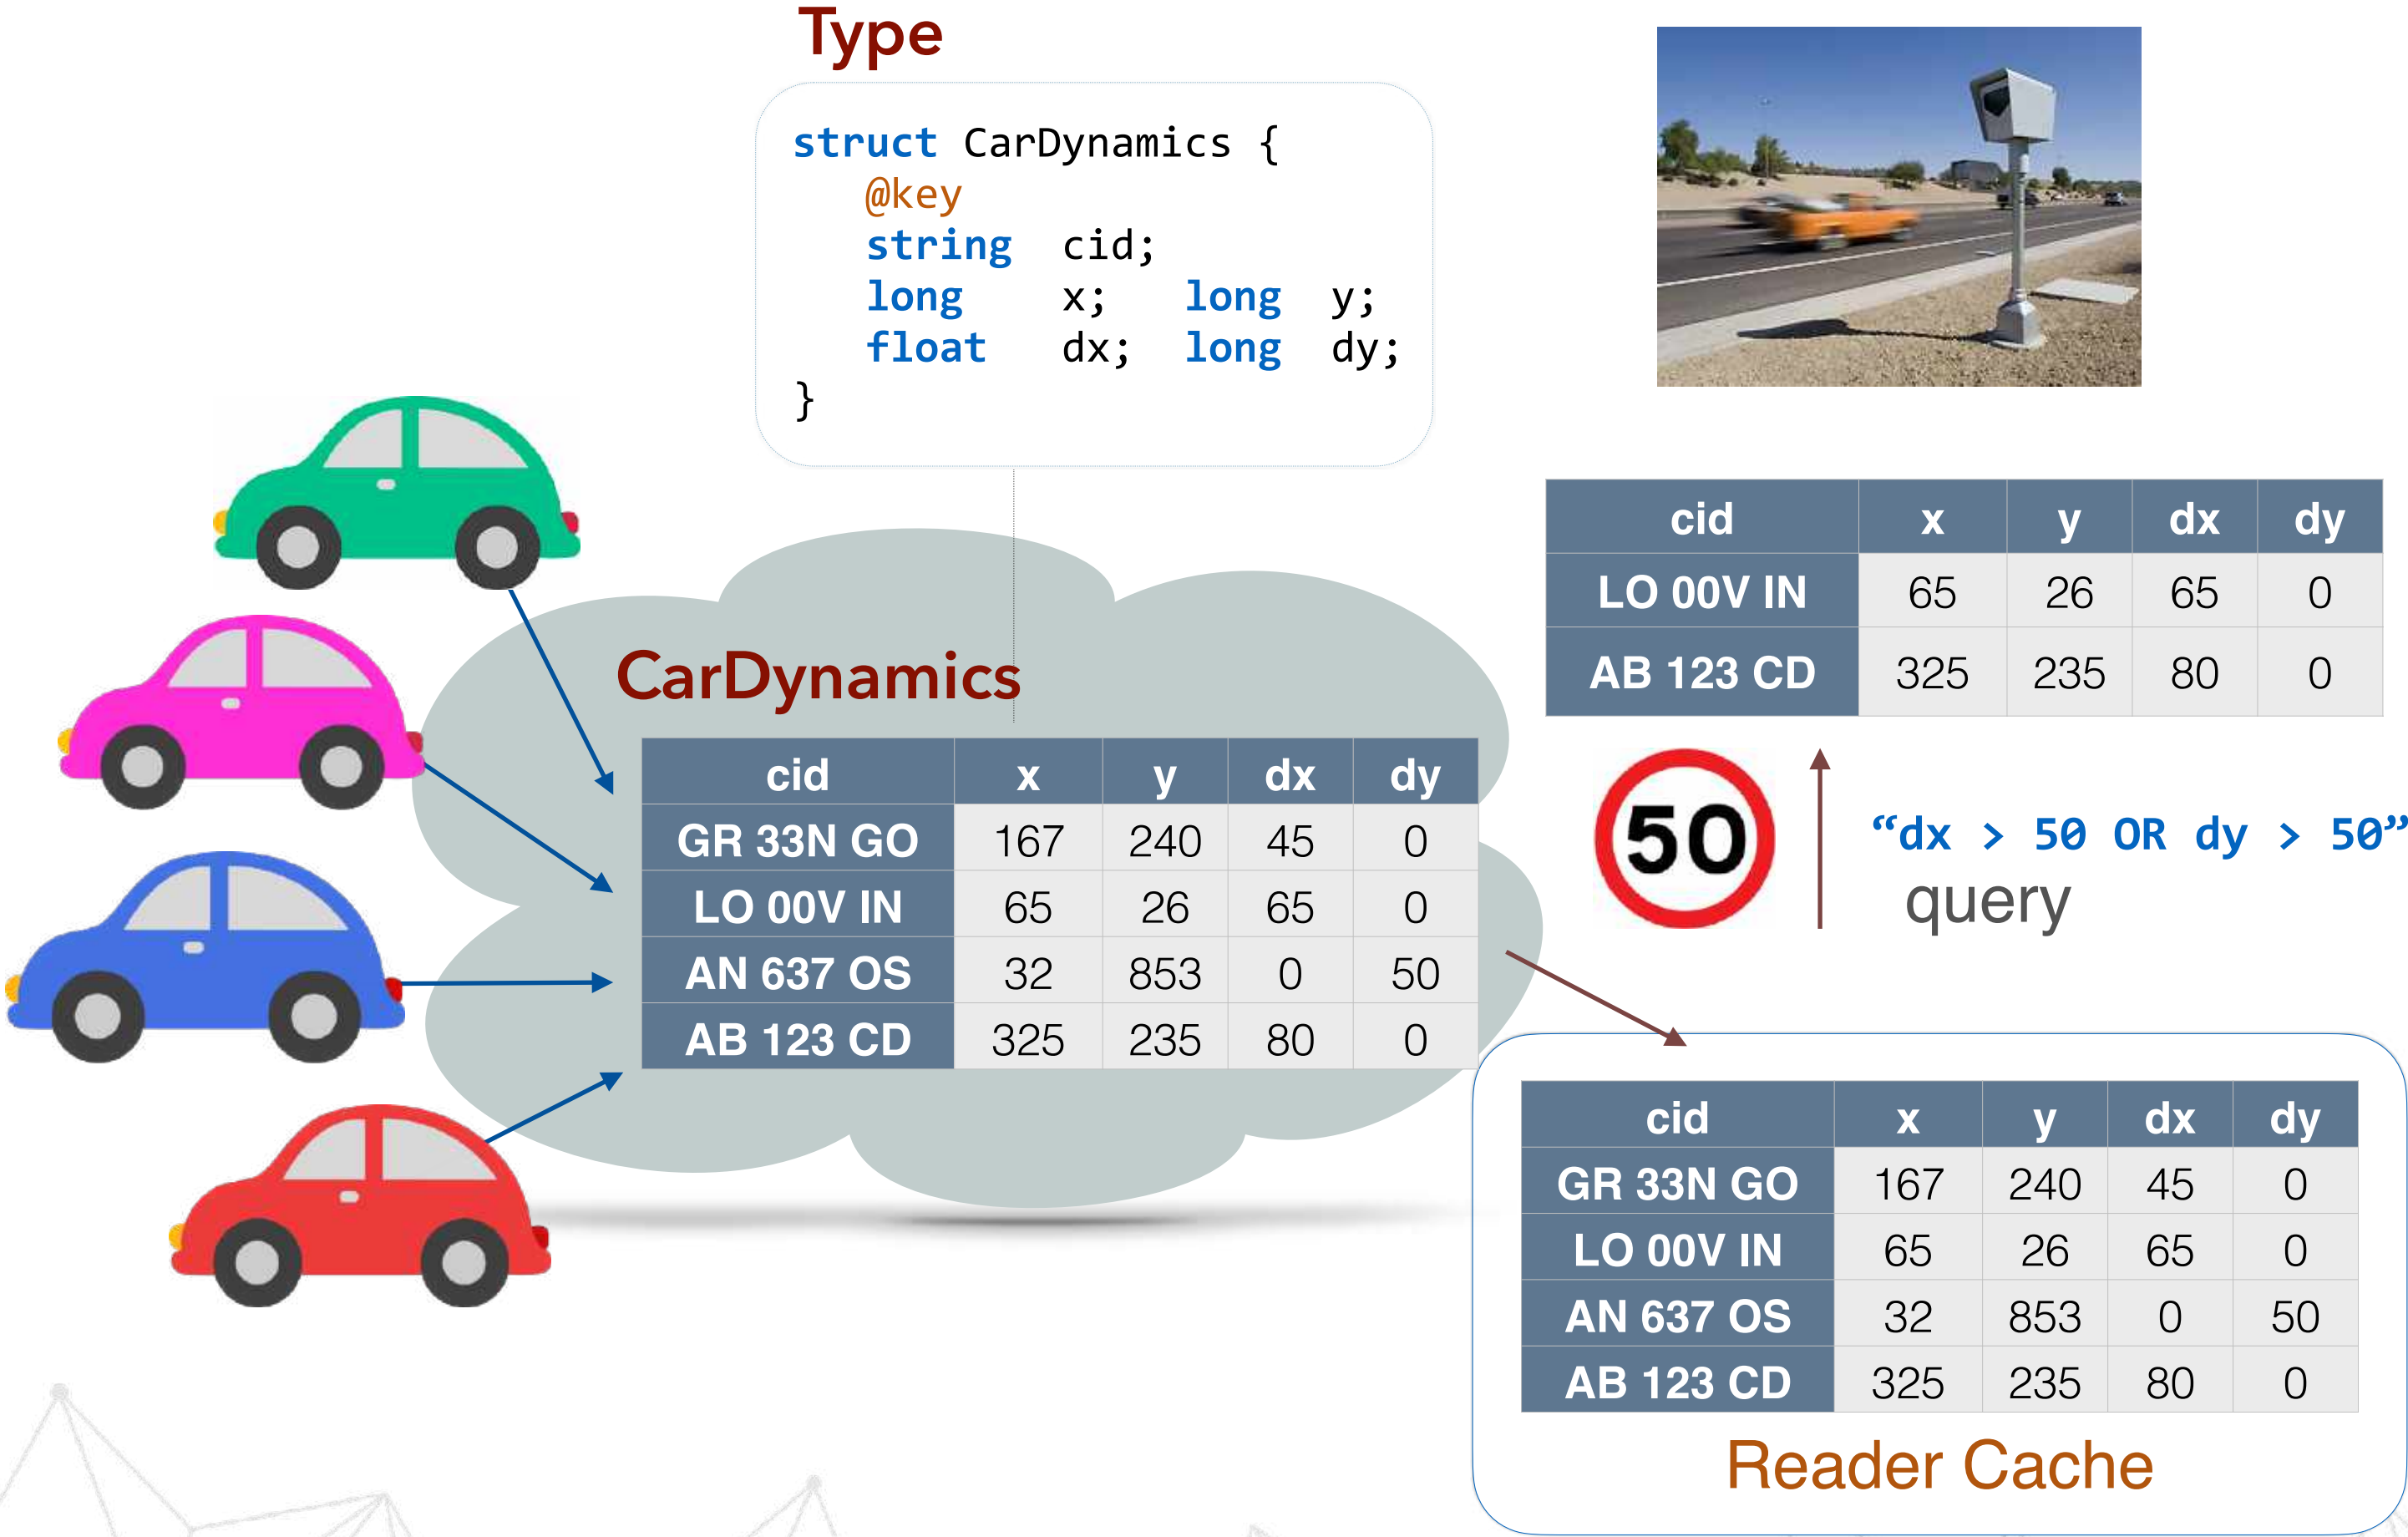
“dx > 50 OR dy > 50”

cid	x	y	dx	dy
LO 00V IN	65	26	65	0
AB 123 CD	325	235	80	0

Reader Cache

Queries

Queries can be used to select out of the local cache the data matching a given predicate




Stream Durability

Stream Durability

Through QoS settings it is possible to control which subset of the stream data will be **retained** and **made available** (*replayed*) **to late joiners**

DDS can store the last **n** samples ($n=1$ is a special case) or all the samples written for a topic



Stream Durability


DDS provides three kinds of durability:

Volatile: i.e. no durability


Transient (Local): data is available for late joiners (re-play) as far as the system (data source) is running

Durable: data is available for late joiners (re-play) as far as the system/source is running

Volatile Durability



Transient Durability



Stream Reliability

Best Effort

DDS will deliver an arbitrary subsequence of the samples written against a Topic Instance

Samples may be dropped because of network loss or because of flow-control

Best Effort




Last n-values Reliability

Under stationary conditions an application is guaranteed to receive the last n-samples written for a Topic Instance

Samples falling outside the history may be dropped at the sending or receiving side for flow/resource control


Notice that this kind of reliability behaves as a circuit breaker for slow consumers

Last 3-values Reliability



Reliable

All samples written against a Topic Instance are delivered. Since from a theoretical perspective reliability in asynchronous systems either violate progress or requires infinite memory, DDS provides QoS to control both resources as well as blocking time



Fault-Tolerance



Failure Detection

DDS provides mechanism for detecting traditional faults as well as performance failures

The Fault-Detection mechanism is controlled by means of the DDS

Liveliness policy


Performance Failures can be detected using the Deadline Policy which allows to receive notification when data is not received within the expected delays



Fault-Masking

DDS provides a built-in fault-masking mechanism that allow to replicate **Sources** and transparently switch over when a failure occurs

At any point in time the “active” source is the one with the highest strength. Where the strength is an integer parameter controller by the user

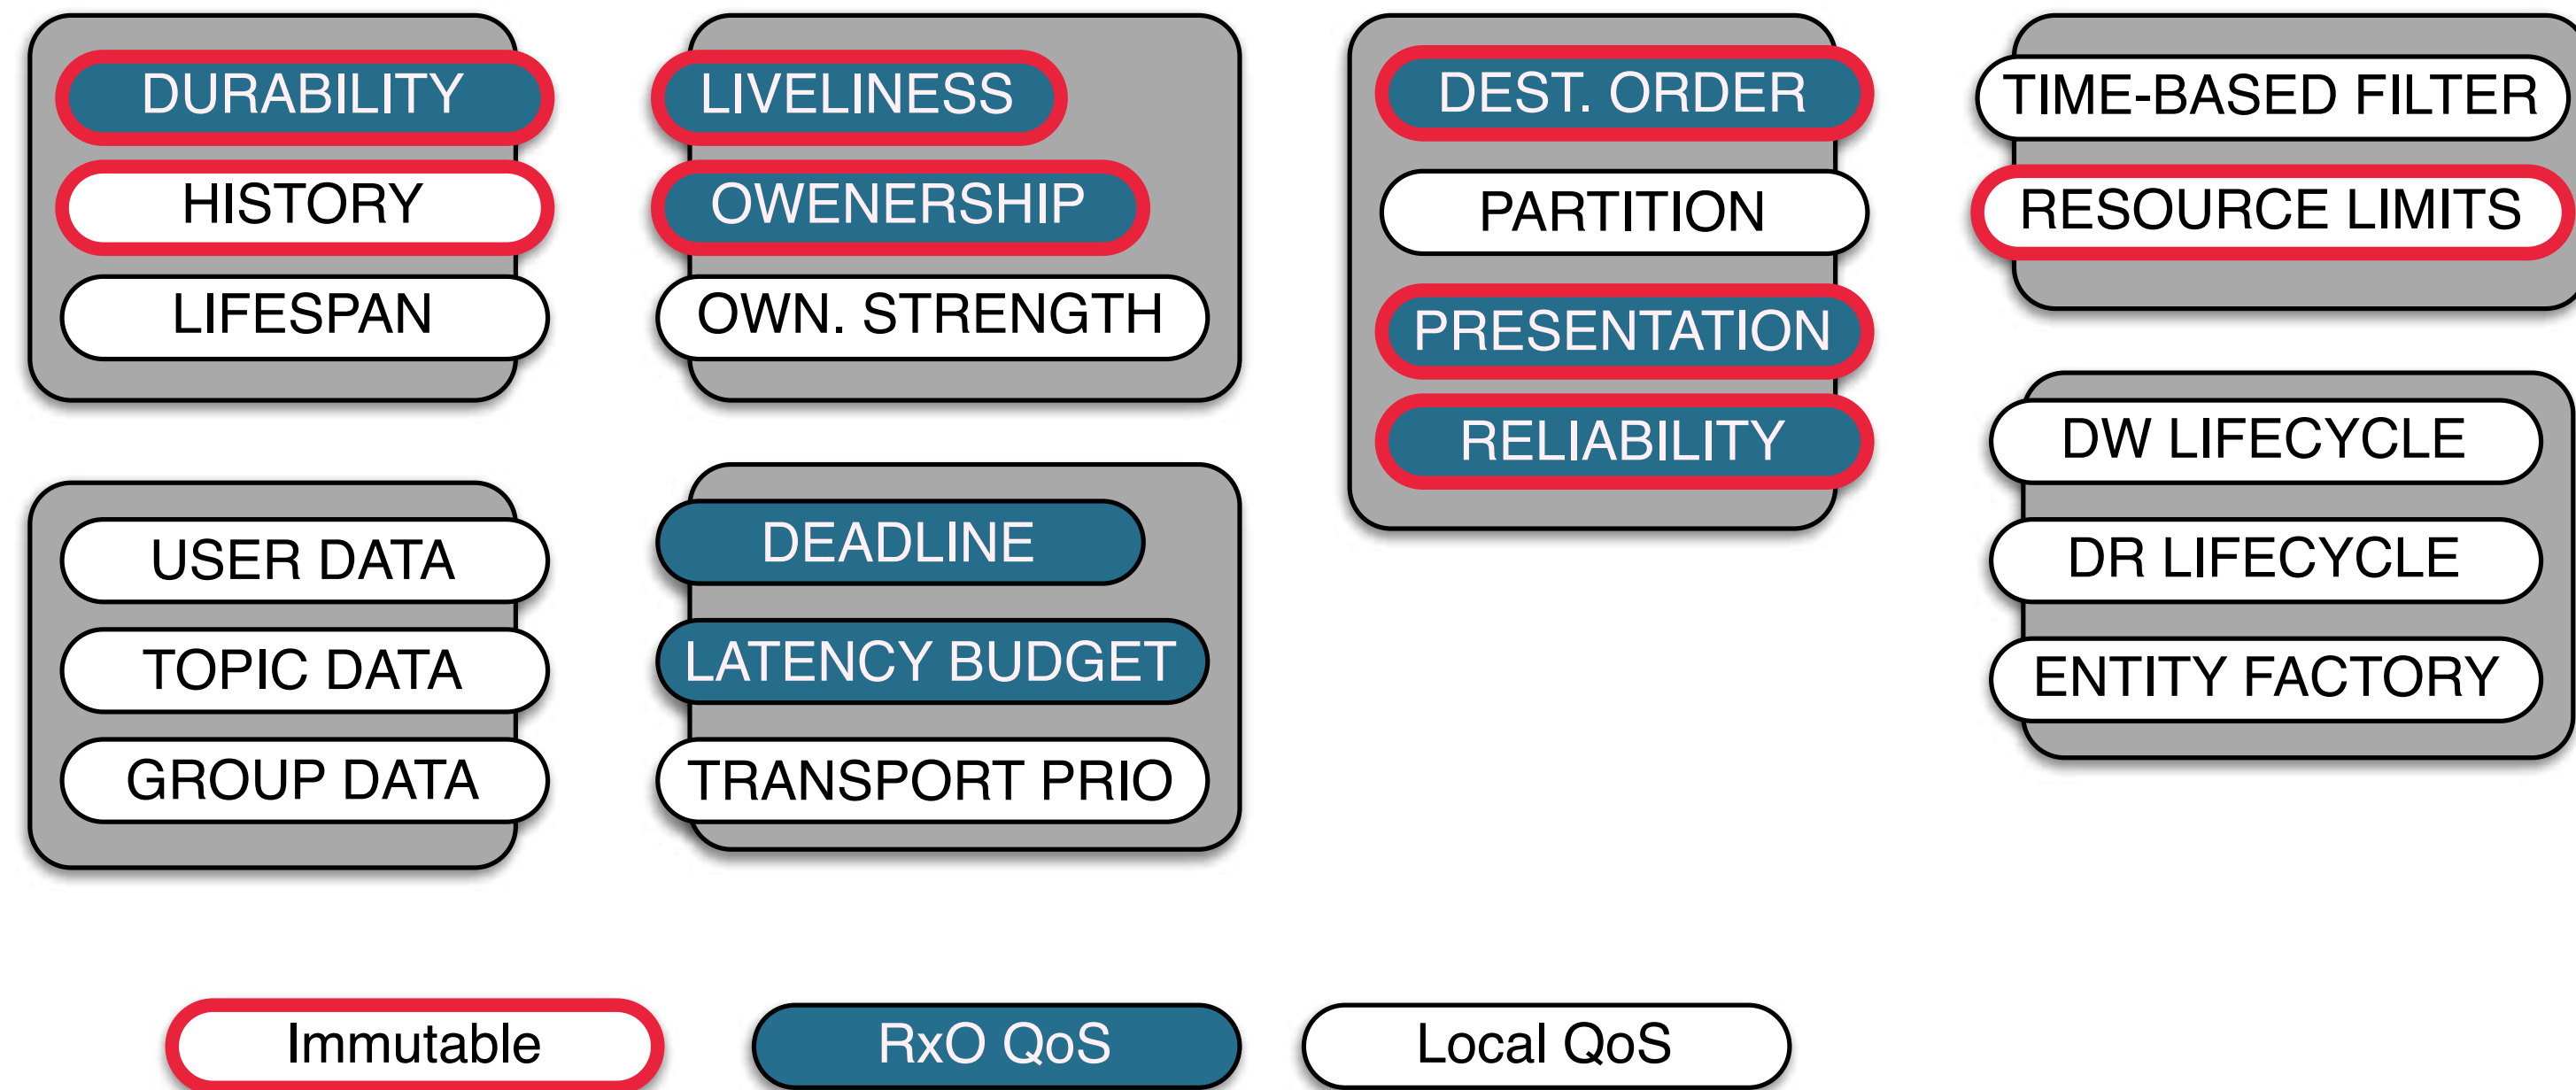


QoS Policies

DDS QoS Policies

DDS provides 20+ standard QoS Policies

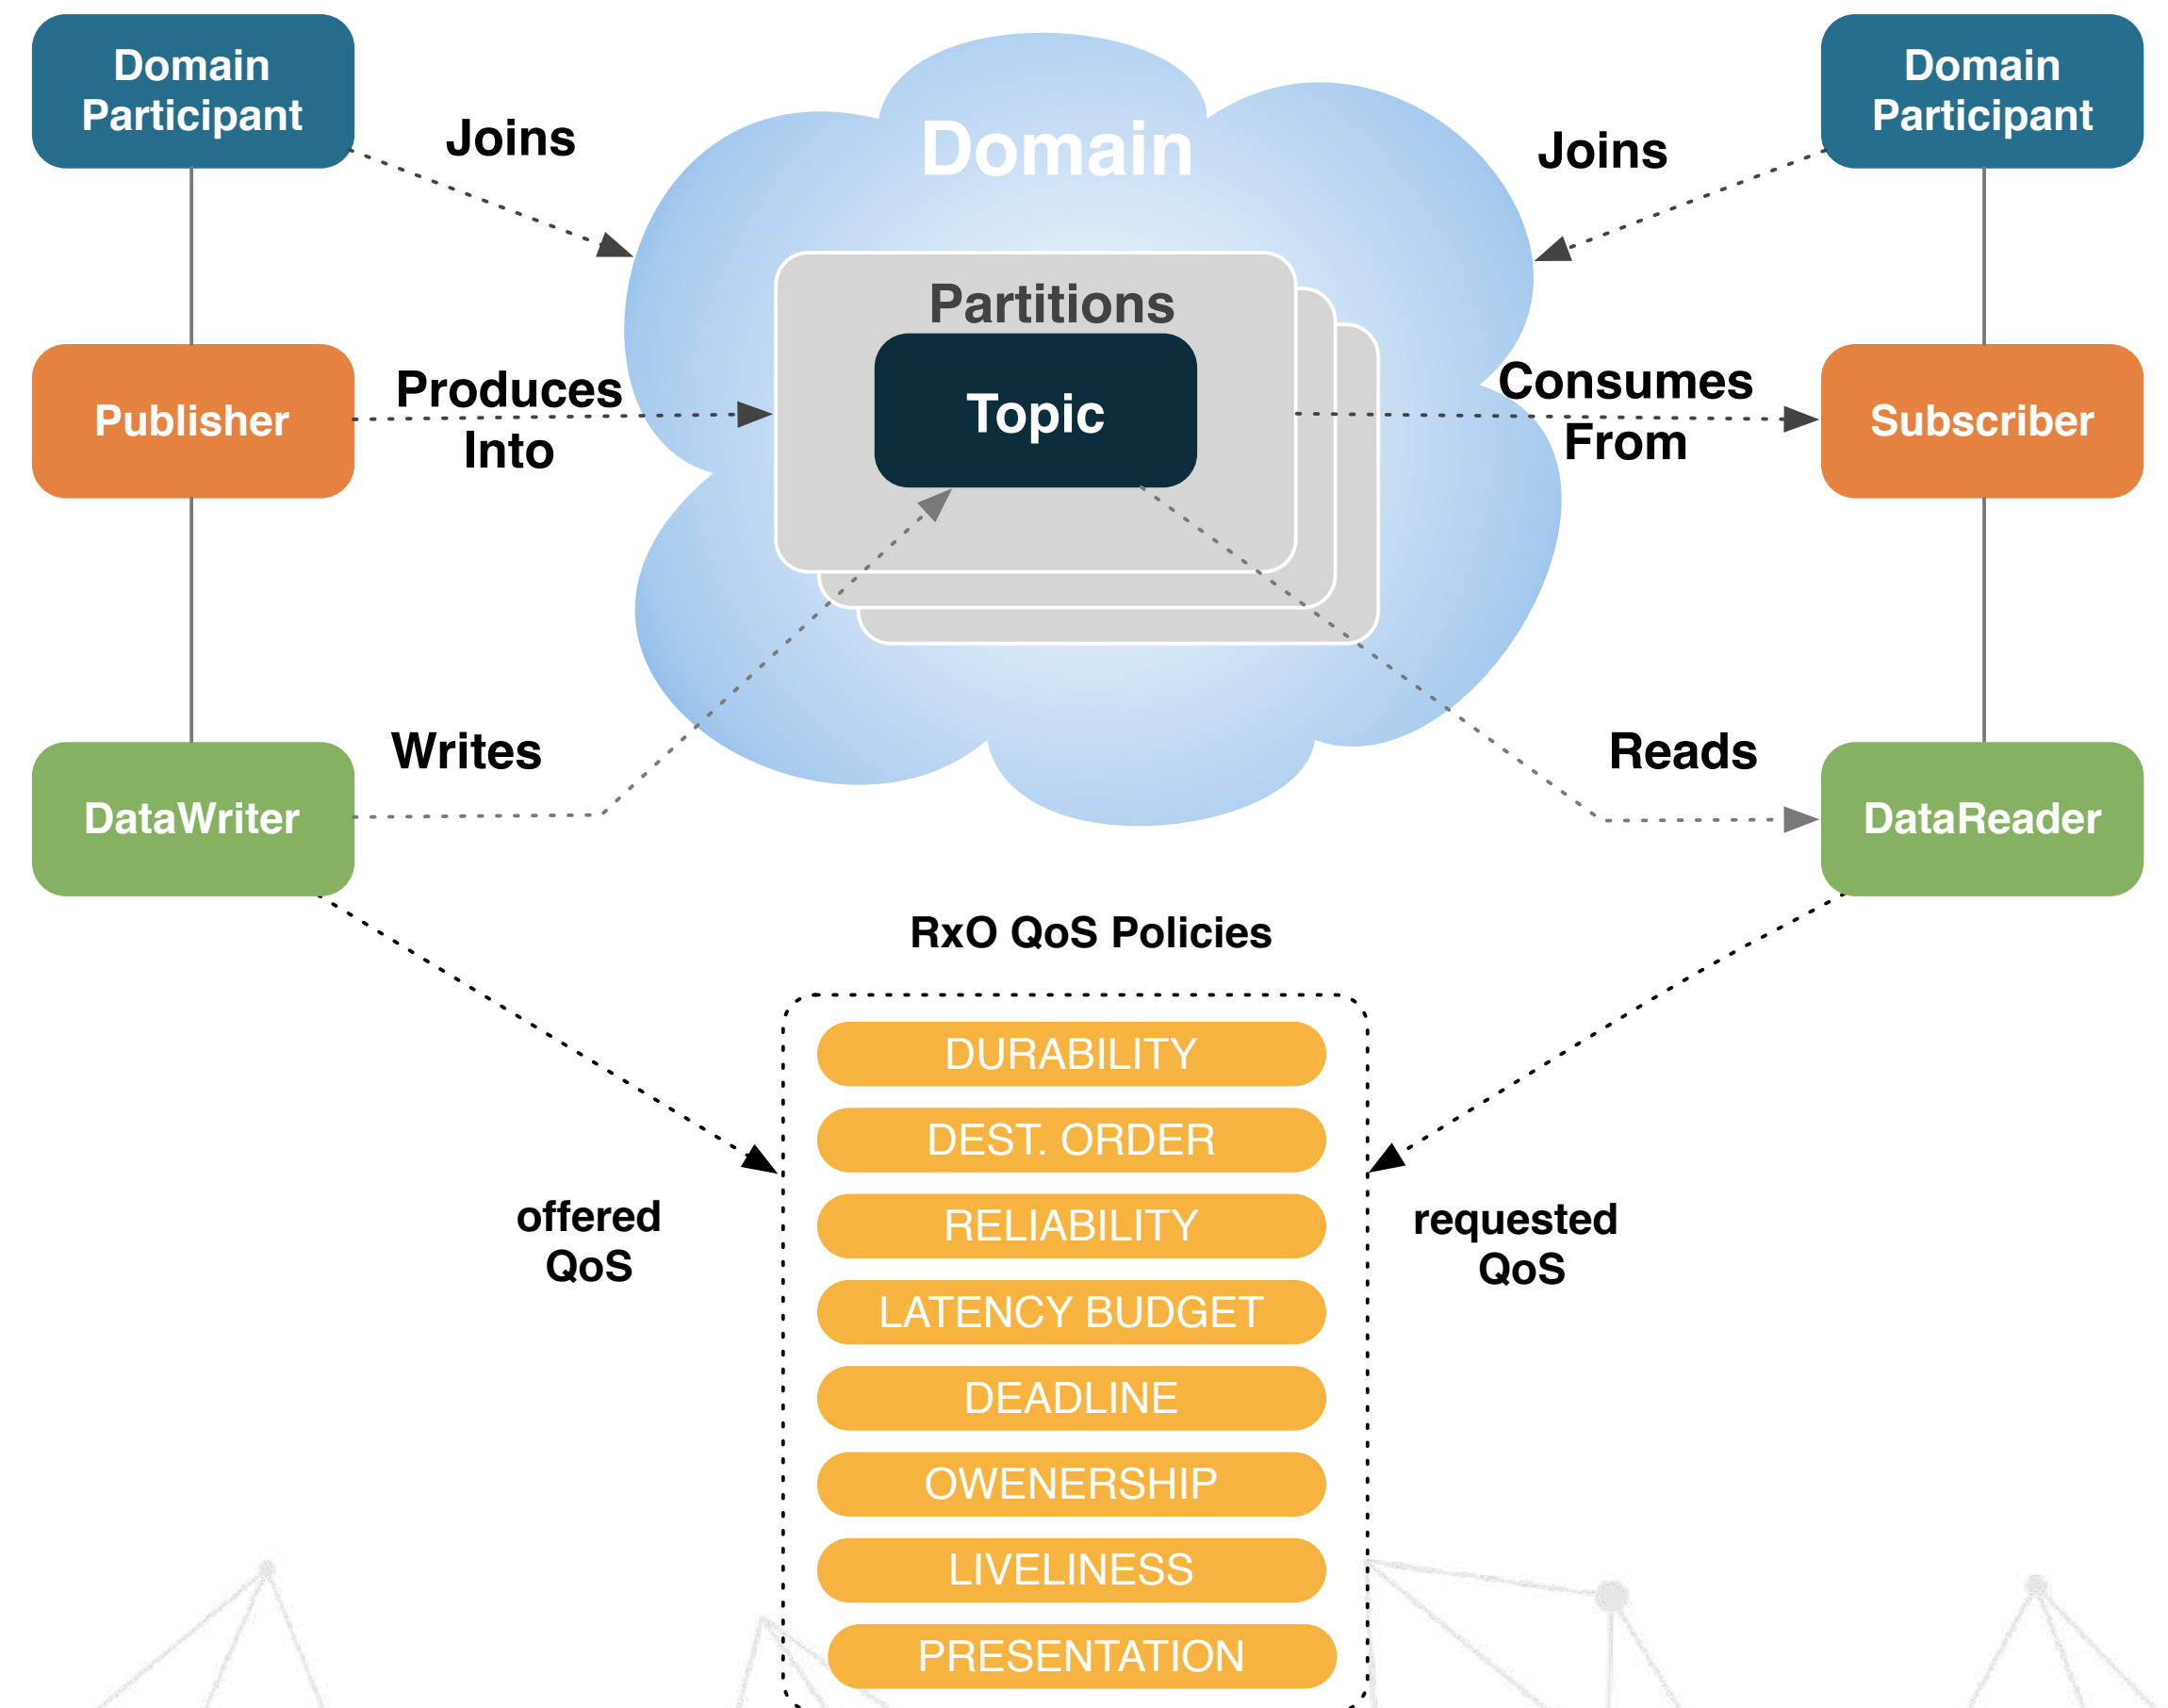
While this may seem a lot of complexity they are often used in combination to achieve certain patterns



RxO Model

A subset of the QoS Policies impact the matching

Most of the QoS relate to end-to-end properties of data




DDS in Robotics

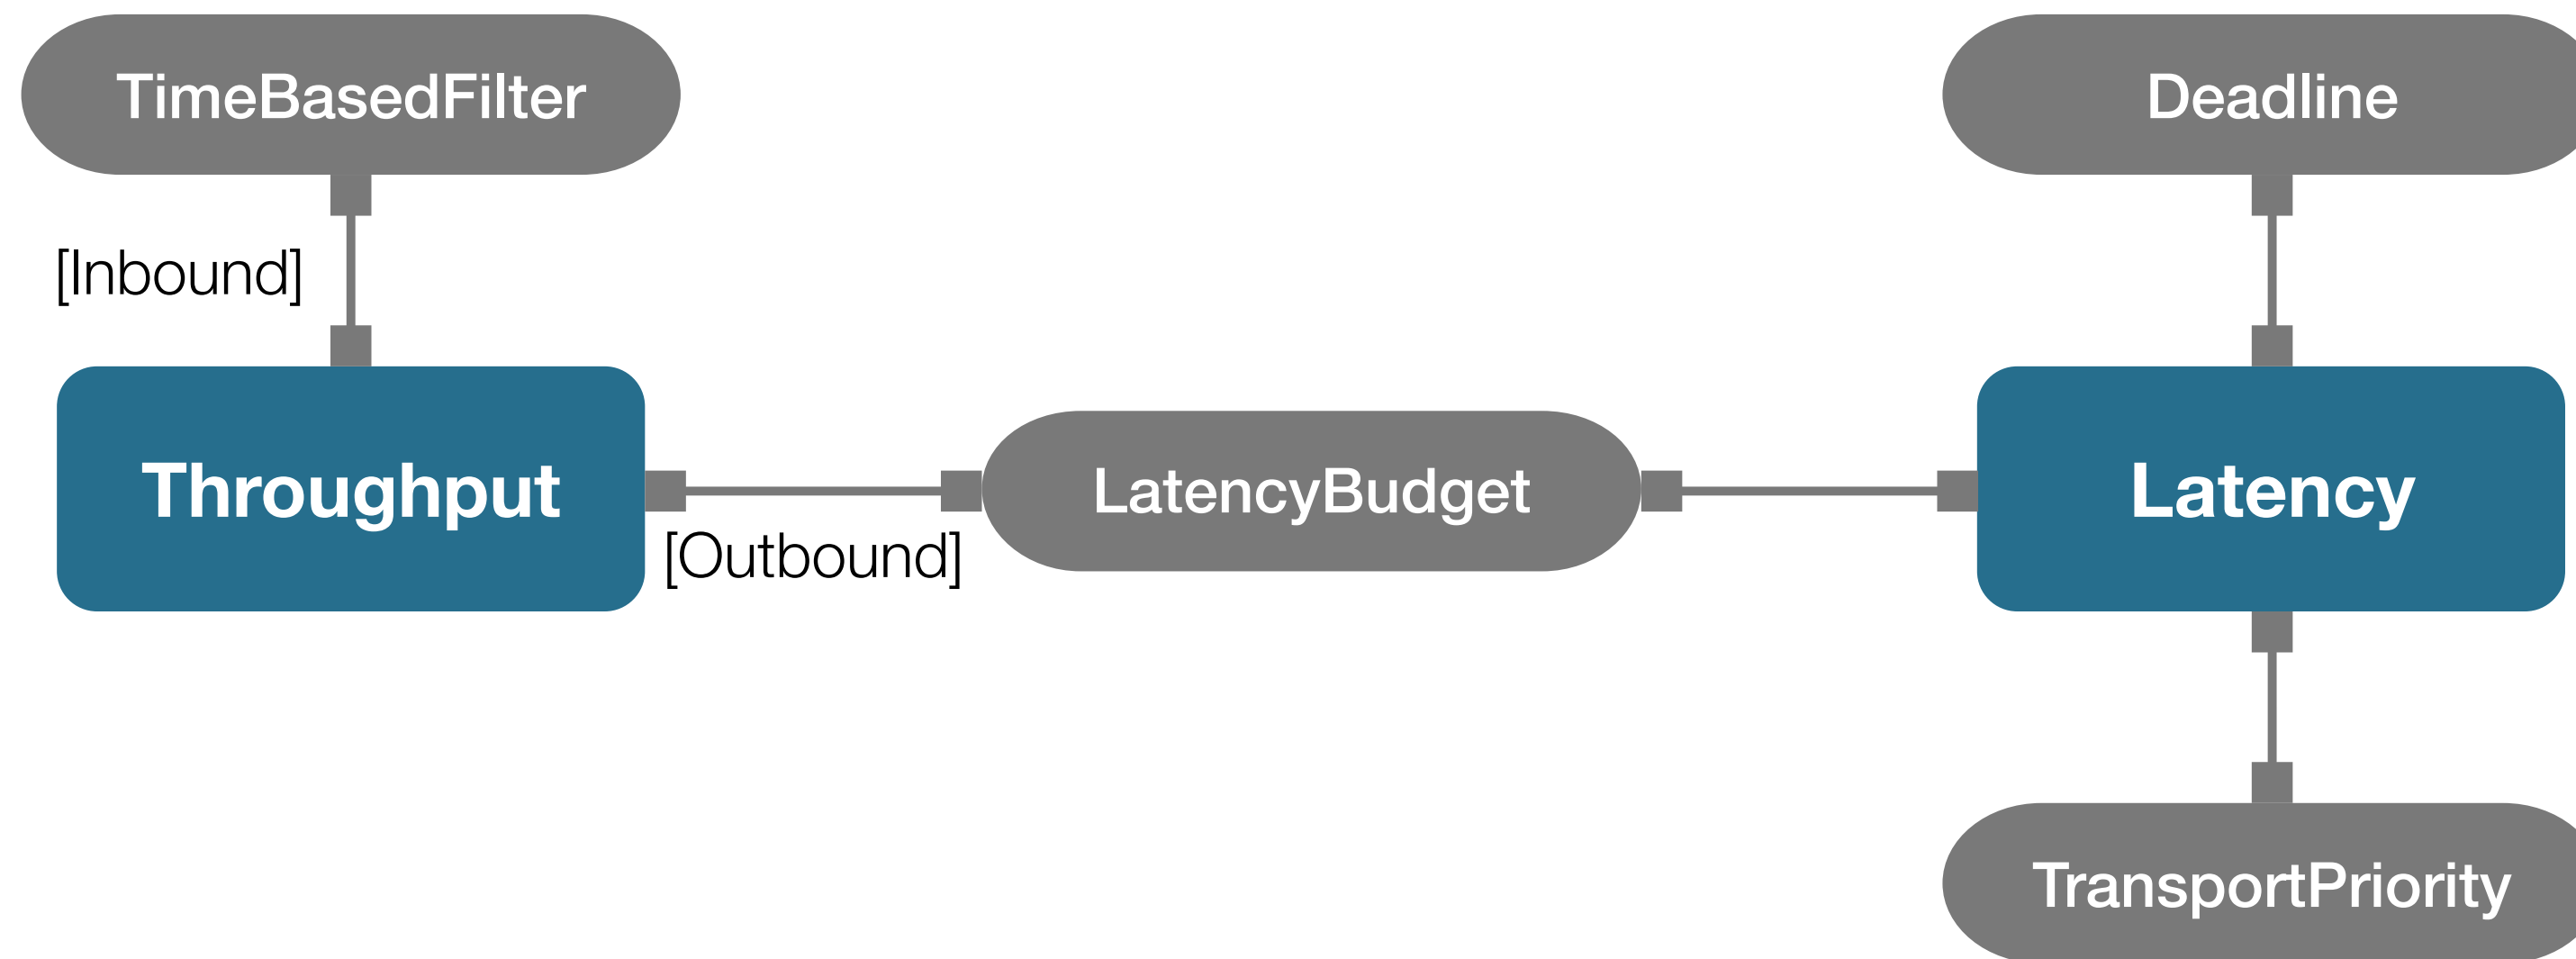
Happy Marriage

DDS is a perfect match for Robotics applications and frameworks, such as ROS2 because of:


- Data Centric Abstraction
- Decentralised Architecture
- Real-Time Behaviour
- Performance
- QoS-Driven non-functionals



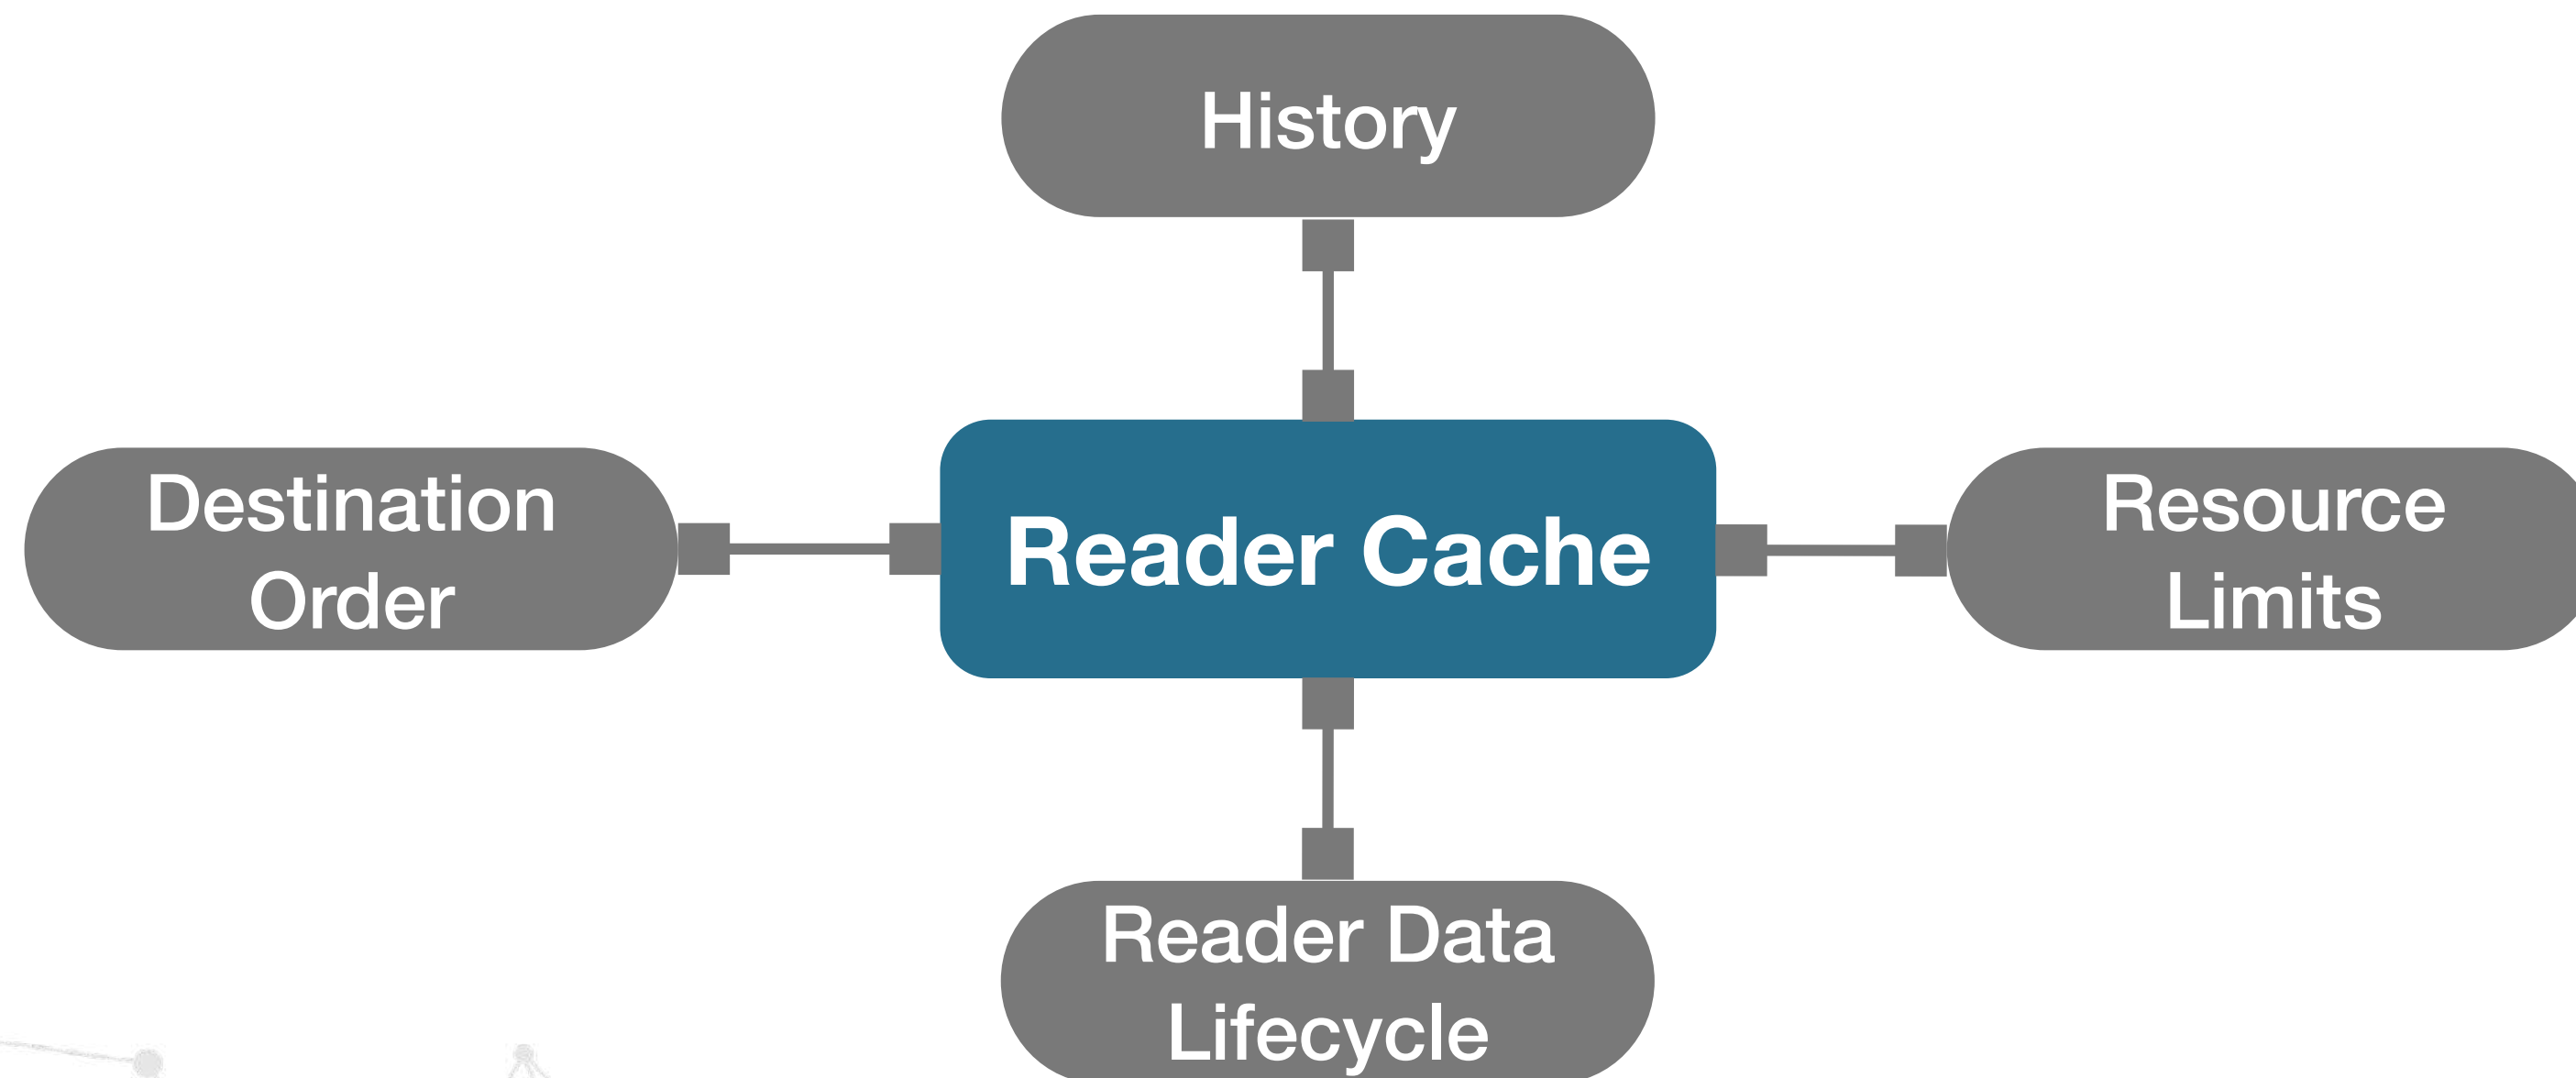
Temporal Properties



Data Availability



Reader Cache




QoS Modeling Idioms

State vs. Events

DDS provides first class support for modelling distributed state and events

Different QoS combination should be used to associate with the topic representing a state or an event the proper semantics



Soft State

In distributed systems you often need to model **soft-state** -- a **state that is periodically updated**

Examples are the reading of a sensor (e.g. Temperature Sensor), the position of a vehicle, etc.

The QoS combination to model **Soft-State** is the following:

Reliability	=>	BestEffort
Durability	=>	Volatile
History	=>	KeepLast(n) <i>[with n = 1 in most of the cases]</i>
Deadline	=>	updatePeriod
LatencyBudget	=>	updatePeriod/3 <i>[rule of thumb]</i>
DestinationOrder	=>	SourceTimestamp <i>[if multiple writers per instance]</i>

Hard State

In distributed systems you often need to model **hard-state** -- a **state** that is **sporadically updated** and that often has **temporal persistence requirements**

Examples are system configuration, a price estimate, etc.

The QoS combination to model Hard-State is the following:

Reliability	=>	Reliable
Durability	=>	Transient Persistent
History	=>	KeepLast(n) <i>[with n = 1 in most of the cases]</i>
DestinationOrder	=>	SourceTimestamp <i>[if multiple writers per instance]</i>
WriterDataLifecycle	=>	autodispose_unregistered_instances = false

Event

In distributed systems you often need to model **events** -- the **occurrence of something noteworthy for our system**

Examples are a collision alert, the temperature beyond a given threshold, etc.


The **QoS** combination to model Events is the following:

Reliability	=>	Reliable
Durability	=>	any <i>[depends on system requirements]</i>
History	=>	KeepAll <i>[on both DataWriter and DataReader!]</i>
DestinationOrder	=>	SourceTimestamp
WriterDataLifecycle	=>	autodispose_unregistered_instances = false
ResourceLimits	=>	<i>[define appropriate bounds]</i>

Performance


Cyclone Latency

Cyclone DDS has already
DDS implementation



Cyclone Throughput

Cyclone DDS has already
performing DDS imp



Lab 3: Performance Evaluation

C++ tsub.cpp X

```
code > cpp > 01 > C++ tsub.cpp > main(int, char * [])
19 using namespace std;
20 using namespace dds::core;
21 using namespace dds::core::policy;
22 using namespace dds::topic;
23 using namespace dds::pub;
24 using namespace dds::sub;
25
26 using namespace org::eclipse;
27
28 int
29 main(int argc, char* argv[])
30 {
31
32     try {
33         // Parse the command line args
34         tsub_options opt = parse_tsub_args(argc, argv);
35
36         // Create the domain participant
37         dds::domain::DomainParticipant dp(cyclonedds::domain::default_id());
38
39         // Initialize random number generation with a seed
40         srand(clock());
41
42         dds::topic::qos::TopicQos tqos =
43             dp.default_topic_qos() << LatencyBudget(Duration(2,0)) << Deadline(Duration(4,0));
44
45         auto topic = Topic<tutorial::TempSensorType>(dp, "TempSensorTopic", tqos);
46         auto pub = Publisher(dp);
47         auto dw = DataWriter<tutorial::TempSensorType>(pub, topic, tqos);
48
49     }
```

PROBLEMS 8 OUTPUT DEBUG CONSOLE TERMINAL

CMake/Build

```
[variant] Loaded new set of variants
[kit] Successfully loaded 1 kits from /Users/kydos/.local/share/CMakeTools/cmake-tools-kits.json
[main] Configuring folder: dds-tutorial
```


Summing Up

Concluding Remarks

DDS provides an extremely powerful set of abstractions and mechanism for data sharing in large scale distributed systems

DDS appears to be a natural fit for a large class of distributed and real-time applications, including robotics and autonomous driving

