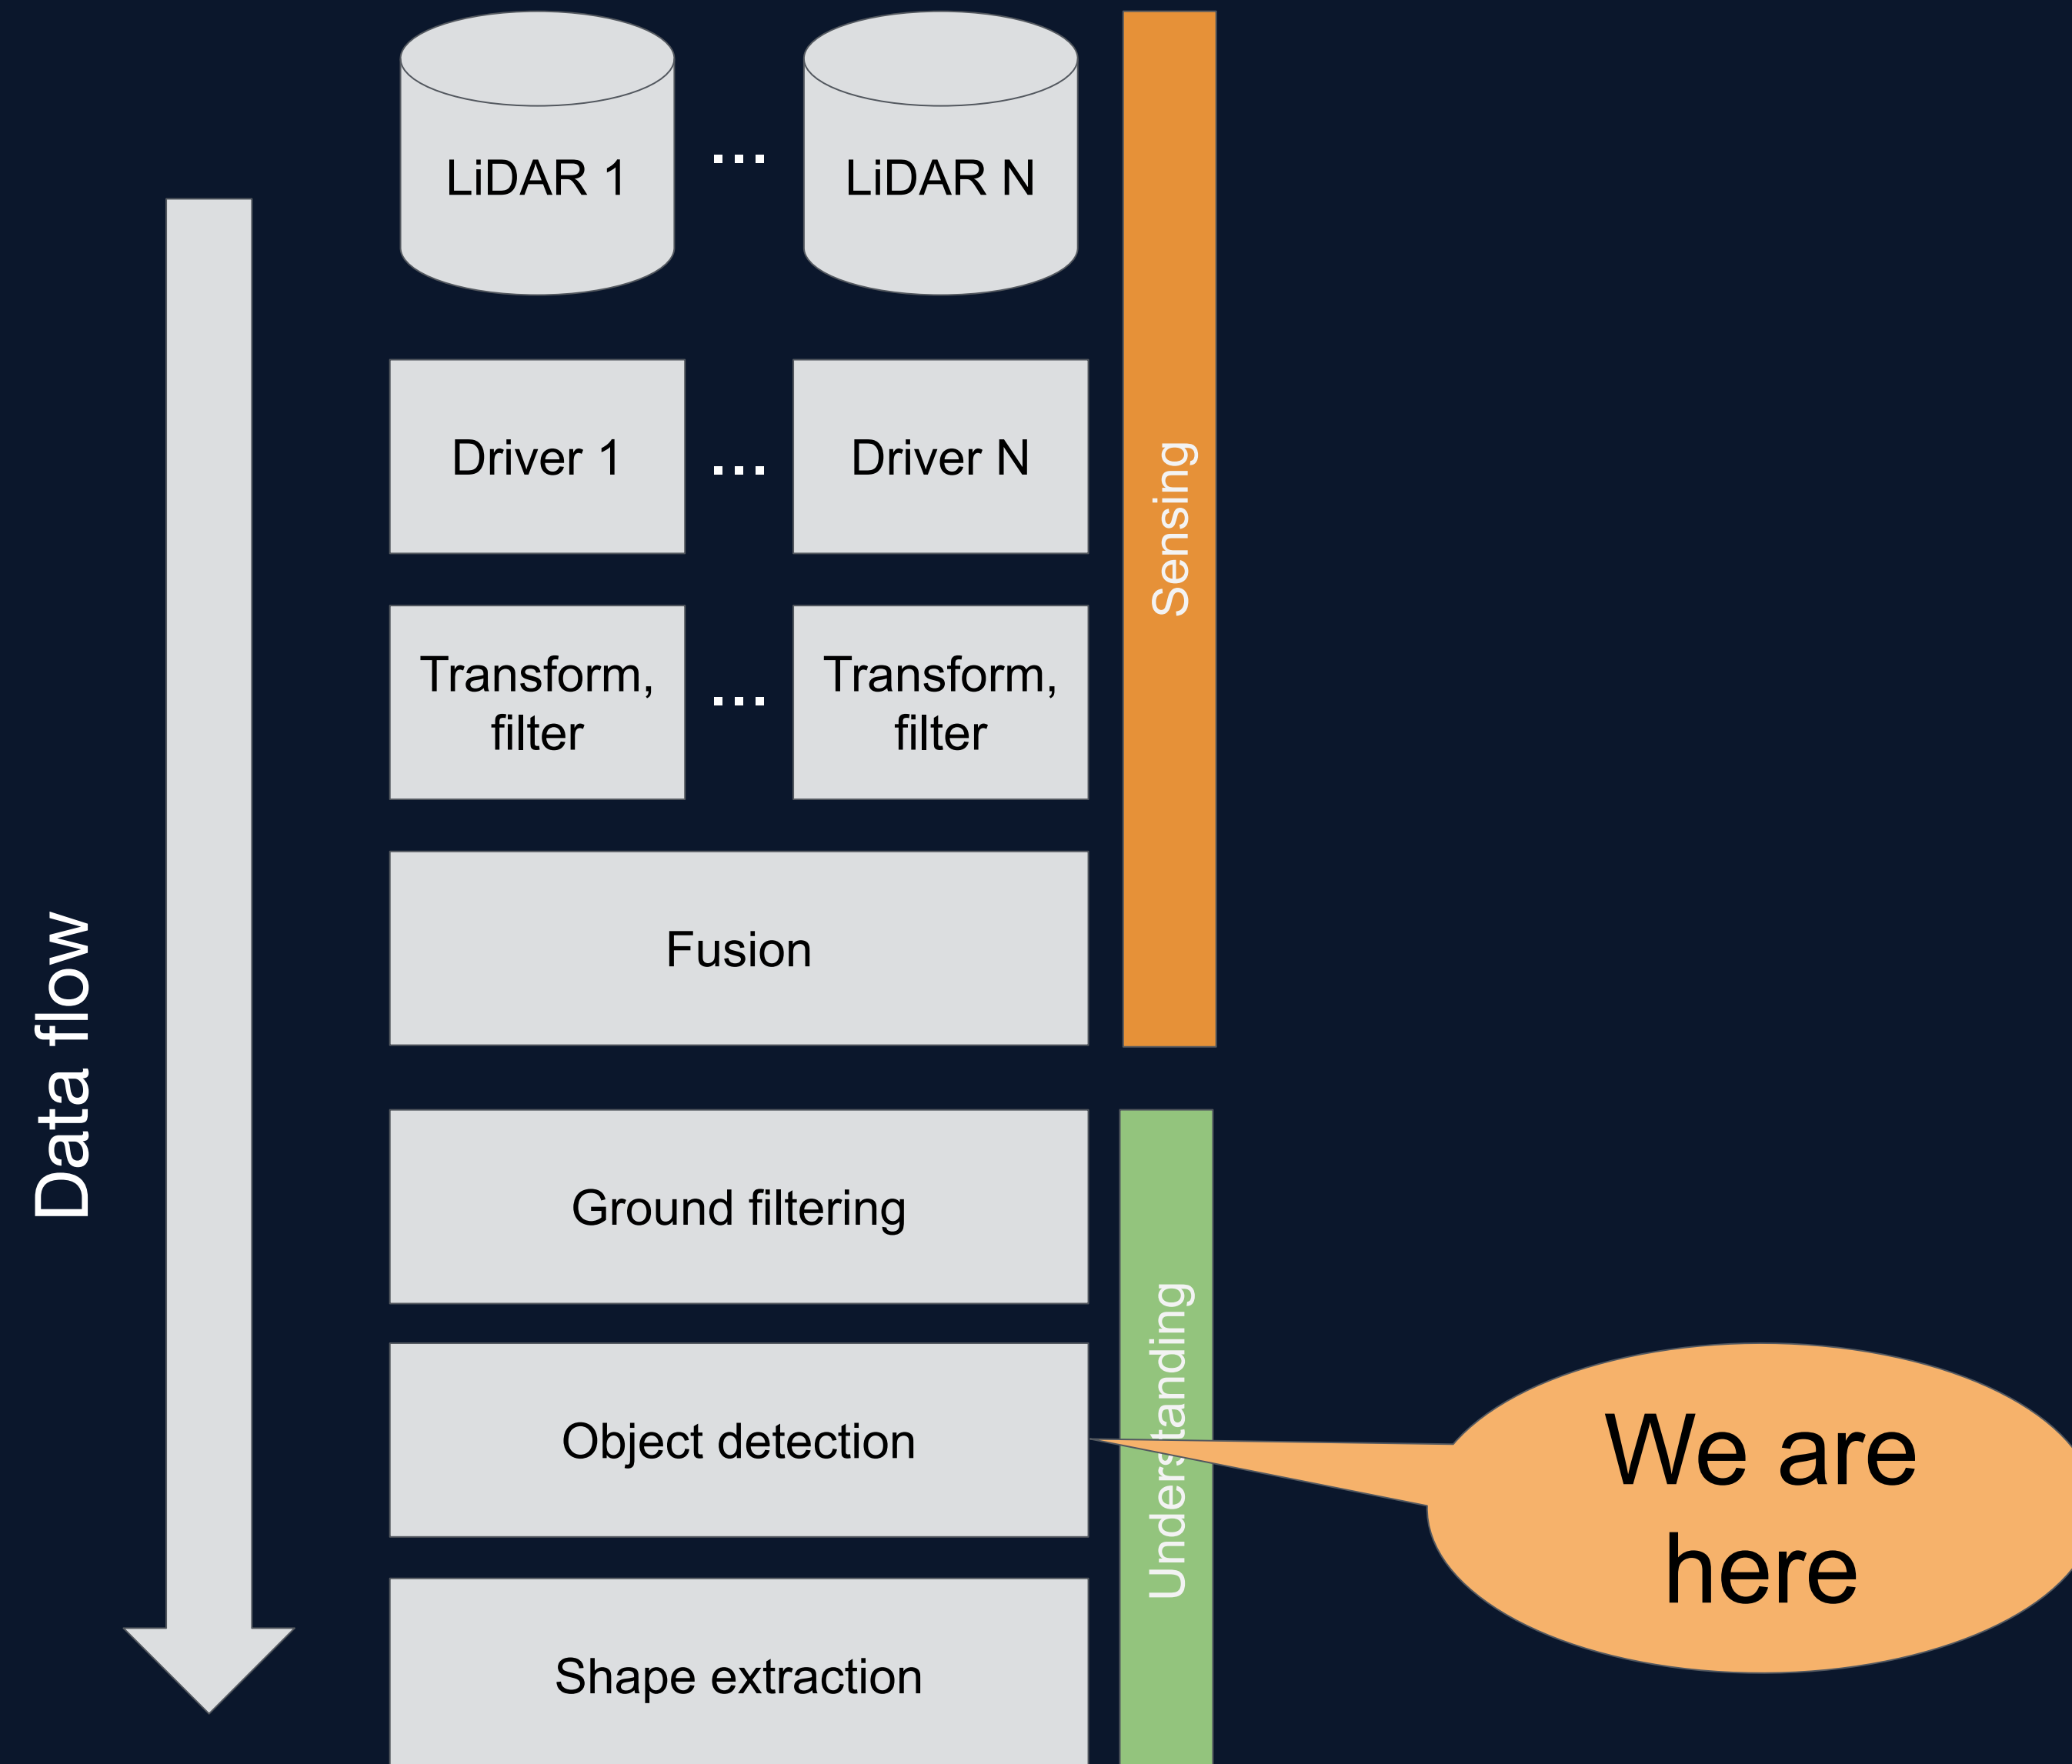




05 / Object Detection

Object Detection in the Classical LiDAR Processing Stack



The Problem of Object Detection

Why?

- Input nonground point cloud is cumbersome alone
- Need to discriminate between separate objects
 - e.g. *segment* nonground point cloud
- Partition point cloud into objects
- Remove some noise

How?

- Fundamentally group point together
 - Somehow
 - Based on some metric

Euclidean Clustering

One of a small handful of classical clustering/segmentation algorithms:

1. create a kd-tree representation for the input point cloud dataset P ;
2. set up an empty list of clusters C , and a queue of the points that need to be checked Q ;
3. then for every point $p_i \in P$, perform the following steps:
 - add p_i to the current queue Q ;
 - for every point $p_j \in Q$ do:
 - search for the set P_i^k of point neighbors of p_i in a sphere with radius $r < d_{th}$;
 - for every neighbor $p_j^k \in P_i^k$, check if the point has already been processed, and if not add it to Q ;
 - when the list of all points in Q has been processed, add Q to the list of clusters C , and reset Q to an empty list
4. the algorithm terminates when all points $p_i \in P$ have been processed and are now part of the list of point clusters C .

[1] Semantic 3D Object Maps for Everyday Manipulation in Human Living Environments, Rusu, 2009

[2] PCL

Two other ways to look at euclidean clustering

In other words:

1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point

Alternatively, the graph view

- Start with point cloud
- Each point is a vertex in the graph
- Connect each point *near* one another
- A cluster is a connected subgraph

Euclidean clustering, illustrated

In other words:


1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point



Euclidean clustering, illustrated

In other words:

1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point



Euclidean clustering, illustrated

In other words:

1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point




Euclidean clustering, illustrated

In other words:

1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point

- Unclassified
- Cluster 1, Final
- Cluster 1, New
- Cluster 2, Final
- Cluster 2, New




Euclidean clustering, illustrated

In other words:

1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point


- Unclassified
- Cluster 1, Final
- Cluster 1, New
- Cluster 2, Final
- Cluster 2, New



Euclidean clustering, illustrated

In other words:

1. Start with an empty cluster
2. Add some point to the cluster
3. Find all points *near* this point, add to cluster
4. Repeat (2) and (3) for each new point in the cluster
5. When there's no more points, accept/reject cluster based on number of points
6. Start a new cluster with a new point




Near, not nearest

Near:

- All points with distance less than some threshold
- Different from nearest
- Different from k-nearest
- Best data structure for (k-nearest) is k-d tree
 - But not for near-neighbor lookups


Near



Nearest:

- Use k-d tree
- $O(\log N)$ for lookup

Nearest



Using the right data structure

Integer lattice/Spatial hash

1. Subdivide space into voxels: $O(n)$
2. Find voxel for query point: $O(1)$
3. Find all voxels a *near* point could possibly fall in: $O(1)$
4. Iterate over all points in all *near* voxels: $O(k)$
 - Add to output each point that is *near* the query point


Average query complexity: $O(1) * O(1) * O(k) = O(k)$

Let n be total number of points, k be the average number of neighbors

Even has its own [wiki page](#)


The spatial hash, illustrated

1. Subdivide space into voxels: $O(n)$
2. Find voxel for query point: $O(1)$
3. Find all voxels a *near* point could possibly fall in: $O(1)$
4. Iterate over all points in all *near* voxels: $O(k)$
 - Add to output each point that is *near* the query point




The spatial hash, illustrated

1. Subdivide space into voxels: $O(n)$
2. Find voxel for query point: $O(1)$
3. Find all voxels a *near* point could possibly fall in: $O(1)$
4. Iterate over all points in all *near* voxels: $O(k)$
 - Add to output each point that is *near* the query point




The spatial hash, illustrated

1. Subdivide space into voxels: $O(n)$
2. Find voxel for query point: $O(1)$
3. Find all voxels a *near* point could possibly fall in: $O(1)$
4. Iterate over all points in all *near* voxels: $O(k)$
 - Add to output each point that is *near* the query point




The spatial hash, illustrated

1. Subdivide space into voxels: $O(n)$
2. Find voxel for query point: $O(1)$
3. Find all voxels a *near* point could possibly fall in: $O(1)$
4. Iterate over all points in all *near* voxels: $O(k)$
 - Add to output each point that is *near* the query point



The spatial hash, illustrated

1. Subdivide space into voxels: $O(n)$
2. Find voxel for query point: $O(1)$
3. Find all voxels a *near* point could possibly fall in: $O(1)$
4. Iterate over all points in all *near* voxels: $O(k)$
 - Add to output each point that is *near* the query point



(Average-case) Complexity Analysis

$O(n)$

$O(1)$

1. create a kd-tree representation for the input point cloud dataset P ;
2. set up an empty list of clusters C , and a queue of the points that need to be checked Q ;
3. then for every point $p_i \in P$, perform the following steps:
 - add p_i to the current queue Q ;
 - for every point $p_i \in Q$ do:
 - search for the set P_i^k of point neighbors of p_i in a sphere with radius $r < d_{th}$;
 - for every neighbor $p_i^k \in P_i^k$, check if the point has already been processed, and if not add it to Q ;
 - when the list of all points in Q has been processed, add Q to the list of clusters C , and reset Q to an empty list
4. the algorithm terminates when all points $p_i \in P$ have been processed and are now part of the list of point clusters C .

$O(1)$

$O(n(k + Q(k)))$

$O(k + Q(k))$

$O(Q(k))$

$O(k)$

$O(1)$

$O(1)$

Total complexity: $O(n + n(k + Q(k))) = O(n(k + Q(k)))$

Original algorithm: $Q(k) = k \log n \rightarrow O(kn(1 + \log n))$

Integer lattice: $Q(k) = k \rightarrow O(kn)$

If k is a problem defined constant, we went from
linearithmic to linear!

Geometric Object Detection - Summary

- Object detection algorithms are needed to distinguish individual collidable objects
- Most object detection algorithms use some kind of region growing method
 - On voxels
 - In euclidean space
 - In angles in the range-image space
- Autoware.Auto uses a version of euclidean clustering

Important optimizations

- Using the right data structure:
 - $O(n \log n) \rightarrow O(n)$
 - Important since n can get big: 10k-1m!
- Concretely:
 - Autoware.ai: Barely 100ms runtime with aggressive downsampling
 - Autoware.Auto: Comfortable 10ms runtime with no* downsampling