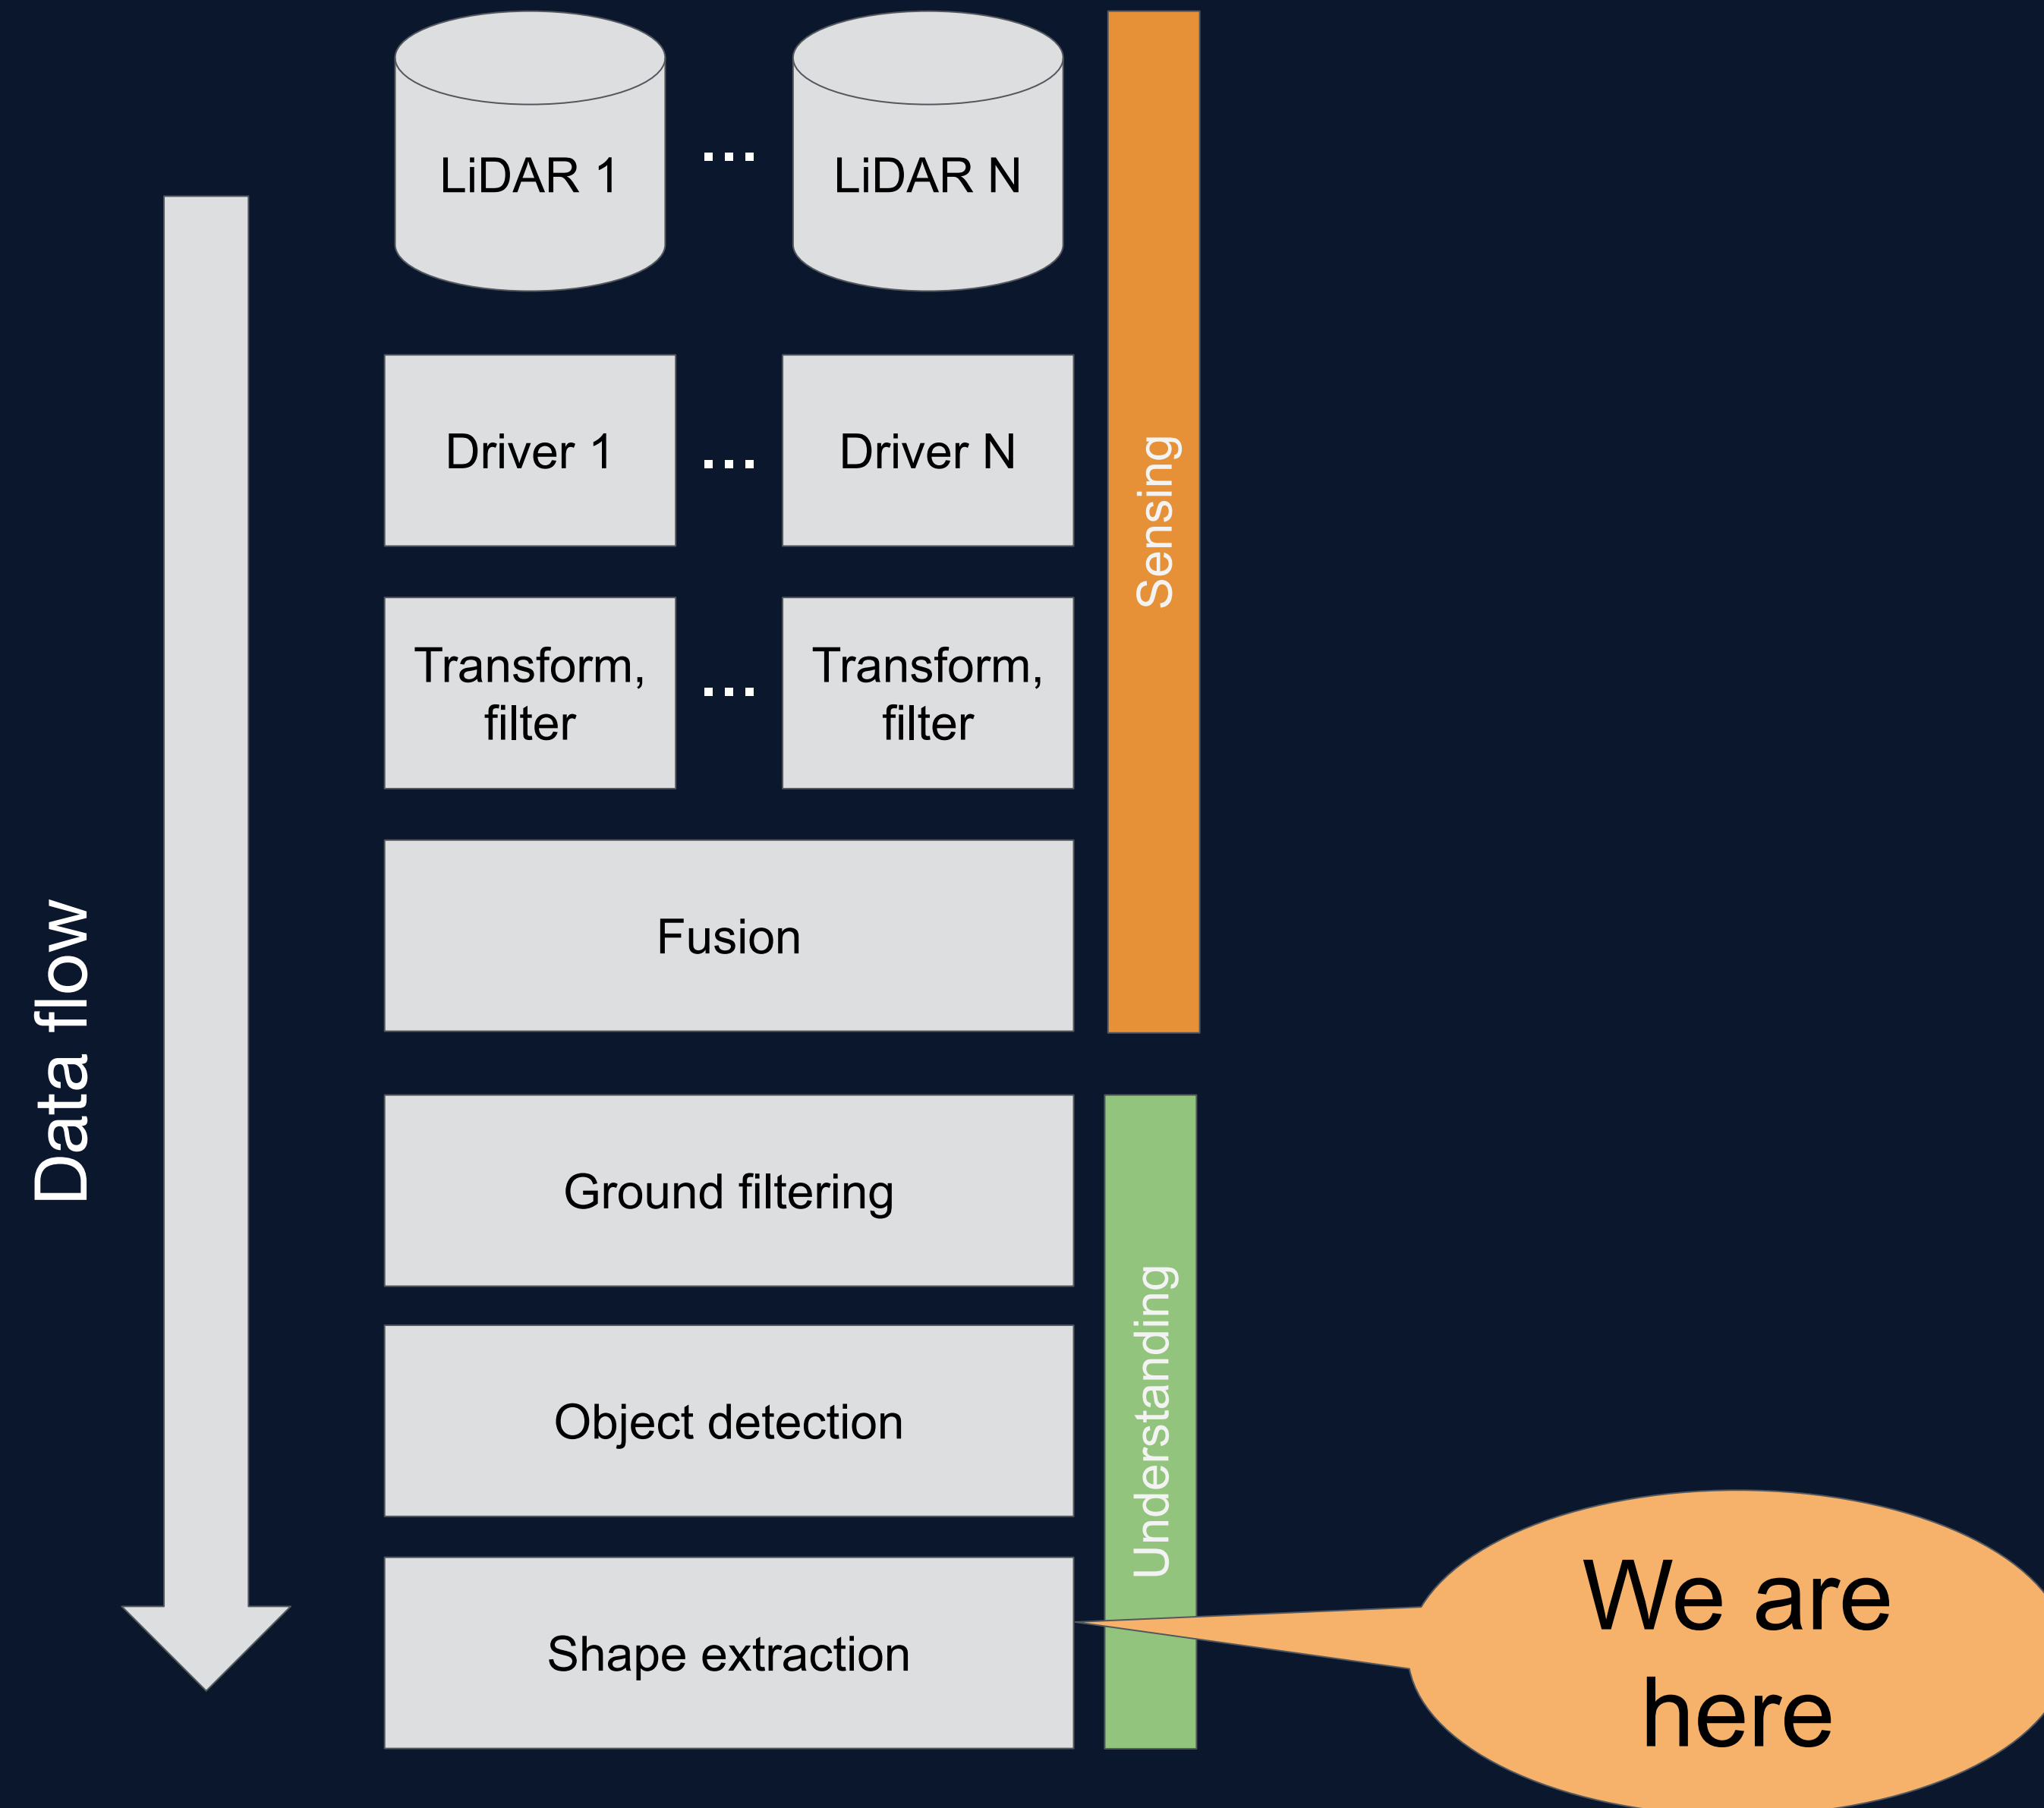




06 / Shape Extraction

Shape Extraction in the Classical LiDAR Processing Stack



An overview on shape extraction techniques

- Why?
- Different ways to compute bounding boxes:
 - Rotating calipers
 - Eigenvalues
 - Optimization
- Other shape extraction techniques

Why not use point blobs directly?

Recall sufficient statistics:

- We only want the minimum of information to produce the right results
- Point blobs are big and unbounded
 - Additional communication and memory overhead
 - Additional computational overhead for collision checking, other use cases

Shape extraction and the collision detection use case

Object detection is generally to inform collision detection

The final object representation should:

- be easy to collision check
- be easy to compute
- reasonably represent the actual object

Oriented bounding boxes are a good compromise

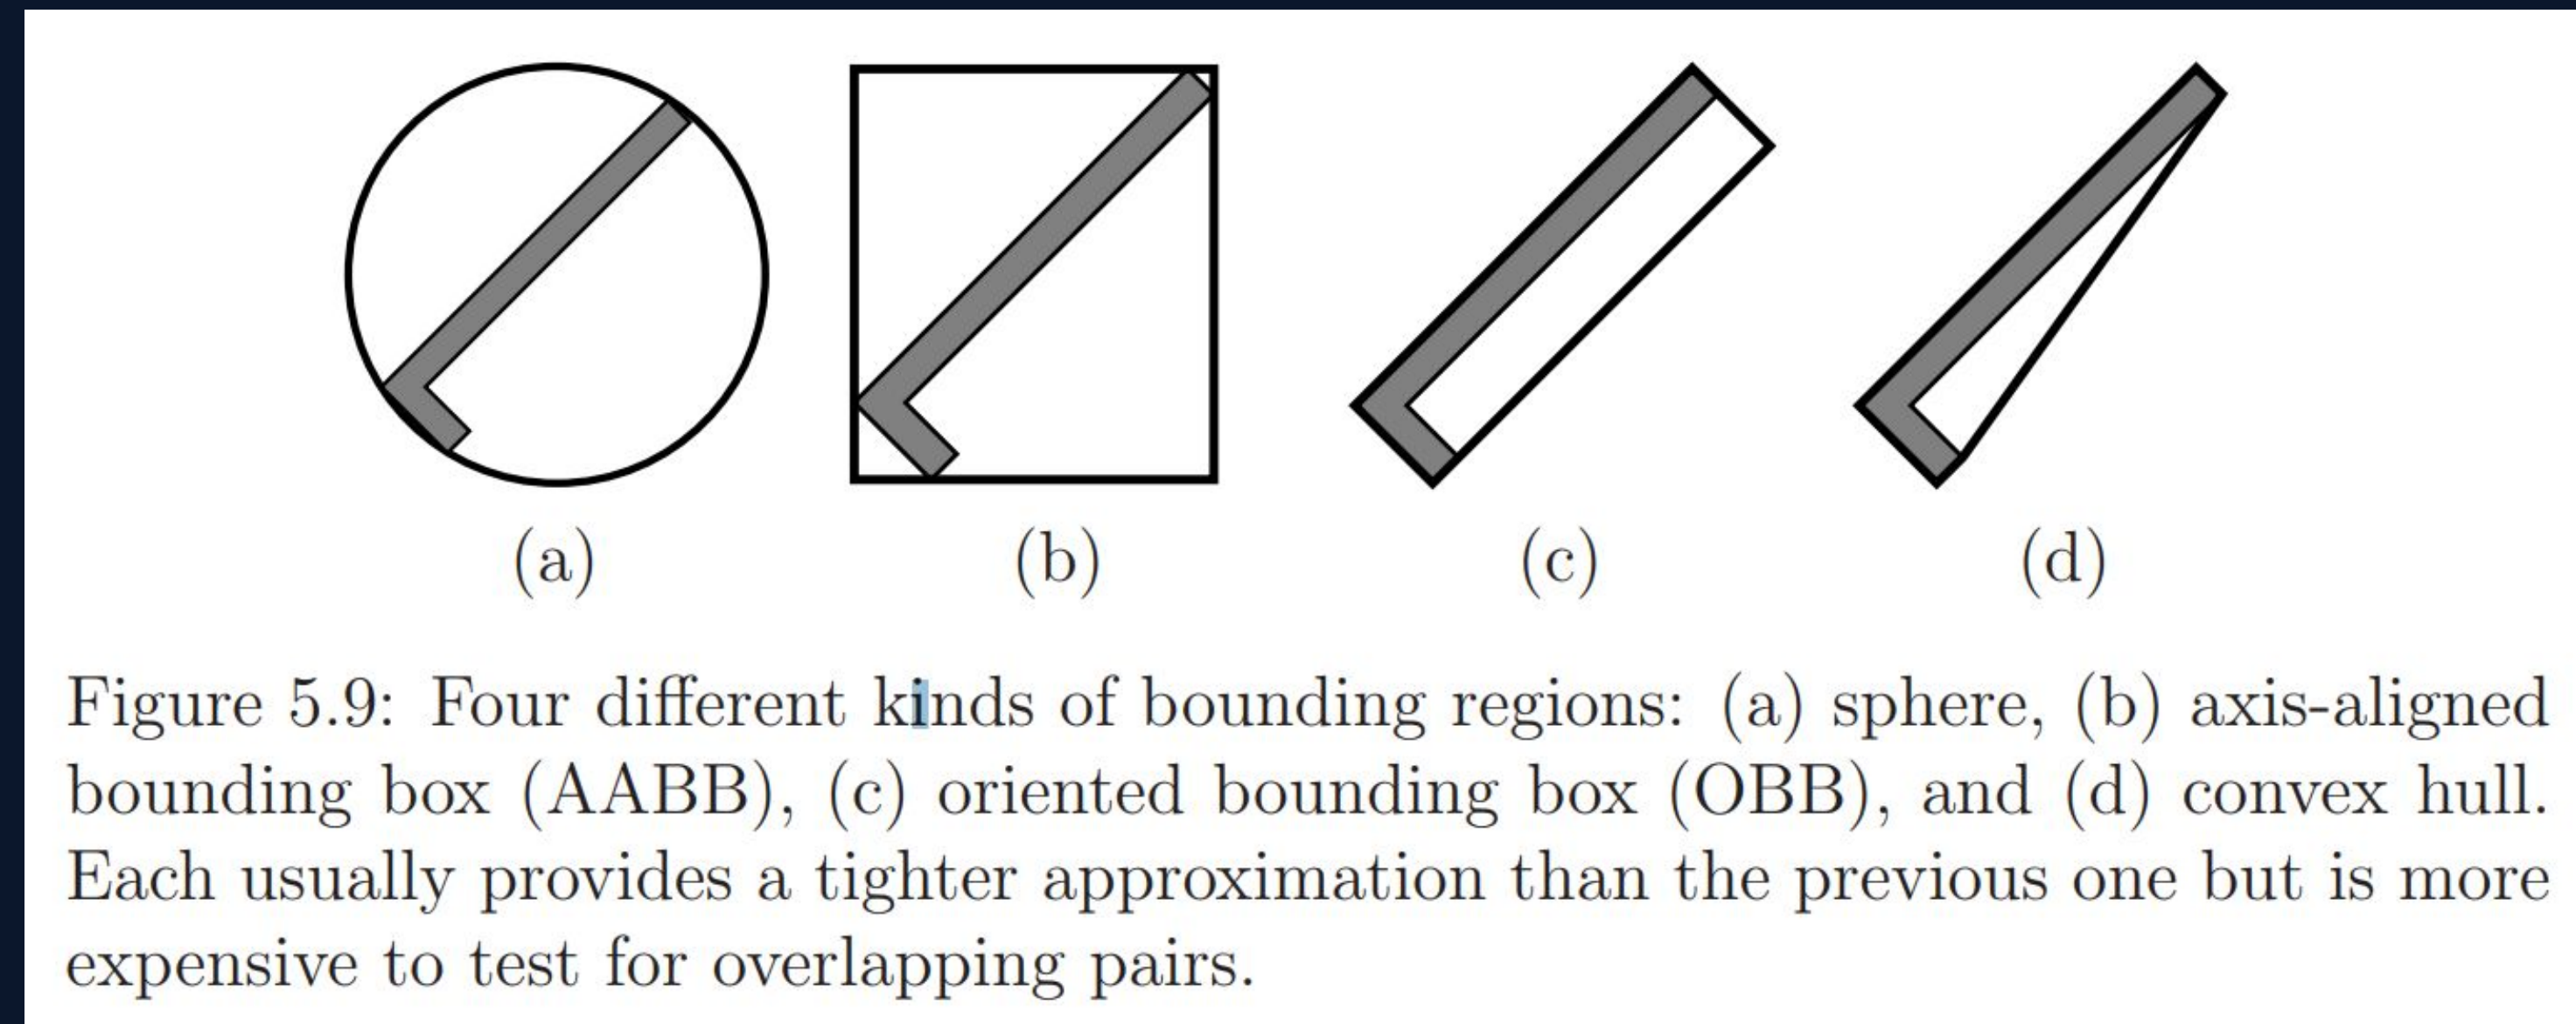




Image credit:

Planning Algorithms, Ch 5, Steven LaValle


Minimum Bounding Boxes



Points



Convex Hull



Bounding Box

- Rotating calipers: compute 2D bounding boxes with minimum perimeter or area in linear time
 - Requires a convex hull (computable in $O(n \log n)$ time)

Minimum Bounding Boxes

...but:

- only considers boundary points
 - approach is sensitive to noise
- it's trivially easy to construct examples that won't semantically fit data


Minimum Bounding Boxes

...but:


- only considers boundary points
 - approach is sensitive to noise
- it's trivially easy to construct examples that won't semantically fit data, e.g.:



Points



Convex Hull



Bounding Box

Eigenbox

Idea: use PCA on all points to compute major and minor axes of bounding box

- Linear in time
- Uses information from entire object
- Bounded suboptimality guarantees

Eigenbox

Idea: use PCA on all points to compute major and minor axes of bounding box

- Linear in time
- Uses information from entire object
- Bounded suboptimality guarantees
- ...but still doesn't produce semantically correct results



Optimal Bounding Boxes

Can instead explicitly solve an optimization problem

- Shen et al
 - Box and partition that best fits L-shape
 - ...but relies on a single scanning LiDAR
- Zhang et al
 - Relaxes assumption of single scanning LiDAR
 - ...but introduces discretization error

$$\begin{aligned} \underset{P, \theta, c_1, c_2}{\text{minimize}} \quad & \sum_{i \in P} (x_i \cos \theta + y_i \sin \theta - c_1)^2 \\ & + \sum_{i \in Q} (-x_i \sin \theta + y_i \cos \theta - c_2)^2 \quad (1) \\ \text{subject to} \quad & P \cup Q = \{1, 2, \dots, m\} \\ & c_1, c_2 \in R \quad 0^\circ \leq \theta < 90^\circ \end{aligned}$$

Bounding Boxes in Autoware.Auto

Use approach from Shen et al, but:

- Compute principal component of blob
- Sort points along this axis
- Extra work, but:
 - No discretization error
 - Works on arbitrary point clouds



Further Approaches

Many other ways to do shape extraction

- Convex hulls
 - Only loosely bounded memory complexity
- Minimum Volume Bounding Ellipsoids
 - e.g. Kumar and Yildirim
 - Collision detection requires solving an optimization problem
- Super quadratics (i.e. Pascoal et al):
 - Handles concave shapes
 - Requires solving an optimization problem

Some open problems

- Handling large, concave objects, e.g.
 - Inside of a small garage

Shape Extraction - Summary

Shape extraction is necessary to reduce complexity:

- Communication
- Semantic
- Computational

There are many approaches:

- Bounding boxes
- Ellipsoids
- Convex hulls
- Others

And there are still some open problems:

- Large concave objects